

Big Data Computing

Master's Degree in Computer Science

2020-2021

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

tolomei@di.uniroma1.it



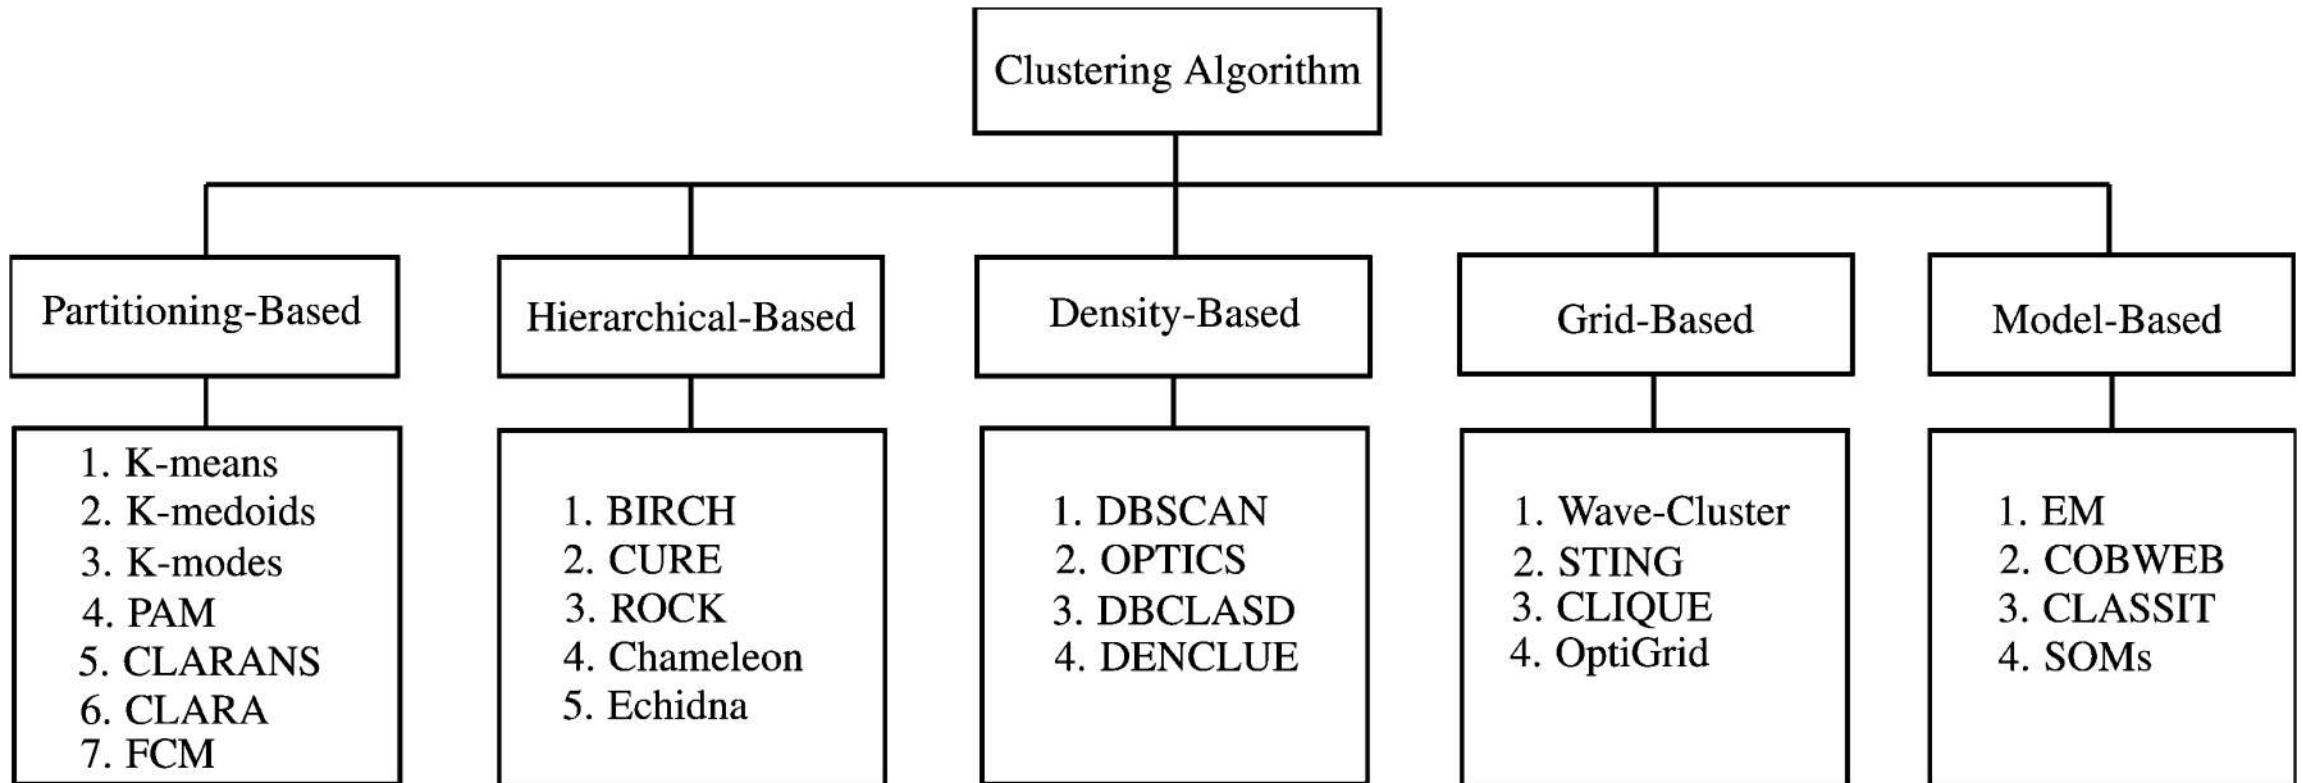
SAPIENZA
UNIVERSITÀ DI ROMA

Recap from Last Lecture(s)

- Clustering is an unsupervised learning technique to group "similar" data objects together
- Depends on:
 - object representation
 - similarity measure
- Harder when data dimensionality gets large (**curse of dimensionality**)
- Number of output clusters is part of the problem itself!

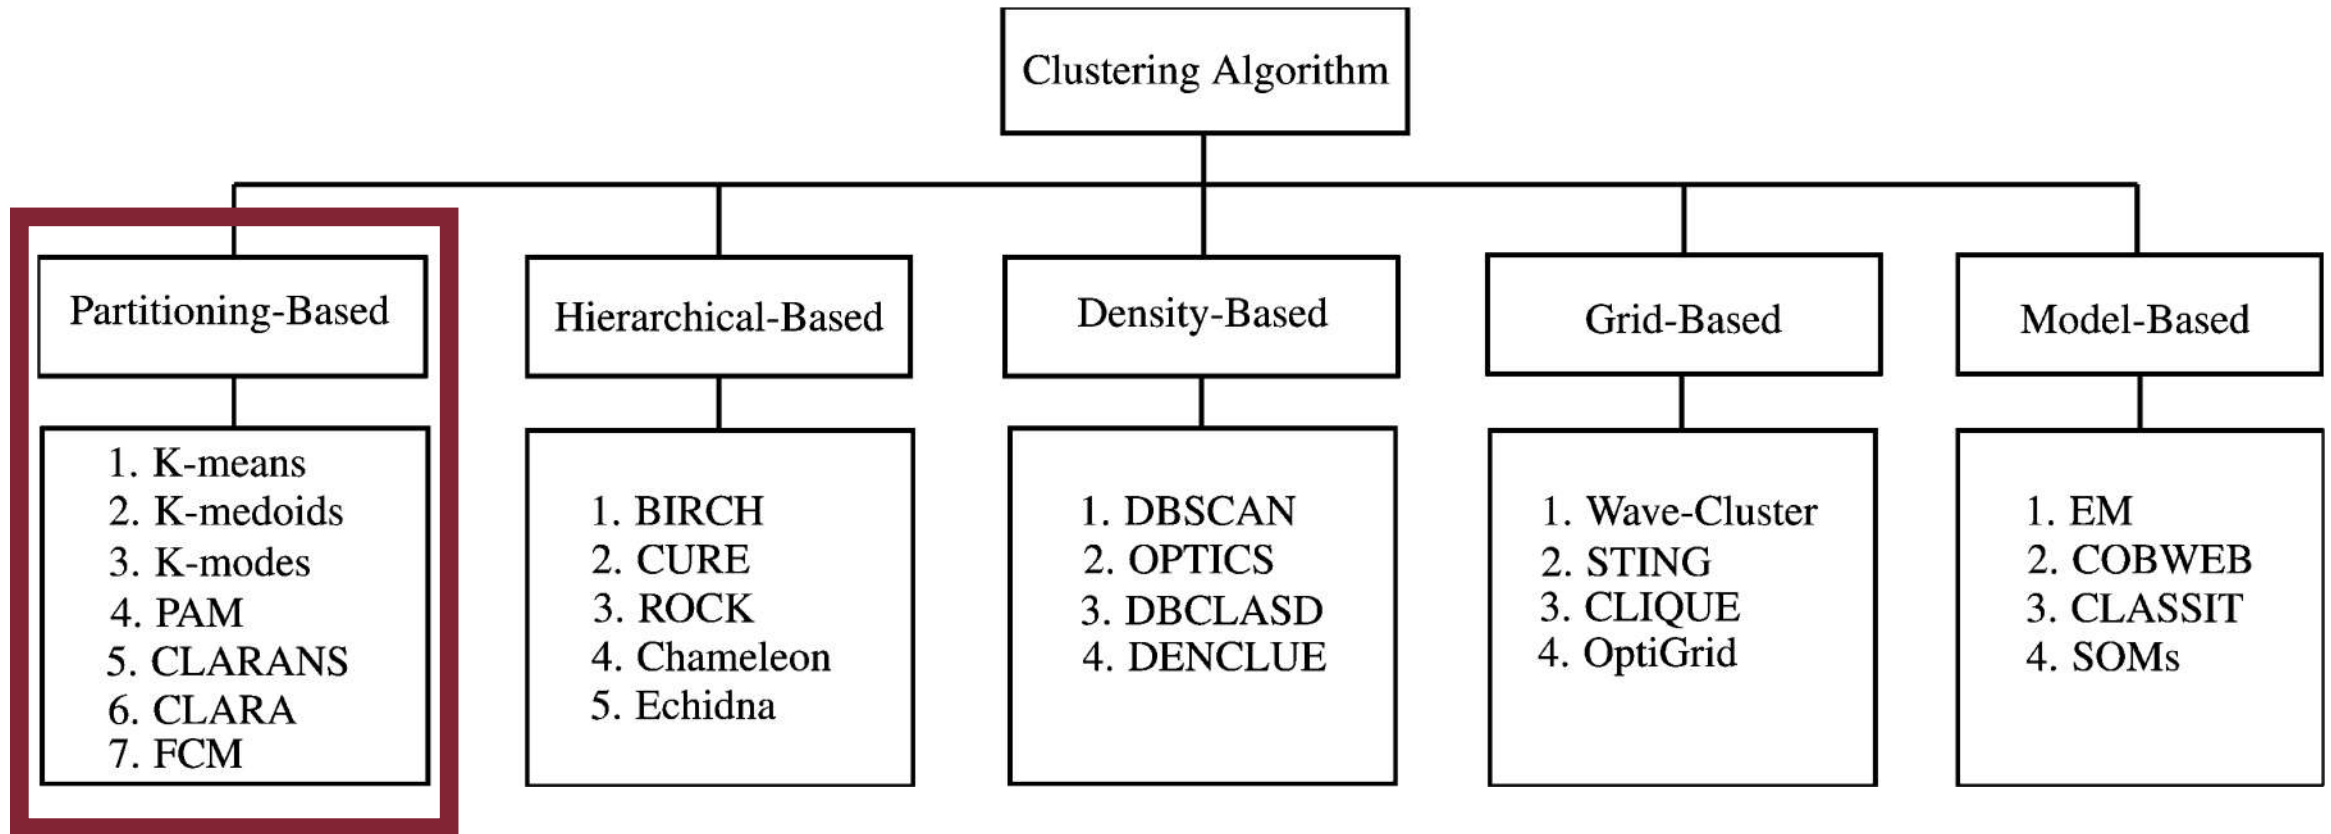
Clustering Algorithms

Clustering Algorithms: Taxonomy



source: <https://www.computer.org/csdl/journal/ec/2014/03/06832486/13rRUEgs2xB>

Clustering Algorithms: Taxonomy



source: <https://www.computer.org/csdl/journal/ec/2014/03/06832486/13rRUEgs2xB>

Partitioning: Hard Clustering

- **Input:** A set of N data points and a number K ($K < N$)

Partitioning: Hard Clustering

- **Input:** A set of N data points and a number K ($K < N$)
- **Output:** A partition of the N data points into K clusters

Partitioning: Hard Clustering

- **Input:** A set of N data points and a number K ($K < N$)
- **Output:** A partition of the N data points into K clusters
- **Goal:** Find the partition which optimizes a certain criterion

Partitioning: Hard Clustering

- **Input:** A set of N data points and a number K ($K < N$)
- **Output:** A partition of the N data points into K clusters
- **Goal:** Find the partition which optimizes a certain criterion
 - Global optimum \rightarrow Intractable for many objective function (might require to enumerate all the possible partitions)*

*Kleinberg, J., "An Impossibility Theorem for Clustering" (NIPS 2002)

Partitioning: Hard Clustering

- **Input:** A set of N data points and a number K ($K < N$)
- **Output:** A partition of the N data points into K clusters
- **Goal:** Find the partition which optimizes a certain criterion
 - Global optimum \rightarrow Intractable for many objective function (might require to enumerate all the possible partitions)*
 - $S(K, N) \sim O(K^N) \rightarrow$ K-way non-empty partitions of N elements

Stirling partition number

*Kleinberg, J., "An Impossibility Theorem for Clustering" (NIPS 2002)

Partitioning: Hard Clustering

- **Input:** A set of N data points and a number K ($K < N$)
- **Output:** A partition of the N data points into K clusters
- **Goal:** Find the partition which optimizes a certain criterion
 - Global optimum \rightarrow Intractable for many objective function (might require to enumerate all the possible partitions)*
 - $S(K, N) \sim O(K^N) \rightarrow$ K-way non-empty partitions of N elements
 - Effective heuristics \rightarrow K-means, K-medoids, K-means++, etc.

Stirling partition number

*Kleinberg, J., "An Impossibility Theorem for Clustering" (NIPS 2002)

Flat Hard Clustering: General Framework

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ the set of N input data points
 $\{C_1, \dots, C_K\}$ the set of K output clusters
 C_k the generic k -th cluster
 $\boldsymbol{\theta}_k$ is the *representative* of cluster C_k

Flat Hard Clustering: General Framework

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ the set of N input data points
 $\{C_1, \dots, C_K\}$ the set of K output clusters
 C_k the generic k -th cluster
 $\boldsymbol{\theta}_k$ is the *representative* of cluster C_k

Note:

At this stage we haven't yet specified what a cluster representative actually is

Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$$

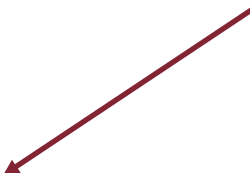
where:

- A is an $N \times K$ matrix s.t. $\alpha_{n,k} = 1$ iff \mathbf{x}_n is assigned to cluster C_k , 0 otherwise
- $\Theta = \{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K\}$ are the cluster representatives
- $\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$ is a function measuring the distance between \mathbf{x}_n and $\boldsymbol{\theta}_k$

Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \theta_k)$$

$\sum_{k=1}^K \alpha_{n,k} = 1$ hard clustering



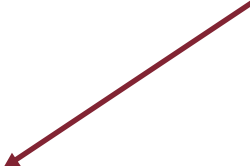
where:

- A is an $N \times K$ matrix s.t. $\alpha_{n,k} = 1$ iff \mathbf{x}_n is assigned to cluster C_k , 0 otherwise
- $\Theta = \{\theta_1, \dots, \theta_K\}$ are the cluster representatives
- $\delta(\mathbf{x}_n, \theta_k)$ is a function measuring the distance between \mathbf{x}_n and θ_k

Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \theta_k)$$

$\sum_{k=1}^K \alpha_{n,k} = 1$ hard clustering



where:

- A is an $N \times K$ matrix s.t. $\alpha_{n,k} = 1$ iff \mathbf{x}_n is assigned to cluster C_k , 0 otherwise
- $\Theta = \{\theta_1, \dots, \theta_K\}$ are the cluster representatives
- $\delta(\mathbf{x}_n, \theta_k)$ is a function measuring the distance between \mathbf{x}_n and θ_k

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \theta_k)}_{L(A, \Theta)}$$

Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

exact solution must explore
exponential search space
 $S(K, N) \sim O(K^N)$



NP-hard

Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

exact solution must explore
exponential search space
 $S(K, N) \sim O(K^N)$



NP-hard

non-convex due to the discrete
assignment matrix A



multiple local minima

Iterative Solution: Lloyd-Forgy Algorithm

- NP-hardness doesn't allow us to compute the exact solution

Iterative Solution: Lloyd-Forgy Algorithm

- NP-hardness doesn't allow us to compute the exact solution
- Non-convexity doesn't allow us to rely on nice property of convex optimization with numerical methods (unique global minimum)

Iterative Solution: Lloyd-Forgy Algorithm

- NP-hardness doesn't allow us to compute the exact solution
- Non-convexity doesn't allow us to rely on nice property of convex optimization with numerical methods (unique global minimum)
- **Lloyd-Forgy Algorithm**: 2-step **iterative** approximated solution
 - Assignment step
 - Update step

Iterative Solution: Lloyd-Forgy Algorithm

- NP-hardness doesn't allow us to compute the exact solution
- Non-convexity doesn't allow us to rely on nice property of convex optimization with numerical methods (unique global minimum)
- **Lloyd-Forgy Algorithm**: 2-step **iterative** approximated solution
 - Assignment step
 - Update step

Does not guarantee to find the global optimum as it may stuck to a local optimum or a saddle point

2-Step Optimization: Assignment Step

Minimize L w.r.t. A by fixing Θ

$L(\Theta; A)$ fixed Θ parametrized by A

2-Step Optimization: Assignment Step

Minimize L w.r.t. A by fixing Θ

$L(\Theta; A)$ fixed Θ parametrized by A

Note:

Can't take the gradient of L w.r.t. A
since A is discrete!

2-Step Optimization: Assignment Step

Minimize L w.r.t. A by fixing Θ

$L(\Theta; A)$ fixed Θ parametrized by A

Intuitively, given a set of fixed representatives, L is minimized if each data point is assigned to the closest centroid according to δ

(L is just the summation of all the distances from each data point to its assigned representative)

2-Step Optimization: Assignment Step

Minimize L w.r.t. A by fixing Θ

$L(\Theta; A)$ fixed Θ parametrized by A

Intuitively, given a set of fixed representatives, L is minimized if each data point is assigned to the closest centroid according to δ

(L is just the summation of all the distances from each data point to its assigned representative)

$$\alpha_{n,k} = \begin{cases} 1 & \text{if } \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) = \min_{1 \leq j \leq K} \{\delta(\mathbf{x}_n, \boldsymbol{\theta}_j)\} \\ 0 & \text{otherwise} \end{cases}$$

2-Step Optimization: Update Step

Minimize L w.r.t. Θ by fixing A

$L(A; \Theta)$ fixed A parametrized by Θ

2-Step Optimization: Update Step

Minimize L w.r.t. Θ by fixing A

$L(A; \Theta)$ fixed A parametrized by Θ

We can minimize L by taking the **gradient** of L w.r.t Θ (i.e., the vector of partial derivatives), set it to 0 and solve it for Θ

2-Step Optimization: Update Step

$$\nabla L(A; \Theta) = \left(\frac{\partial L(A; \Theta)}{\partial \theta_1}, \dots, \frac{\partial L(A; \Theta)}{\partial \theta_K} \right)$$

2-Step Optimization: Update Step

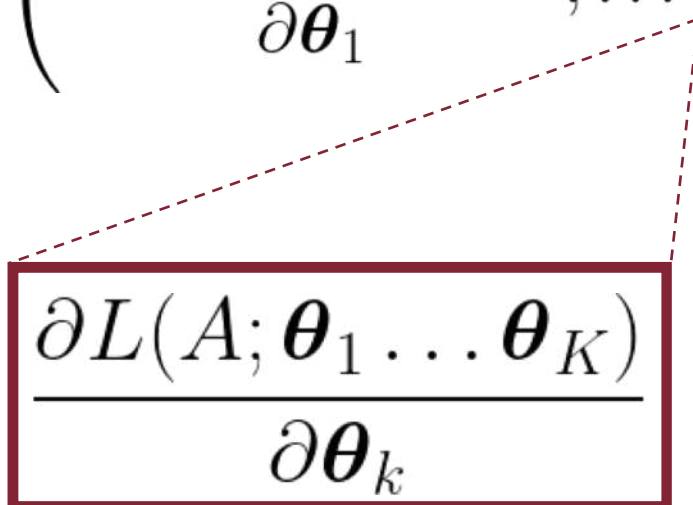
$$\nabla L(A; \Theta) = \left(\frac{\partial L(A; \Theta)}{\partial \theta_1}, \dots, \frac{\partial L(A; \Theta)}{\partial \theta_K} \right)$$

$$\nabla L(A; \Theta) = \left(\frac{\partial L(A; \theta_1 \dots \theta_K)}{\partial \theta_1}, \dots, \frac{\partial L(A; \theta_1 \dots \theta_K)}{\partial \theta_K} \right)$$

2-Step Optimization: Update Step

$$\nabla L(A; \Theta) = \left(\frac{\partial L(A; \Theta)}{\partial \theta_1}, \dots, \frac{\partial L(A; \Theta)}{\partial \theta_K} \right)$$

$$\nabla L(A; \Theta) = \left(\frac{\partial L(A; \theta_1 \dots \theta_K)}{\partial \theta_1}, \dots, \frac{\partial L(A; \theta_1 \dots \theta_K)}{\partial \theta_K} \right)$$


$$\frac{\partial L(A; \theta_1 \dots \theta_K)}{\partial \theta_k}$$

The general k -th partial derivative

2-Step Optimization: Update Step

$$\nabla L(A; \Theta) = \mathbf{0} \Leftrightarrow \frac{\partial L(A; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K)}{\partial \boldsymbol{\theta}_k} = 0 \quad \forall k \in \{1, \dots, K\}$$

2-Step Optimization: Update Step

$$\nabla L(A; \Theta) = \mathbf{0} \Leftrightarrow \frac{\partial L(A; \theta_1, \dots, \theta_K)}{\partial \theta_k} = 0 \quad \forall k \in \{1, \dots, K\}$$

$$\frac{\partial L(A; \theta_1 \dots \theta_K)}{\partial \theta_k} = \frac{\partial}{\partial \theta_k} \left[\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \theta_k) \right]$$

2-Step Optimization: Update Step

$$\nabla L(A; \Theta) = \mathbf{0} \Leftrightarrow \frac{\partial L(A; \boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K)}{\partial \boldsymbol{\theta}_k} = 0 \quad \forall k \in \{1, \dots, K\}$$

$$\frac{\partial L(A; \boldsymbol{\theta}_1 \dots \boldsymbol{\theta}_K)}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right]$$

↓


$$\frac{\partial L}{\partial \boldsymbol{\theta}_k}$$

To make the notation easier!

2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$


When computing the partial derivative w.r.t. $\boldsymbol{\theta}_k$ any other term $\boldsymbol{\theta}_j$ of the inner summation is treated as constant!

2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) = 0$$

2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \frac{\partial}{\partial \boldsymbol{\theta}_k} \left[\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) = 0$$

Solve for each $\boldsymbol{\theta}_k$ independently

Depends on the distance function δ

A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)

A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)
- The centroid of a cluster is the **mean** of the instances assigned to that cluster

A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)
- The centroid of a cluster is the **mean** of the instances assigned to that cluster
- (Re)Assignment of instances to clusters is based on distance/similarity to the current cluster centroids

A Special Case: K-means

- Each cluster representative is its center of mass (i.e., **centroid**)
- The centroid of a cluster is the **mean** of the instances assigned to that cluster
- (Re)Assignment of instances to clusters is based on distance/similarity to the current cluster centroids
- The basic idea is constructing clusters so that the total within-cluster **Sum of Square Distances (SSD)** is minimized

K-means: Setup

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ the set of N input data points

$\{C_1, \dots, C_K\}$ the set of K output clusters

C_k the generic k -th cluster

$$\boldsymbol{\theta}_k = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{n=1}^N \mathbf{x}_n$$

$$\text{where } |C_k| = \sum_{n=1}^N \alpha_{n,k}$$

K-means: Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \underbrace{(\|\mathbf{x}_n - \boldsymbol{\theta}_k\|_2)^2}_{\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}$$

Euclidean space

K-means: Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \underbrace{(\|\mathbf{x}_n - \boldsymbol{\theta}_k\|_2)^2}_{\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}$$

$$\begin{aligned} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) &= (\|\mathbf{x}_n - \boldsymbol{\theta}_k\|_2)^2 = \\ &= \left[\sqrt{(\mathbf{x}_n - \boldsymbol{\theta}_k)^2} \right]^2 = (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 \end{aligned}$$

Sum of Square Distances
(SSD)

K-means: Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \underbrace{(\|\mathbf{x}_n - \boldsymbol{\theta}_k\|_2)^2}_{\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}$$

$$\begin{aligned} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) &= (\|\mathbf{x}_n - \boldsymbol{\theta}_k\|_2)^2 = \\ &= \left[\sqrt{(\mathbf{x}_n - \boldsymbol{\theta}_k)^2} \right]^2 = (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 \end{aligned}$$

Sum of Square Distances
(SSD)

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2$$

K-means: Assignment Step

Minimize L w.r.t. A by fixing Θ

Intuitively, given a set of fixed centroids, L is minimized if each data point is assigned to the centroid with the smallest SSD
(L is just the SSD from each data point to its assigned centroid)

$$\alpha_{n,k} = \begin{cases} 1 & \text{if } (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 = \min_{1 \leq j \leq K} \{(\mathbf{x}_n - \boldsymbol{\theta}_j)^2\} \\ 0 & \text{otherwise} \end{cases}$$

K-means: Update Step

Minimize L w.r.t. A by fixing Θ

$$\Theta^* = \operatorname{argmin}_{\Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2}_{L(A, \Theta)}$$

Compute the gradient w.r.t. Θ , set it to 0 and solve it for Θ

K-means: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 = 0 \quad \forall k \in \{1, \dots, K\}$$

K-means: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 = 0 \quad \forall k \in \{1, \dots, K\}$$

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N -2\alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)$$

K-means: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N \alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)^2 = 0 \quad \forall k \in \{1, \dots, K\}$$

$$\frac{\partial L}{\partial \boldsymbol{\theta}_k} = \sum_{n=1}^N -2\alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k)$$

$$\text{Find } \boldsymbol{\theta}_k^* \text{ s.t. } \sum_{n=1}^N -2\alpha_{n,k} (\mathbf{x}_n - \boldsymbol{\theta}_k^*) = 0$$

K-means: Update Step

$$\begin{aligned} \sum_{n=1}^N -2\alpha_{n,k}(\mathbf{x}_n - \boldsymbol{\theta}_k^*) &= 0 \Leftrightarrow \\ 2 \sum_{n=1}^N \alpha_{n,k} \boldsymbol{\theta}_k^* &= 2 \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n \\ \boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} &= \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n \end{aligned}$$

K-means: Update Step

$\boldsymbol{\theta}_k^*$ does not depend on N,
therefore it can be factored out

$$\begin{aligned} \sum_{n=1}^N -2\alpha_{n,k}(\mathbf{x}_n - \boldsymbol{\theta}_k^*) &= 0 \Leftrightarrow \\ 2 \sum_{n=1}^N \alpha_{n,k} \boldsymbol{\theta}_k^* &= 2 \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n \\ \boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} &= \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n \end{aligned}$$

K-means: Update Step

$$\boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} = \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$

$$\boldsymbol{\theta}_k^* = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{n=1}^N \mathbf{x}_n$$

K-means: Update Step

$$\boldsymbol{\theta}_k^* \sum_{n=1}^N \alpha_{n,k} = \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$

$$\boldsymbol{\theta}_k^* = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{n=1}^N \mathbf{x}_n$$

K-means: Update Step

$$\theta_k^* \sum_{n=1}^N \alpha_{n,k} = \sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n$$

$$\theta_k^* = \frac{\sum_{n=1}^N \alpha_{n,k} \mathbf{x}_n}{\sum_{n=1}^N \alpha_{n,k}} = \mu_k = \frac{1}{|C_k|} \sum_{n=1}^N \mathbf{x}_n$$

The cluster centroid (i.e., **mean**) minimizes the objective
(for a fixed assignment A)

K-means: Lloyd-Forgy Algorithm

- I. Specify the number of output clusters K

K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids

K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids
3. **Assignment step:** Assign each observation to the closest centroid based on the distance measure chosen

K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids
3. **Assignment step:** Assign each observation to the closest centroid based on the distance measure chosen
4. **Update step:** For each of the K clusters update the centroid by computing the new mean values of all the data points now in the cluster

K-means: Lloyd-Forgy Algorithm

1. Specify the number of output clusters K
2. Select K observations **at random** from the N data points as the initial cluster centroids
3. **Assignment step:** Assign each observation to the closest centroid based on the distance measure chosen
4. **Update step:** For each of the K clusters update the centroid by computing the new mean values of all the data points now in the cluster
5. Iteratively repeat steps 3-4 until a **stopping criterion** is met

Stopping Criterion

- Several options to choose from:
 - Fixed number of iterations
 - Cluster assignments stop changing (beyond some threshold)
 - Centroid doesn't change (beyond some threshold)

Lloyd-Forgy's Convergence

- How/Why are we guaranteed the K-means algorithm ever reaches a fixed point?
 - A state in which clusters do not change

Lloyd-Forgy's Convergence

- How/Why are we guaranteed the K-means algorithm ever reaches a fixed point?
 - A state in which clusters do not change
- Intuitively, in both steps we either improve the objective or not

Lloyd-Forgy's Convergence

- How/Why are we guaranteed the K-means algorithm ever reaches a fixed point?
 - A state in which clusters do not change
- Intuitively, in both steps we either improve the objective or not
- It is an instance of more general **Expectation Maximization (EM)**
 - EM is known to converge (although not necessarily to a global optimum)

Lloyd-Forgy's Relationship with EM

- E-step = Assignment step
 - Each object is assigned to the closest centroid, i.e., to the most likely cluster
 - Monotonically decreases SSD

Lloyd-Forgy's Relationship with EM

- E-step = Assignment step

- Each object is assigned to the closest centroid, i.e., to the most likely cluster
- Monotonically decreases SSD

- M-step = Update step

- The model (i.e., centroids) are updated (i.e., SSD optimization)
- Monotonically decreases each SSD_k

Lloyd-Forgy's Complexity Analysis

- Computing the distance between two d -dimensional data points takes $O(d)$

Lloyd-Forgy's Complexity Analysis

- Computing the distance between two d -dimensional data points takes $O(d)$
- (Re-)Assigning clusters [E-step]: $O(KN)$ distance computations or $O(KNd)$

Lloyd-Forgy's Complexity Analysis

- Computing the distance between two d -dimensional data points takes $O(d)$
- (Re-)Assigning clusters [E-step]: $O(KN)$ distance computations or $O(KNd)$
- Computing centroids [M-step]: $O(Nd)$ as there are $O(N)$ average computations since each data point is added to a cluster exactly once *at each iteration*, each one taking $O(d)$

Lloyd-Forgy's Complexity Analysis

- Computing the distance between two d -dimensional data points takes $O(d)$
- (Re-)Assigning clusters [E-step]: $O(KN)$ distance computations or $O(KNd)$
- Computing centroids [M-step]: $O(Nd)$ as there are $O(N)$ average computations since each data point is added to a cluster exactly once *at each iteration*, each one taking $O(d)$
- Overall: $O(RKNd)$ assuming the 2 steps above are repeated R times

K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**
 - Forgy method **randomly** chooses K data points as the initial means
 - Random Partition method **randomly** assigns a cluster to each observation

K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**
 - Forgy method **randomly** chooses K data points as the initial means
 - Random Partition method **randomly** assigns a cluster to each observation
- Randomness may result in convergence to **sub-optimal** clusterings

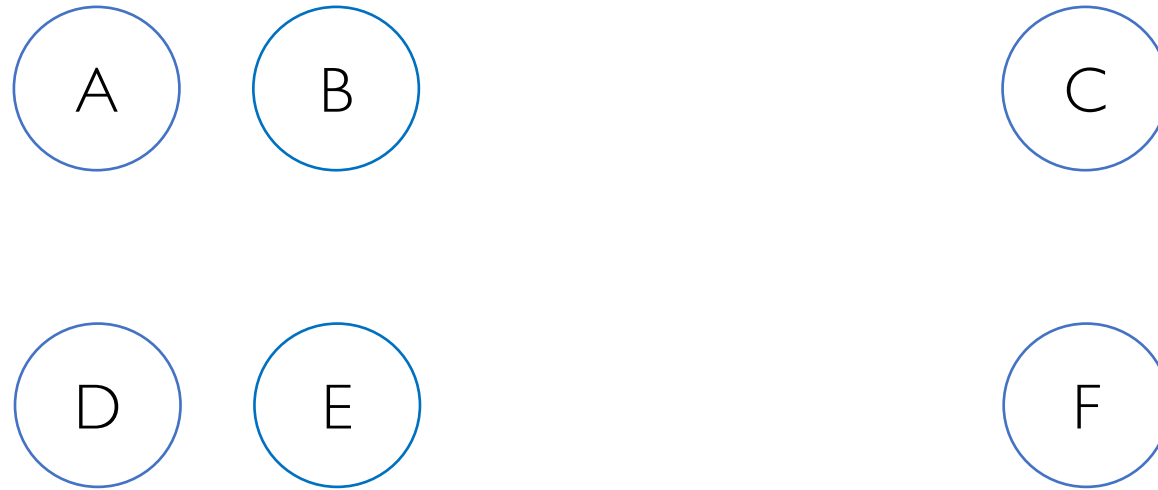
K-means: Seed Choice

- Convergence (rate) and clustering quality depends on the selection of **initial centroids**
 - Forgy method **randomly** chooses K data points as the initial means
 - Random Partition method **randomly** assigns a cluster to each observation
- Randomness may result in convergence to **sub-optimal** clusterings

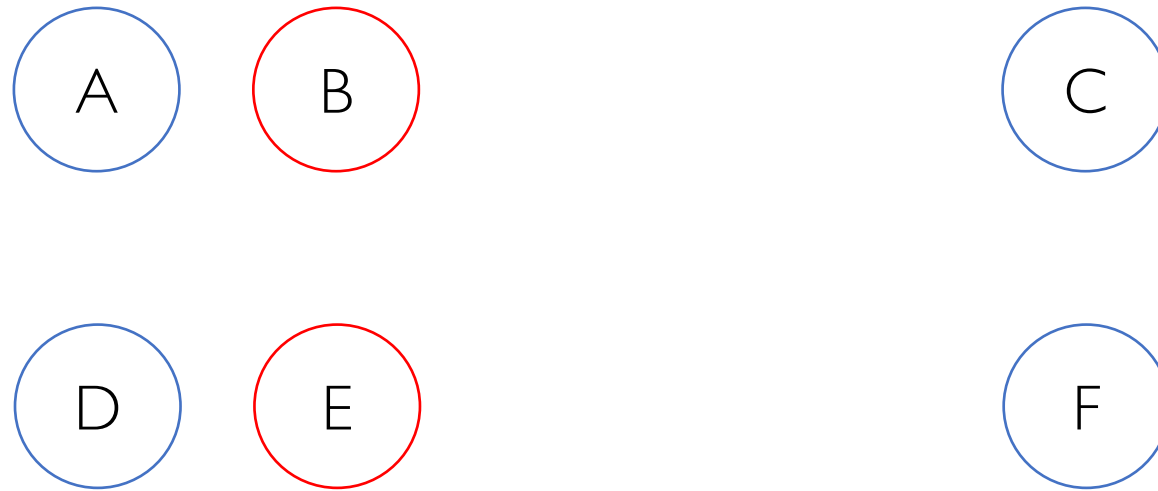
Problem Mitigation:

Execute several runs of the Lloyd-Forgy algorithm with multiple random initialization seeds

K-means: Seed Choice

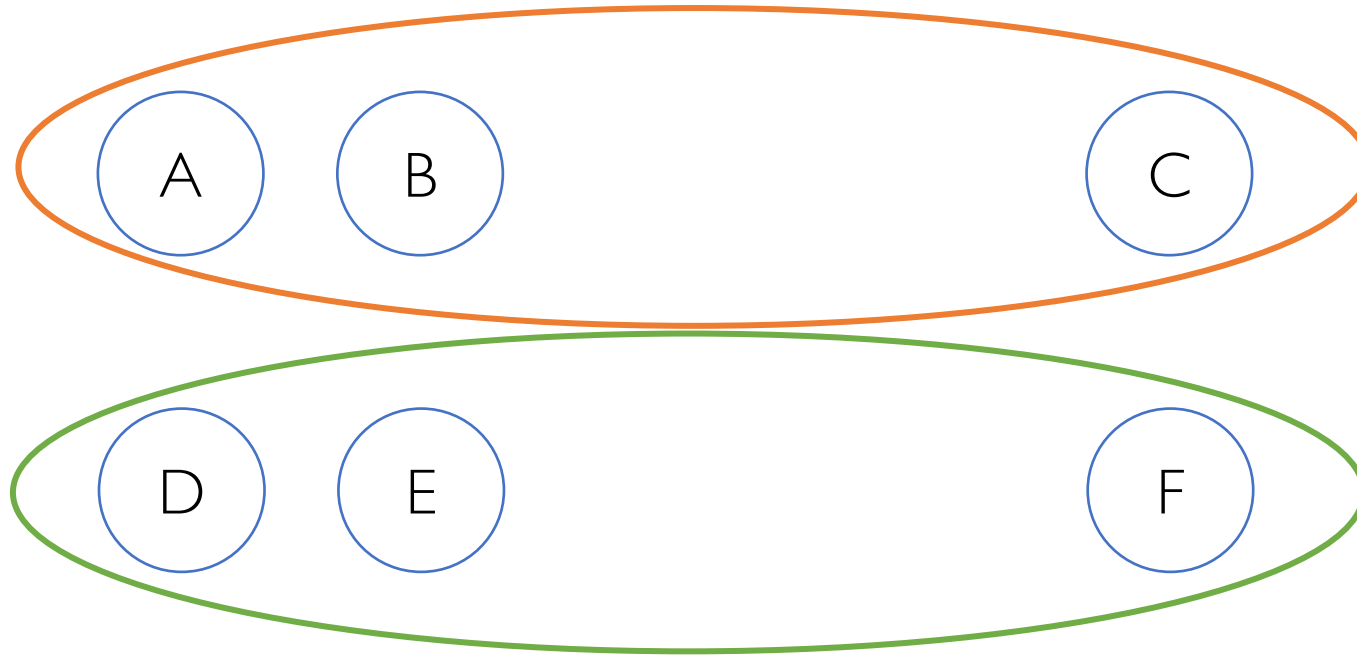


K-means: Bad (Unlucky) Seed Choice



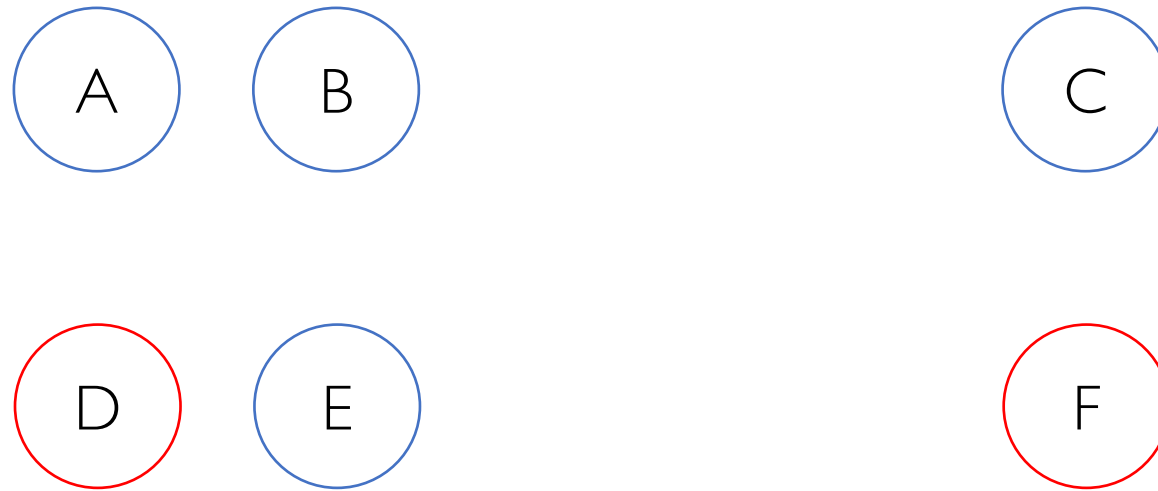
If B and E are randomly chosen as initial centroids...

K-means: Bad (Unlucky) Seed Choice



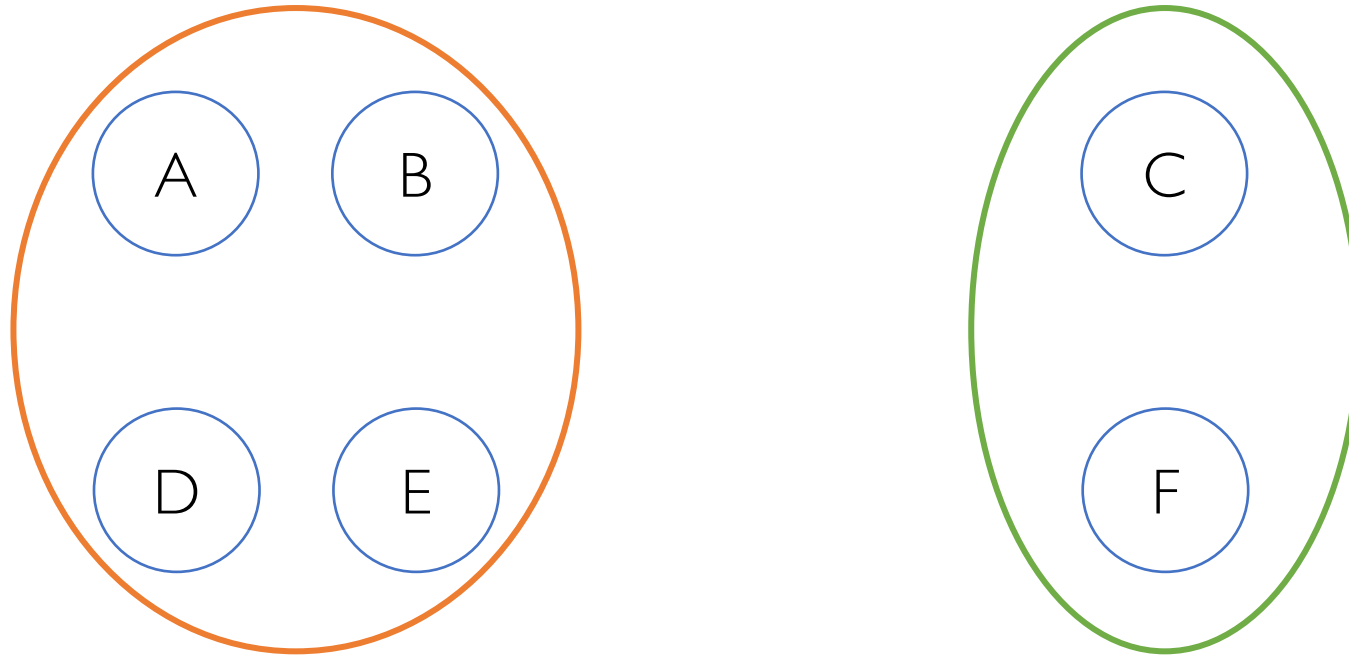
The algorithm converges to the sub-optimal clustering above

K-means: Good (Lucky) Seed Choice



If D and F are randomly chosen as initial centroids instead...

K-means: Good (Lucky) Seed Choice



The algorithm converges to a better clustering

Alternative Seed Choice: K-means++

- A preliminary method to carefully select initial centroids proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]

Alternative Seed Choice: K-means++

- A preliminary method to carefully select initial centroids proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]
- Intuition: spreading out the K initial cluster centers is a good thing

Alternative Seed Choice: K-means++

- A preliminary method to carefully select initial centroids proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]
- Intuition: spreading out the K initial cluster centers is a good thing
 1. Choose one center uniformly at random from among initial data points

Alternative Seed Choice: K-means++

- A preliminary method to carefully select initial centroids proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]
- Intuition: spreading out the K initial cluster centers is a good thing
 1. Choose one center uniformly at random from among initial data points
 2. For each data point \mathbf{x} , compute $D(\mathbf{x})$ as the distance between \mathbf{x} and the nearest center that has already been chosen

Alternative Seed Choice: K-means++

- A preliminary method to carefully select initial centroids proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]
- Intuition: spreading out the K initial cluster centers is a good thing
 1. Choose one center uniformly at random from among initial data points
 2. For each data point \mathbf{x} , compute $D(\mathbf{x})$ as the distance between \mathbf{x} and the nearest center that has already been chosen
 3. Choose one new data point at random as a new center with probability proportional to $D(\mathbf{x})^2$

Alternative Seed Choice: K-means++

- A preliminary method to carefully select initial centroids proposed in 2007 by Arthur and Vassilvitskii [[paper](#)]
- Intuition: spreading out the K initial cluster centers is a good thing
 1. Choose one center uniformly at random from among initial data points
 2. For each data point \mathbf{x} , compute $D(\mathbf{x})$ as the distance between \mathbf{x} and the nearest center that has already been chosen
 3. Choose one new data point at random as a new center with probability proportional to $D(\mathbf{x})^2$
 4. Repeat steps 2. and 3. until K centers are chosen, then run Lloyd-Forgy

"Vanilla" K-means vs. K-means++

- Random initialization used with "vanilla" K-means may produce clusters that are **arbitrarily worse** than optimum

"Vanilla" K-means vs. K-means++

- Random initialization used with "vanilla" K-means may produce clusters that are **arbitrarily worse** than optimum
- K-means++ provides an upper-bound to the approximation obtained w.r.t. the optimal solution

"Vanilla" K-means vs. K-means++

- Random initialization used with "vanilla" K-means may produce clusters that are **arbitrarily worse** than optimum
- K-means++ provides an upper-bound to the approximation obtained w.r.t. the optimal solution
- At most, clusters obtained with K-means++ initialization are $O(\log K)$ worse than the optimal partitioning

K-means: How Many Clusters?

- Number of clusters K is given
 - Great! Partition N data points into a predetermined number K of clusters
 - Unfortunately, it is very uncommon to know K in advance

K-means: How Many Clusters?

- Number of clusters K is given
 - Great! Partition N data points into a predetermined number K of clusters
 - Unfortunately, it is very uncommon to know K in advance
- Finding the “right” number K of clusters is part of the problem!
 - Trade-off between having too few and too many clusters
 - Total benefit vs. Total cost

K-means: Total Benefit

- Given a clustering, define the benefit b_i for a data point \mathbf{x}_i to be the similarity to its assigned centroid

K-means: Total Benefit

- Given a clustering, define the benefit b_i for a data point \mathbf{x}_i to be the similarity to its assigned centroid
- Define the total benefit B to be the sum of the individual benefits

K-means: Total Benefit

- Given a clustering, define the benefit b_i for a data point x_i to be the similarity to its assigned centroid
- Define the total benefit B to be the sum of the individual benefits

NOTE

There is always a clustering whose total benefit $B=N$
(where N is the number of data points)

Why?

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P=Kp$

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P = Kp$
- Define the value V of a clustering to be total benefit-total cost

$$V = B - P$$

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P = Kp$
- Define the value V of a clustering to be total benefit-total cost

$$V = B - P$$

Goal:

Find the clustering which **maximizes** V , over all choices of K

K-means: Total Cost

- Assign a cost p to each cluster, thereby a clustering with K clusters has a total cost $P=Kp$
- Define the value V of a clustering to be total benefit-total cost

$$V = B - P$$

Goal:

Find the clustering which **maximizes** V , over all choices of K

B increases with larger values of K , but P allows to stop that

K-means: "Elbow" Method

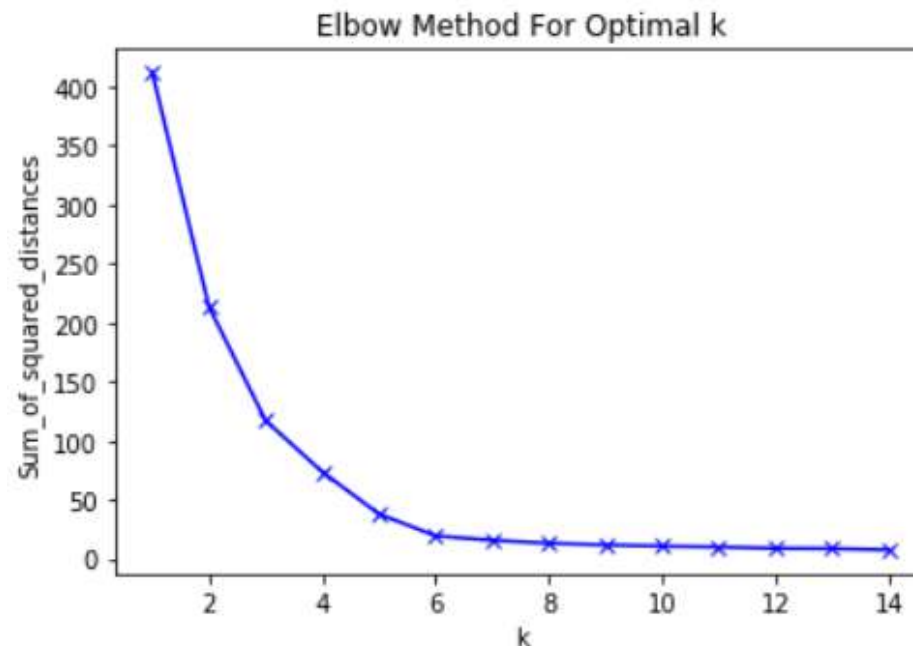
- Empirical method to figure out the right number K of clusters

K-means: "Elbow" Method

- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost

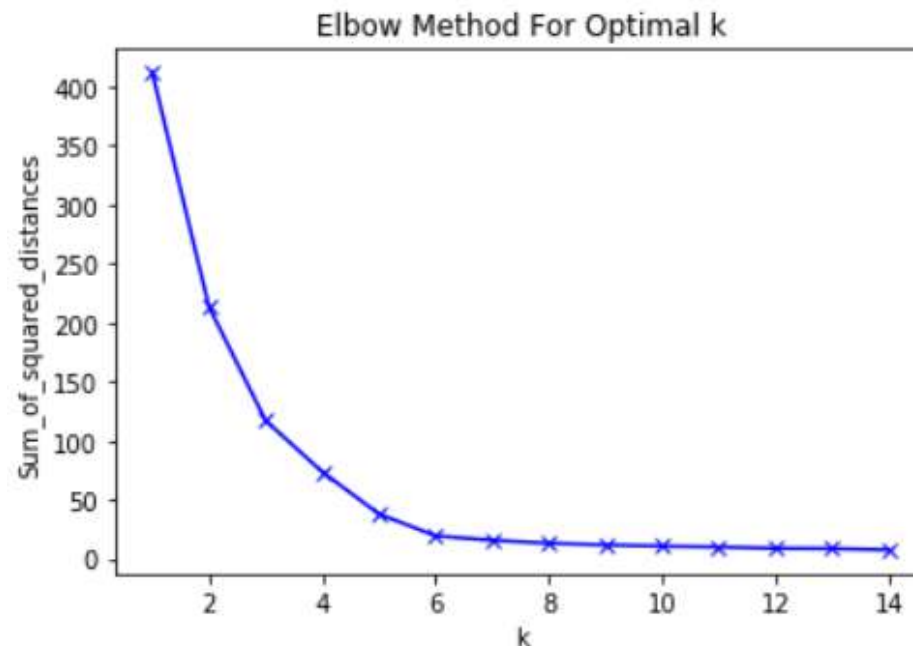
K-means: "Elbow" Method

- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost
- Try multiple values of K and look at the change of the SSD



K-means: "Elbow" Method


- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost
- Try multiple values of K and look at the change of the SSD



As K increases, SSD sharply decreases

K-means: "Elbow" Method

- Empirical method to figure out the right number K of clusters
- Trade-off between total benefit and total cost
- Try multiple values of K and look at the change of the SSD



As K increases, SSD sharply decreases
up to a certain value

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)
- The same hard clustering framework can be used with other δ

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)
- The same hard clustering framework can be used with other δ
- Some of them just resemble Euclidean distance, and centroids (i.e., means) still minimize those
 - $\delta = \text{Cosine distance}$ = Euclidean distance on normalized input points
 - $\delta = \text{Correlation}$ = Euclidean distance on standardized input points

Non-Euclidean Distances

- So far, we have focused on **Euclidean distance** (i.e., $\delta = L^2$ -Norm)
- The same hard clustering framework can be used with other δ
- Some of them just resemble Euclidean distance, and centroids (i.e., means) still minimize those
 - $\delta = \text{Cosine distance}$ = Euclidean distance on normalized input points
 - $\delta = \text{Correlation}$ = Euclidean distance on standardized input points
- Others, require specific minimizers
 - $\delta = \text{Manhattan distance}$ (L^1 -Norm) \rightarrow median is the minimizer (**K-medians**)

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster
- Works with **any arbitrary distance** δ

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster
- Works with **any arbitrary distance** δ
- **PAM** (**P**artitioning **A**round **M**edoids) greedy Algorithm, introduced by Kaufman and Rousseeuw in 1987 [[paper](#)] vs. Lloyd-Forgy

Alternative Formulations: K-medoids

- Similar to K-means yet chooses input data points as centers (**medoids**)
- A medoid is the closest object to any other point in the cluster
- Works with **any arbitrary distance** δ
- **PAM** (**P**artitioning **A**round **M**edoids) greedy Algorithm, introduced by Kaufman and Rousseeuw in 1987 [[paper](#)] vs. Lloyd-Forgy
- Robust to outliers yet computationally expensive $O(K(N-K)^2)$

Bradley-Fayyad-Reina (BFR) K-means

- A variant of K-means explicitly thought for large datasets

Bradley-Fayyad-Reina (BFR) K-means

- A variant of K-means explicitly thought for large datasets
- Works better in high-dimensional Euclidean space

Bradley-Fayyad-Reina (BFR) K-means

- A variant of K-means explicitly thought for large datasets
- Works better in high-dimensional Euclidean space
- (Strong) Assumption on the shape of clusters:
 - Normally distributed around the centroid
 - Independence between data dimensions

Bradley-Fayyad-Reina (BFR) K-means

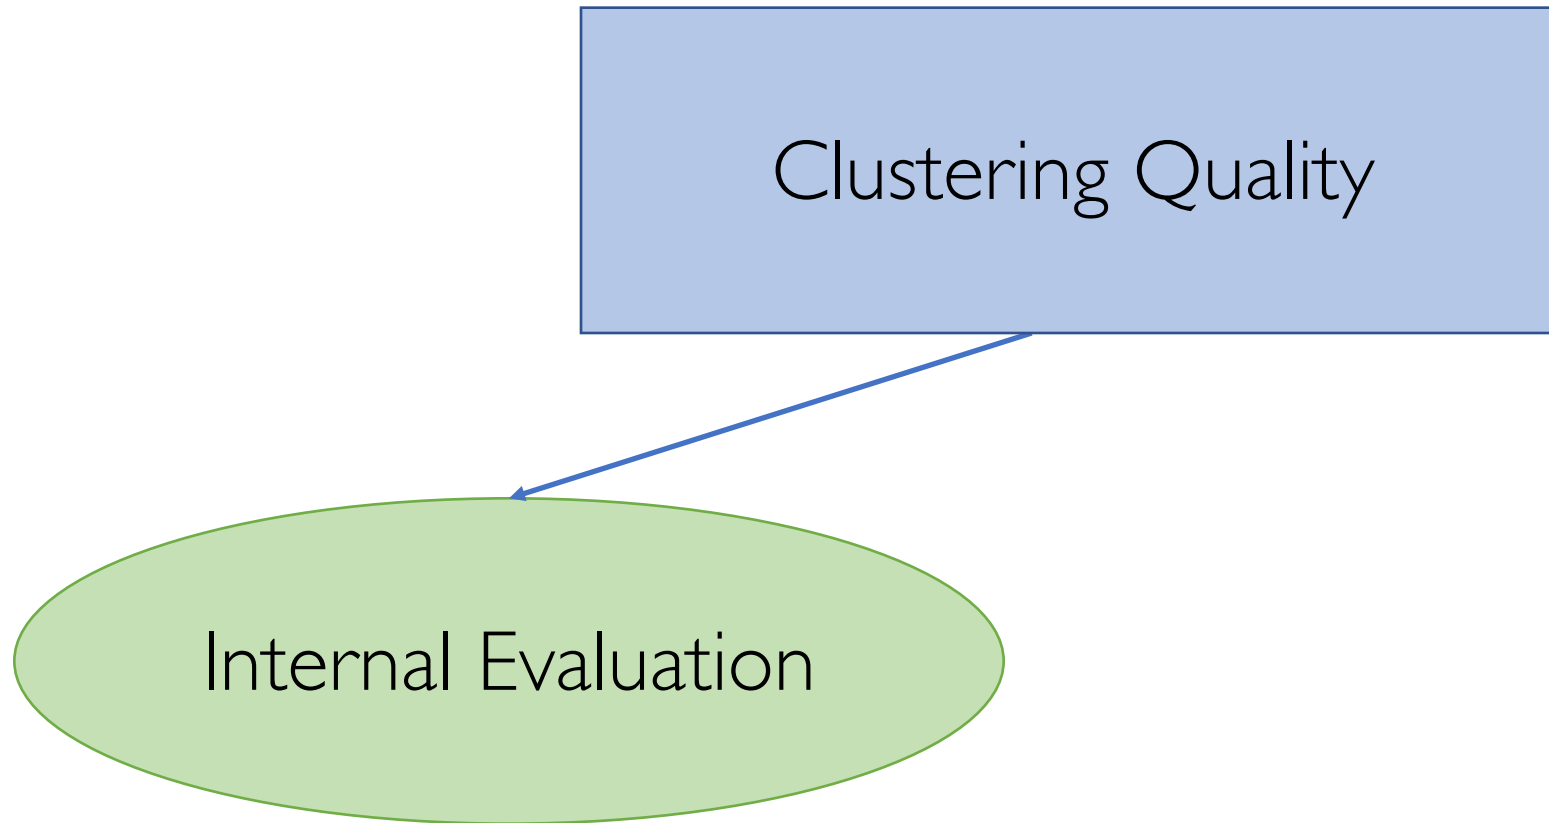
- A variant of K-means explicitly thought for large datasets
- Works better in high-dimensional Euclidean space
- (Strong) Assumption on the shape of clusters:
 - Normally distributed around the centroid
 - Independence between data dimensions
- Reference to the original [paper](#)

Measures of Clustering Quality

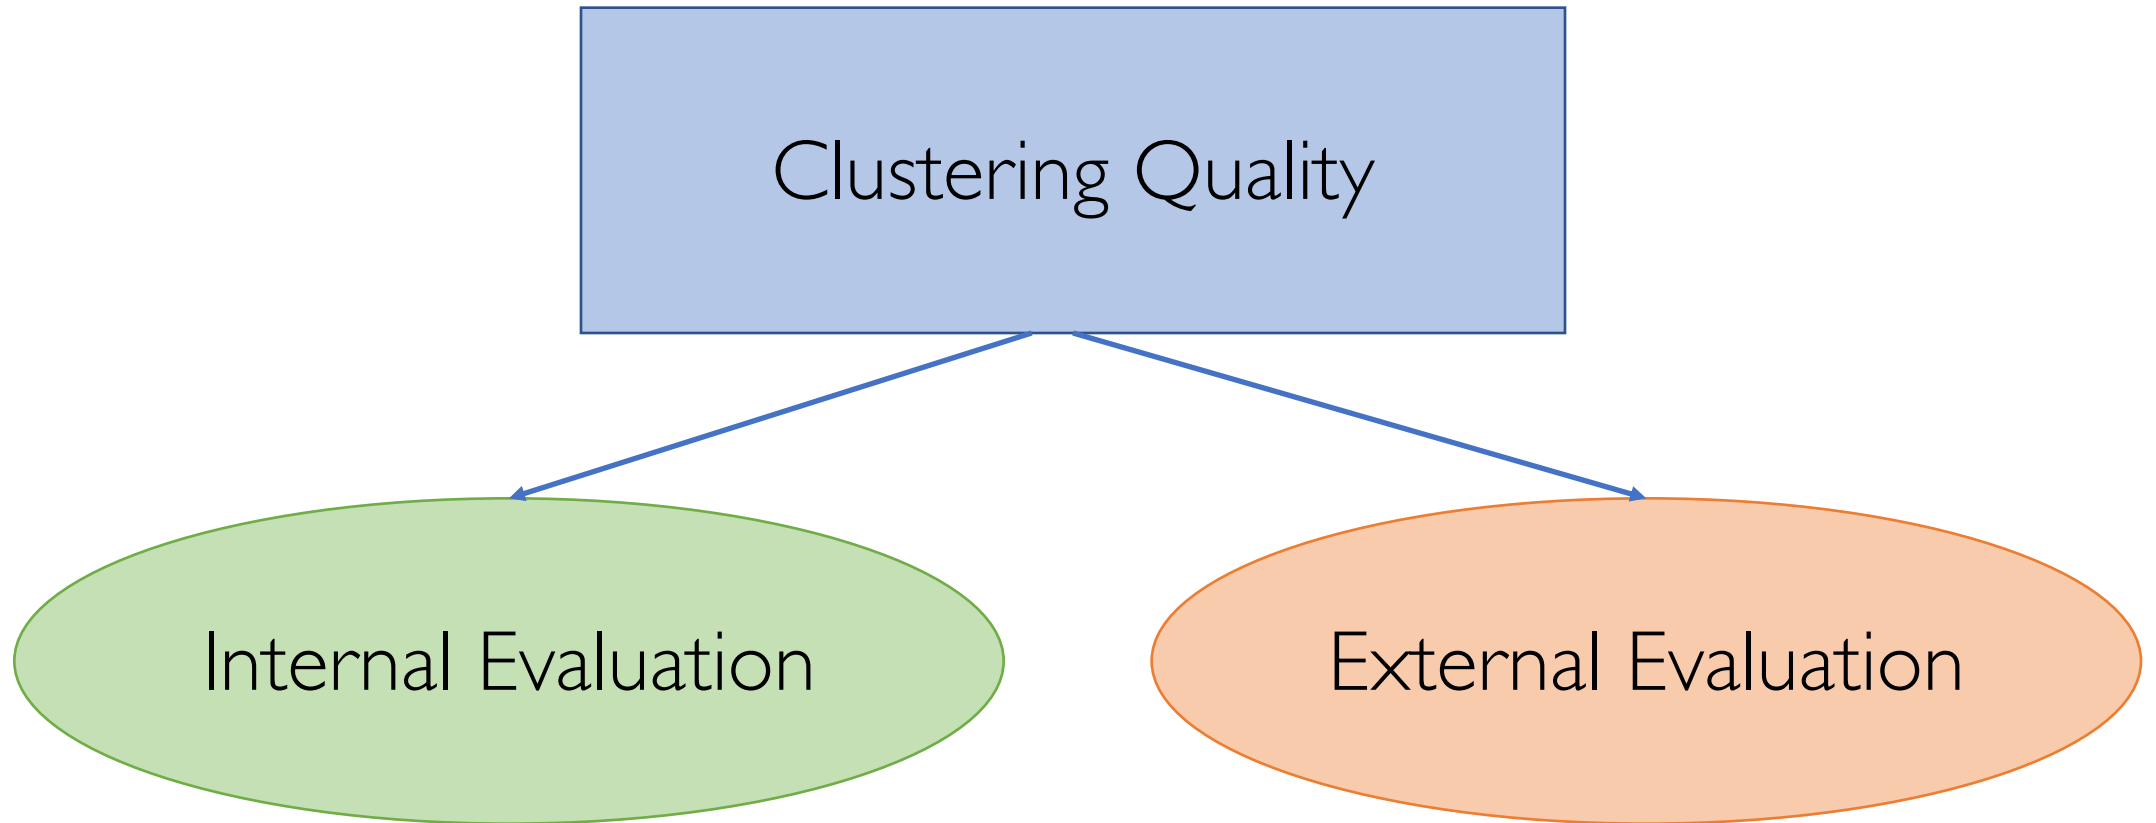


Clustering Quality

Measures of Clustering Quality



Measures of Clustering Quality



Internal Evaluation

- Clustering is evaluated based on the data that was clustered itself

Internal Evaluation

- Clustering is evaluated based on the data that was clustered itself
- A good clustering will produce high quality clusters with:
 - high intra-cluster similarity
 - low inter-cluster similarity

Internal Evaluation

- Clustering is evaluated based on the data that was clustered itself
- A good clustering will produce high quality clusters with:
 - high intra-cluster similarity
 - low inter-cluster similarity
- The measured quality of a clustering depends on
 - data representation
 - similarity measure

Internal Evaluation: Davies-Bouldin Index

$$DB = \frac{1}{K} \sum_{i=1}^K \max_{j \neq i} \left(\frac{\sigma_i + \sigma_j}{\delta(\boldsymbol{\mu}_i, \boldsymbol{\mu}_j)} \right)$$

K = number of clusters

$\boldsymbol{\mu}_k$ = centroid of cluster C_k

σ_k = avg. distance of all elements of cluster C_k from its centroid $\boldsymbol{\mu}_k$

$\delta(\boldsymbol{\mu}_i, \boldsymbol{\mu}_j)$ = distance between centroids of C_i and C_j

The smaller the better

Internal Evaluation: Dunn Index

$$D = \frac{\min_{1 \leq i < j \leq K} \delta(C_i, C_j)}{\max_{1 \leq k \leq K} \delta'(C_k)}$$

K = number of clusters

$\delta(C_i, C_j)$ = distance between cluster C_i and C_j

$\delta'(C_k)$ = intra-cluster distance of cluster C_k

Distance between centroids

Max distance between any pair of objects

The higher the better

Internal Evaluation: Silhouette Coefficient

mean distance between i and all other data points in the same cluster C_i

$$a(i) = \frac{1}{|C_i| - 1} \sum_{j \in C_i, j \neq i} \delta(i, j)$$

smallest mean distance of i to all points in any other cluster $C_k \neq C_i$

$$b(i) = \min_{k \neq i} \frac{1}{|C_k|} \sum_{j \in C_k} \delta(i, j)$$

$$s(i) = \begin{cases} 1 - a(i)/b(i) & \text{if } a(i) < b(i) \\ 0 & \text{if } a(i) = b(i) \\ b(i)/a(i) - 1 & \text{if } a(i) > b(i) \end{cases}$$

The higher the better

External Evaluation

- Clustering is evaluated based on data that was not used for clustering, yet pre-classified (**gold standard** data)

External Evaluation

- Clustering is evaluated based on data that was not used for clustering, yet pre-classified (**gold standard** data)
- Quality measured by the ability to discover some or all of the hidden patterns in gold standard data

External Evaluation

- Clustering is evaluated based on data that was not used for clustering, yet pre-classified (**gold standard** data)
- Quality measured by the ability to discover some or all of the hidden patterns in gold standard data
- Hard as it requires labeled data typically provided by human experts

External Evaluation: Purity

$C_1 \dots, C_K$ = set of K clusters

$L_1 \dots, L_J$ = set of J labels

$n_{i,j}$ = number of items with label L_j clustered in C_i

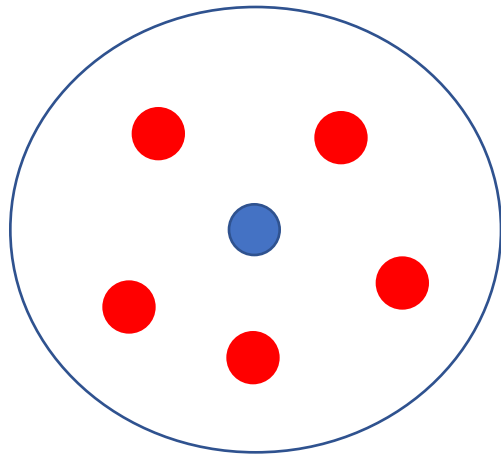
$n_i = \sum_{j=1}^J n_{i,j}$ number of items clustered in C_i

$$\text{purity}(C_i) = \frac{1}{n_i} \max_{j \in \{1, \dots, J\}} n_{i,j}$$

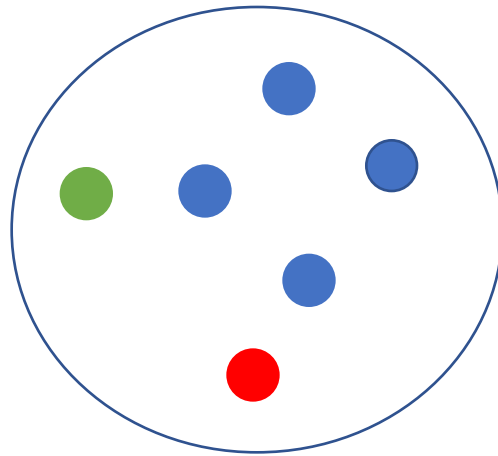
$$\text{purity} = \frac{1}{K} \sum_{i=1}^K \text{purity}(C_i)$$

Biased because having as many clusters as items maximizes purity

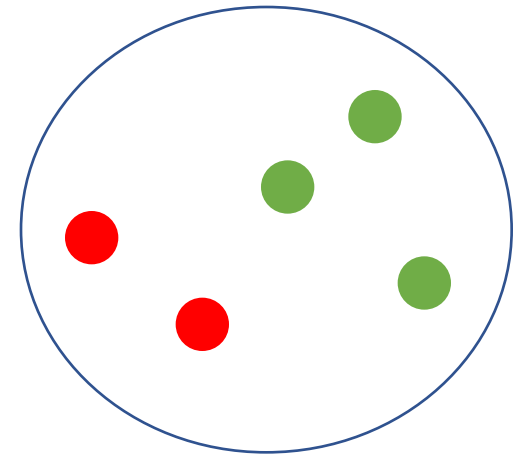
External Evaluation: Purity Example



C_1



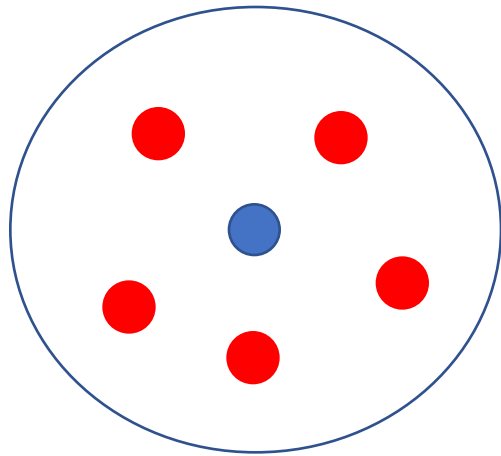
C_2



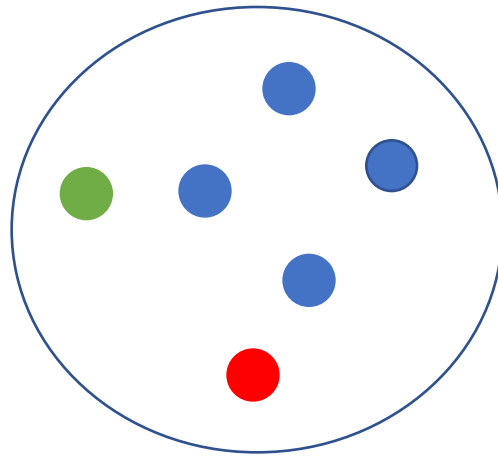
C_3

● L_1 ● L_2 ● L_3

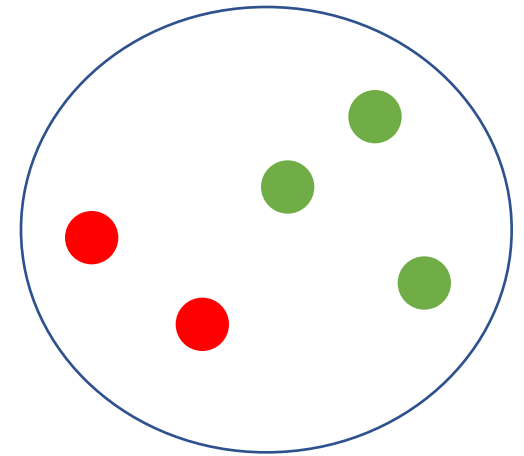
External Evaluation: Purity Example



C_1



C_2

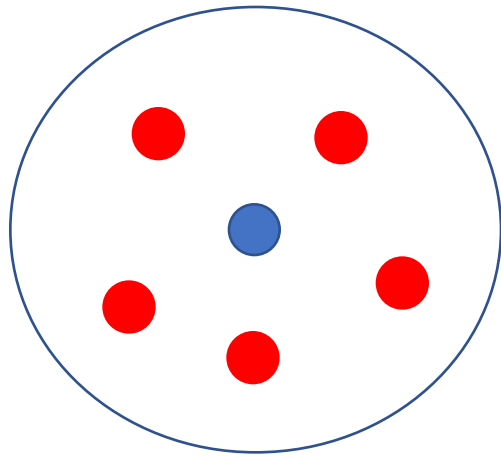


C_3

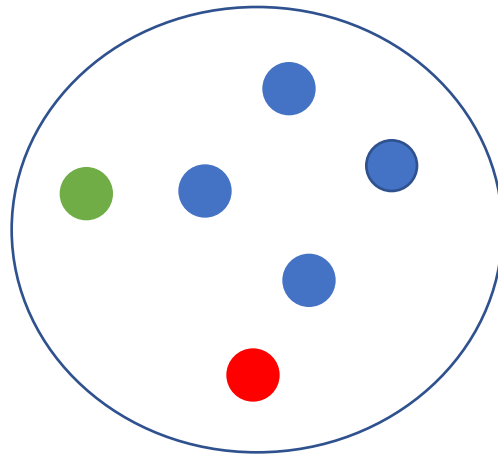
● L_1 ● L_2 ● L_3

$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

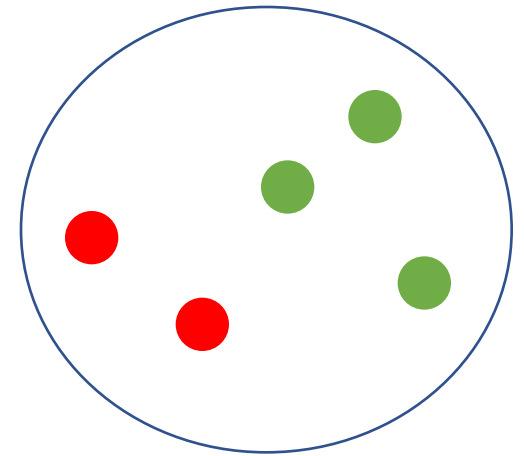
External Evaluation: Purity Example



C_1



C_2




C_3

● L_1 ● L_2 ● L_3


$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

$$\text{purity}(C_2) = 1/6 * \max\{1, 4, 1\} = 4/6 = 2/3$$


External Evaluation: Purity Example



C_1



C_2



C_3

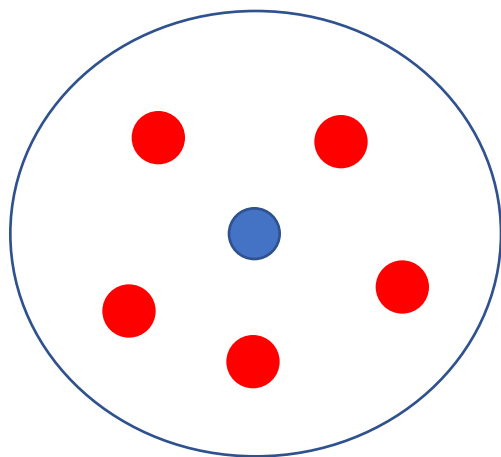
● L_1 ● L_2 ● L_3

$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

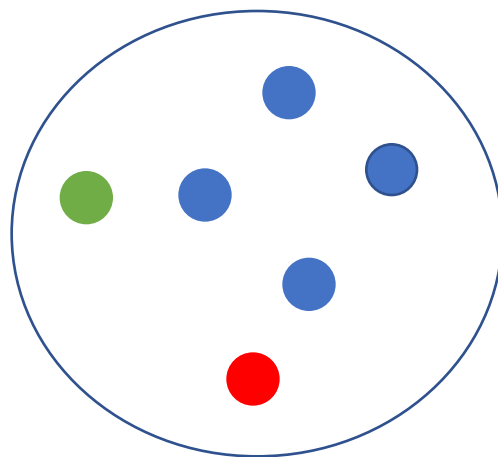
$$\text{purity}(C_2) = 1/6 * \max\{1, 4, 1\} = 4/6 = 2/3$$

$$\text{purity}(C_3) = 1/5 * \max\{2, 0, 3\} = 3/5$$

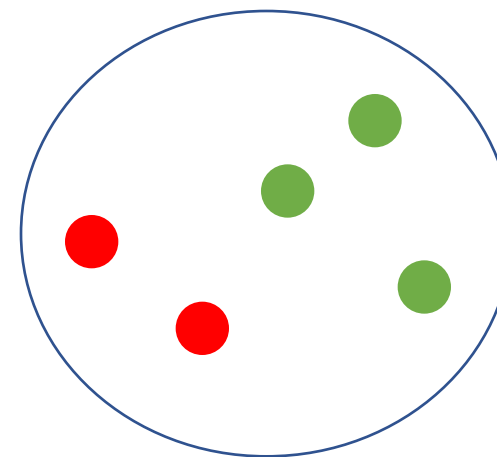
External Evaluation: Purity Example



C_1



C_2



C_3

● L_1 ● L_2 ● L_3

$$\text{purity}(C_1) = 1/6 * \max\{5, 1, 0\} = 5/6$$

$$\text{purity}(C_2) = 1/6 * \max\{1, 4, 1\} = 4/6 = 2/3$$

$$\text{purity}(C_3) = 1/5 * \max\{2, 0, 3\} = 3/5$$

$$\text{purity} = 1/3 * \text{purity}(C_1) + \text{purity}(C_2) + \text{purity}(C_3) = 7/10$$

External Evaluation: Rand Index

$$\text{Rand} = \frac{TP + TN}{TP + TN + FP + FN}$$

TP = number of *true positives*

TN = number of *true negatives*

FP = number of *false positives*

FN = number of *false negatives*

All computed from **pairs** of elements

Measures the level of agreement between
clustering and ground truth

External Evaluation: Rand Index

n. of pairs	Same Cluster in Clustering	Different Clusters in Clustering
Same Cluster in Ground-Truth	TRUE POSITIVES (TP)	FALSE NEGATIVES (FN)
Different Clusters in Ground-Truth	FALSE POSITIVES (FP)	TRUE NEGATIVES (TN)

Confusion Matrix

External Evaluation: Precision, Recall, F-measure

$$P = \frac{TP}{TP + FP} \quad R = \frac{TP}{TP + FN}$$

$$F_{\beta} = \frac{(\beta^2 + 1) \cdot P \cdot R}{\beta^2 \cdot P + R}$$

$$F_1 = \frac{2 \cdot P \cdot R}{P + R}$$

Balances the contribution of false negatives by weighting recall through a parameter β

External Evaluation: Many Other Measures

- Jaccard index
- Dice index
- Fowlkes-Mallows index
- Mutual information
- etc.

Take-Home Message of Today

- Formulate clustering as a (**non-convex**) optimization problem
 - Focus on flat partitioning (hard)

Take-Home Message of Today

- Formulate clustering as a (**non-convex**) optimization problem
 - Focus on flat partitioning (hard)
- Computing exact solution is **NP-hard** due to exponential search space

Take-Home Message of Today

- Formulate clustering as a (**non-convex**) optimization problem
 - Focus on flat partitioning (hard)
- Computing exact solution is **NP-hard** due to exponential search space
- Iterative (**approximated**) methods converge to local minimum
 - Lloyd-Forgy Algorithm for **K-means** (to minimize Euclidean-related distances)
 - PAM Algorithm for **K-medoids** (to minimize any distance function)

Take-Home Message of Today

- Formulate clustering as a (**non-convex**) optimization problem
 - Focus on flat partitioning (hard)
- Computing exact solution is **NP-hard** due to exponential search space
- Iterative (**approximated**) methods converge to local minimum
 - Lloyd-Forgy Algorithm for **K-means** (to minimize Euclidean-related distances)
 - PAM Algorithm for **K-medoids** (to minimize any distance function)
- Internal vs. External measures of **clustering quality**