

Big Data Computing

Master's Degree in Computer Science

2022-2023

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

tolomei@di.uniroma1.it



SAPIENZA
UNIVERSITÀ DI ROMA

Recap from Last Lecture

- Data often come with redundant and noisy (high-dimensional) representations

Recap from Last Lecture

- Data often come with redundant and noisy (high-dimensional) representations
- **Goal:** Extract the maximum possible information from the data while reducing the noise and ignoring redundancies

Recap from Last Lecture

- Data often come with redundant and noisy (high-dimensional) representations
- **Goal:** Extract the maximum possible information from the data while reducing the noise and ignoring redundancies
- PCA achieves this goal by transforming correlated features in the data into **linearly independent** (i.e., orthogonal) components


Recap from Last Lecture

- Data often come with redundant and noisy (high-dimensional) representations
- **Goal:** Extract the maximum possible information from the data while reducing the noise and ignoring redundancies
- PCA achieves this goal by transforming correlated features in the data into **linearly independent** (i.e., orthogonal) components
- As a result, data dimensionality can be reduced to these components


Eigenvectors of the Covariance Matrix

$$K\mathbf{e} = \lambda\mathbf{e}$$

Eigenvectors of the Covariance Matrix




Eigenvectors of the Covariance Matrix



When you multiply a matrix by an **eigenvector** \mathbf{e} the resulting vector does not change its direction, but it is only scaled by a factor λ (**eigenvalue**)


Eigenvectors of the Covariance Matrix



When you multiply a matrix by an **eigenvector** \mathbf{e} the resulting vector does not change its direction, but it is only scaled by a factor λ (**eigenvalue**)

In other words, eigenvectors encapsulate all the relevant information to describe a linear transformation (in our case, represented by the covariance matrix K)

Eigenvectors of the Covariance Matrix



When you multiply a matrix by an **eigenvector** e the resulting vector does not change its direction, but it is only scaled by a factor λ (**eigenvalue**)

Principal Components

eigenvectors of the covariance matrix with the **largest** eigenvalues

How Do We Compute Eigenvectors?

Remember that we want to solve for \mathbf{e} the following:

$$K\mathbf{e} = \lambda\mathbf{e}$$

How Do We Compute Eigenvectors?

Remember that we want to solve for \mathbf{e} the following:

$$K\mathbf{e} = \lambda\mathbf{e}$$

We can rewrite the system of equations above as:

$$K\mathbf{e} - \lambda\mathbf{e} = 0 \Rightarrow (K - \lambda I)\mathbf{e} = 0$$

I is the identity matrix

How Do We Compute Eigenvectors?

We therefore resort to solving the following **homogeneous system**:

$$(K - \lambda I)\mathbf{e} = 0$$

How Do We Compute Eigenvectors?

We therefore resort to solving the following **homogeneous system**:

$$(K - \lambda I)\mathbf{e} = 0$$

Any homogeneous system always has a **trivial solution**, i.e., the zero vector
 $\mathbf{e} = 0$

How Do We Compute Eigenvectors?

We therefore resort to solving the following **homogeneous system**:

$$(K - \lambda I)\mathbf{e} = 0$$

Any homogeneous system always has a **trivial solution**, i.e., the zero vector
 $\mathbf{e} = 0$

The only way for the homogeneous system above to have a **non-trivial** solution is for its matrix $(K - \lambda I)$ to be **non-invertible**, otherwise:

$$\cancel{(K - \lambda I)} \cancel{(K - \lambda I)^{-1}} \boxed{\mathbf{e} = 0} \cancel{(K - \lambda I)^{-1}}$$

How Do We Compute Eigenvectors?

A square matrix is **invertible** iff its determinant is **not** equal to 0

How Do We Compute Eigenvectors?

A square matrix is **invertible** iff its determinant is **not** equal to 0



If the determinant of the matrix $(K - \lambda I)$ is equal to 0, it is **non-invertible**

How Do We Compute Eigenvectors?

A square matrix is **invertible** iff its determinant is **not** equal to 0



If the determinant of the matrix $(K - \lambda I)$ is equal to 0, it is **non-invertible**



The corresponding homogeneous system will have a **non-trivial** solution

How Do We Compute Eigenvectors?

I. Find the eigenvalues by solving for λ : $\det(K - \lambda I) = 0$

$$\det \left(\underbrace{\begin{bmatrix} 2 - \lambda & 4/5 \\ 4/5 & 3/5 - \lambda \end{bmatrix}}_{K - \lambda I} \right) = 0$$

How Do We Compute Eigenvectors?

I. Find the eigenvalues by solving for λ : $\det(K - \lambda I) = 0$

$$\det \left(\underbrace{\begin{bmatrix} 2 - \lambda & 4/5 \\ 4/5 & 3/5 - \lambda \end{bmatrix}}_{K - \lambda I} \right) = 0$$

$$(2 - \lambda)(3/5 - \lambda) - (4/5)(4/5) = \lambda^2 - 13/5\lambda + 14/25$$

$$\boxed{\lambda^2 - 13/5\lambda + 14/25 = 0} \quad \text{characteristic equation of } K$$

How Do We Compute Eigenvectors?

I. Find the eigenvalues by solving for λ : $\det(K - \lambda I) = 0$

$$\det \left(\underbrace{\begin{bmatrix} 2 - \lambda & 4/5 \\ 4/5 & 3/5 - \lambda \end{bmatrix}}_{K - \lambda I} \right) = 0$$

$$(2 - \lambda)(3/5 - \lambda) - (4/5)(4/5) = \lambda^2 - 13/5\lambda + 14/25$$

$$\lambda^2 - 13/5\lambda + 14/25 = 0$$

characteristic equation of K

$$\lambda_1 = \frac{13 + \sqrt{113}}{10} \approx 2.36; \quad \lambda_2 = \frac{13 - \sqrt{113}}{10} \approx 0.24$$

How Do We Compute Eigenvectors?

2. Plug each eigenvalue in to find the corresponding eigenvector

$$\underbrace{\begin{bmatrix} 2 & 4/5 \\ 4/5 & 3/5 \end{bmatrix}}_K \underbrace{\begin{bmatrix} e_{1,1} \\ e_{1,2} \end{bmatrix}}_{\mathbf{e}_1} = \lambda_1 \underbrace{\begin{bmatrix} e_{1,1} \\ e_{1,2} \end{bmatrix}}_{\mathbf{e}_1}$$
$$\underbrace{\begin{bmatrix} 2 & 4/5 \\ 4/5 & 3/5 \end{bmatrix}}_K \underbrace{\begin{bmatrix} e_{2,1} \\ e_{2,2} \end{bmatrix}}_{\mathbf{e}_2} = \lambda_2 \underbrace{\begin{bmatrix} e_{2,1} \\ e_{2,2} \end{bmatrix}}_{\mathbf{e}_2}$$

How Do We Compute Eigenvectors?

Let's see what happens for λ_1

$$\begin{cases} 2e_{1,1} + 4/5e_{1,2} = \frac{13+\sqrt{113}}{10}e_{1,1} \\ 4/5e_{1,1} + 3/5e_{1,2} = \frac{13+\sqrt{113}}{10}e_{1,2} \end{cases}$$

How Do We Compute Eigenvectors?

Let's see what happens for λ_1

$$\begin{cases} 2e_{1,1} + 4/5e_{1,2} = \frac{13+\sqrt{113}}{10}e_{1,1} \\ 4/5e_{1,1} + 3/5e_{1,2} = \frac{13+\sqrt{113}}{10}e_{1,2} \end{cases}$$



Just replacing $\lambda_1 \sim 2.36$

$$\begin{cases} 2e_{1,1} + 0.8e_{1,2} = 2.36e_{1,1} \\ 0.8e_{1,1} + 0.6e_{1,2} = 2.36e_{1,2} \end{cases}$$

How Do We Compute Eigenvectors?


Let's see what happens for λ_1

$$\begin{cases} 2e_{1,1} + 4/5e_{1,2} = \frac{13+\sqrt{113}}{10}e_{1,1} \\ 4/5e_{1,1} + 3/5e_{1,2} = \frac{13+\sqrt{113}}{10}e_{1,2} \end{cases}$$



Just replacing $\lambda_1 \sim 2.36$

$$\begin{cases} 2e_{1,1} + 0.8e_{1,2} = 2.36e_{1,1} \\ 0.8e_{1,1} + 0.6e_{1,2} = 2.36e_{1,2} \end{cases}$$




$$e_{1,1} \approx 2.2e_{1,2}$$

The system has infinitely many solutions

How Do We Compute Eigenvectors?

Any vector which satisfies the relationship above works!




$$\begin{bmatrix} 1.1 \\ 0.5 \end{bmatrix}$$

$$e_{1,1} \approx 2.2e_{1,2}$$


How Do We Compute Eigenvectors?

Any vector which satisfies the relationship above works!




How Do We Compute Eigenvectors?

Any vector which satisfies the relationship above works!



How Do We Compute Eigenvectors?


By convention, we restrict to $\|\mathbf{e}_i\| = 1$



$$\begin{bmatrix} 0.91 \\ 0.41 \end{bmatrix}$$
$$e_{1,1} \approx 2.2e_{1,2}$$

How Do We Compute Eigenvectors?


The second eigenvector \mathbf{e}_2 can be found by plugging in the smaller eigenvalue λ_2



How Do We Compute Eigenvectors?

The second eigenvector \mathbf{e}_2 can be found by plugging in the smaller eigenvalue λ_2

This is just orthogonal to the previously found \mathbf{e}_1




How Do We Compute Eigenvectors?

The second eigenvector \mathbf{e}_2 can be found by plugging in the smaller eigenvalue λ_2

This is just orthogonal to the previously found \mathbf{e}_1


For any $d \times d$ real symmetric matrix (like K), the eigenvalues are **real** and eigenvectors can be chosen **real** and **orthonormal**



How Do We Compute Eigenvectors?

The second eigenvector \mathbf{e}_2 can be found by plugging in the smaller eigenvalue λ_2

This is just orthogonal to the previously found \mathbf{e}_1



\mathbf{e}_1 and \mathbf{e}_2 are the new coordinate system replacing the original x_1 and x_2

$$\mathbf{e}_1 = \begin{bmatrix} 0.91 \\ 0.41 \end{bmatrix} \quad \mathbf{e}_2 = \begin{bmatrix} -0.41 \\ 0.91 \end{bmatrix}$$

Principal Components

$$\mathbf{e}_1 = \begin{bmatrix} 0.91 \\ 0.41 \end{bmatrix} \quad \mathbf{e}_2 = \begin{bmatrix} -0.41 \\ 0.91 \end{bmatrix}$$

\mathbf{e}_1 is the 1st principal component as it is the eigenvector corresponding to the **largest** eigenvalue

\mathbf{e}_2 is the 2nd principal component as it is the eigenvector corresponding to the **smallest** eigenvalue

Projecting to New Dimensions: 2-d Case

- \mathbf{e}_1 and \mathbf{e}_2 identify our new coordinate system (principal components)

Projecting to New Dimensions: 2-d Case


- \mathbf{e}_1 and \mathbf{e}_2 identify our new coordinate system (principal components)
- Both of them are 2-dimensional vectors

Projecting to New Dimensions: 2-d Case

- \mathbf{e}_1 and \mathbf{e}_2 identify our new coordinate system (principal components)
- Both of them are 2-dimensional vectors
- Let $\mathbf{x} = (x_1, x_2)$ be a point (i.e., a vector) in the original $(\mathbf{x}_1, \mathbf{x}_2)$ -space

Projecting to New Dimensions: 2-d Case

- \mathbf{e}_1 and \mathbf{e}_2 identify our new coordinate system (principal components)
- Both of them are 2-dimensional vectors
- Let $\mathbf{x} = (x_1, x_2)$ be a point (i.e., a vector) in the original $(\mathbf{x}_1, \mathbf{x}_2)$ -space




Goal

We want to represent \mathbf{x} in the new $(\mathbf{e}_1, \mathbf{e}_2)$ -coordinate system

Projecting to New Dimensions: 2-d Case

I. Center \mathbf{x} around the mean of each dimension


$$\mathbf{x}' = \mathbf{x} - \boldsymbol{\mu} = (x_1 - \mu_1, x_2 - \mu_2)$$



Projecting to New Dimensions: 2-d Case

I. Center \mathbf{x} around the mean of each dimension


$$\mathbf{x}' = \mathbf{x} - \boldsymbol{\mu} = (x_1 - \mu_1, x_2 - \mu_2)$$



Projecting to New Dimensions: 2-d Case


2. Project \mathbf{x}' on each dimension \mathbf{e}_1 and \mathbf{e}_2

$$\mathbf{x}' = \underbrace{(x'_1, x'_2)}_{\text{coordinates of } \mathbf{x}' \text{ in the } (\mathbf{e}_1, \mathbf{e}_2)\text{-space}} = (\mathbf{x}'^T \mathbf{e}_1, \mathbf{x}'^T \mathbf{e}_2)$$




Projecting to New Dimensions: 2-d Case

Why the dot product?



Projecting to New Dimensions: 2-d Case


Why the dot product?



Projecting to New Dimensions: 2-d Case

Why the dot product?

$$\cos\theta = (\mathbf{x}'\mathbf{e}_1)/\|\mathbf{x}'\|\|\mathbf{e}_1\|$$




Projecting to New Dimensions: 2-d Case

Why the dot product?

$$\cos\theta = (\mathbf{x}'\mathbf{e}_1)/\|\mathbf{x}'\|\|\mathbf{e}_1\|$$

$$\|\mathbf{x}'\|\cos\theta = \mathbf{x}'\mathbf{e}_1/\|\mathbf{e}_1\|$$




Projecting to New Dimensions: 2-d Case

Why the dot product?

$$\cos\theta = (\mathbf{x}'\mathbf{e}_1)/\|\mathbf{x}'\|\|\mathbf{e}_1\|$$


$$\|\mathbf{x}'\|\cos\theta = \mathbf{x}'\mathbf{e}_1/\|\mathbf{e}_1\|$$

$$\|\mathbf{e}_1\| = 1$$



Projecting to New Dimensions: 2-d Case

Why the dot product?




$$\cos\theta = (\mathbf{x}'\mathbf{e}_1)/\|\mathbf{x}'\|\|\mathbf{e}_1\|$$

$$\|\mathbf{x}'\|\cos\theta = \mathbf{x}'\mathbf{e}_1/\|\mathbf{e}_1\|$$

$$\|\mathbf{x}'\|\cos\theta = \mathbf{x}'\mathbf{e}_1$$

$$\|\mathbf{e}_1\| = 1$$



Projecting to New Dimensions: 2-d Case

The new coordinates of the original data point \mathbf{x} according to the eigenvectors \mathbf{e}_1 and \mathbf{e}_2 are as follows:

$$\mathbf{x}' = \begin{bmatrix} x'_1 \\ x'_2 \end{bmatrix} = \begin{bmatrix} \mathbf{x}'^T \mathbf{e}_1 \\ \mathbf{x}'^T \mathbf{e}_2 \end{bmatrix} = \begin{bmatrix} (x_1 - \mu_1)e_{1,1} + (x_2 - \mu_2)e_{1,2} \\ (x_1 - \mu_1)e_{2,1} + (x_2 - \mu_2)e_{2,2} \end{bmatrix}$$

Projecting to New Dimensions: d -dimensions

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$

Original d -dimensional data point

Projecting to New Dimensions: d -dimensions

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$

Original d -dimensional
data point

$$\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_k$$
$$k \ll d, \mathbf{e}_i \in \mathbb{R}^d$$

$k \ll d$
principal components

Projecting to New Dimensions: d -dimensions

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$

Original d -dimensional
data point

$$\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_k \\ k \ll d, \mathbf{e}_i \in \mathbb{R}^d$$

$k \ll d$
principal components

I. Mean centering

$$\mathbf{x}' = \mathbf{x} - \boldsymbol{\mu} = \begin{bmatrix} x_1 - \mu_1 \\ x_2 - \mu_2 \\ \vdots \\ x_d - \mu_d \end{bmatrix}$$

Projecting to New Dimensions: d -dimensions

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix}$$

Original d -dimensional data point

$\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_k$
 $k \ll d, \mathbf{e}_i \in \mathbb{R}^d$

$k \ll d$
 principal components

2. Projection to principal components

$$\mathbf{x}' = \begin{bmatrix} x'_1 \\ x'_2 \\ \vdots \\ x'_k \end{bmatrix} = \begin{bmatrix} (\mathbf{x} - \boldsymbol{\mu})^T \mathbf{e}_1 \\ (\mathbf{x} - \boldsymbol{\mu})^T \mathbf{e}_2 \\ \vdots \\ (\mathbf{x} - \boldsymbol{\mu})^T \mathbf{e}_k \end{bmatrix} = \begin{bmatrix} (x_1 - \mu_1)e_{1,1} + (x_2 - \mu_2)e_{1,2} + \dots + (x_d - \mu_d)e_{1,d} \\ (x_1 - \mu_1)e_{2,1} + (x_2 - \mu_2)e_{2,2} + \dots + (x_d - \mu_d)e_{2,d} \\ \vdots \\ (x_1 - \mu_1)e_{k,1} + (x_2 - \mu_2)e_{k,2} + \dots + (x_d - \mu_d)e_{k,d} \end{bmatrix}$$

Why Eigenvector with the Largest Eigenvalue

- We have only "visually proved" that eigenvector \mathbf{e} turns towards the direction of the greatest variance

Why Eigenvector with the Largest Eigenvalue

- We have only "visually proved" that eigenvector \mathbf{e} turns towards the direction of the greatest variance
- To actually prove it, we need to show that \mathbf{e} maximizes the variance among all possible projections

Why Eigenvector with the Largest Eigenvalue

- We have only "visually proved" that eigenvector \mathbf{e} turns towards the direction of the greatest variance
- To actually prove it, we need to show that \mathbf{e} maximizes the variance among all possible projections
- Moreover, we pick the eigenvector with the **largest eigenvalue** λ because λ is exactly the variance of the data along that eigenvector

Why Eigenvector with the Largest Eigenvalue

- We have only "visually proved" that eigenvector \mathbf{e} turns towards the direction of the greatest variance
- To actually prove it, we need to show that \mathbf{e} maximizes the variance among all possible projections
- Moreover, we pick the eigenvector with the **largest eigenvalue** λ because λ is exactly the variance of the data along that eigenvector

More details available here:

https://github.com/gtolomei/big-data-computing/raw/master/extra/Notes_on_Principal_Component_Analysis.pdf

How Many Dimensions?

- In a d -dimensional space we may have $\mathbf{e}_1, \dots, \mathbf{e}_d$ length-1 eigenvectors

How Many Dimensions?


- In a d -dimensional space we may have $\mathbf{e}_1, \dots, \mathbf{e}_d$ length-1 eigenvectors
- We want to select k dimensions ($k \ll d$)

How Many Dimensions?

- In a d -dimensional space we may have $\mathbf{e}_1, \dots, \mathbf{e}_d$ length-1 eigenvectors
- We want to select k dimensions ($k \ll d$)
- We know that eigenvalue λ_i is the variance along \mathbf{e}_i and the overall variance of the data is given by the sum of all the eigenvalues

How Many Dimensions?

- In a d -dimensional space we may have $\mathbf{e}_1, \dots, \mathbf{e}_d$ length-1 eigenvectors
- We want to select k dimensions ($k \ll d$)
- We know that eigenvalue λ_i is the variance along \mathbf{e}_i and the overall variance of the data is given by the sum of all the eigenvalues



Pick the subset of k eigenvectors that "explain" the most variance

How Many Dimensions?

I. Sort eigenvectors by eigenvalues such that $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$

How Many Dimensions?


1. Sort eigenvectors by eigenvalues such that $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$



2. Pick the first k eigenvectors that explain $x\%$ of the total variance

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i} \leq 1$$

e.g., $x = 90 \div 95\%$



Practical Issues of PCA

- Covariance and variance are highly sensitive to large values

Practical Issues of PCA

- Covariance and variance are highly sensitive to large values
- If one dimension takes on extremely large values w.r.t. other dimensions the former will be considered as the principal component

Practical Issues of PCA

- Covariance and variance are highly sensitive to large values
- If one dimension takes on extremely large values w.r.t. other dimensions the former will be considered as the principal component

Solution

Normalize each dimension to 0-mean and 1-std-deviation


$$z = \frac{x - \mu}{\sigma}$$

Practical Issues of PCA

- PCA assumes the projection subspace is linear, i.e., an hyperplane:
 - 1-d \rightarrow straight line, 2-d \rightarrow flat surface, ...


Practical Issues of PCA

- PCA assumes the projection subspace is linear, i.e., an hyperplane:
 - 1-d \rightarrow straight line, 2-d \rightarrow flat surface, ...
- If data live in a low-dimensional but not linear space (i.e., manifold), PCA can still be applied but may not work nicely




Practical Issues of PCA

- PCA assumes the projection subspace is linear, i.e., an hyperplane:
 - 1-d \rightarrow straight line, 2-d \rightarrow flat surface, ...
- If data live in a low-dimensional but not linear space (i.e., manifold), PCA can still be applied but may not work nicely



Practical Issues of PCA

- PCA assumes the projection subspace is linear, i.e., an hyperplane:
 - 1-d \rightarrow straight line, 2-d \rightarrow flat surface, ...
- If data live in a low-dimensional but not linear space (i.e., manifold), PCA can still be applied but may not work nicely



PCA will find a straight line and will not mimic non-linearity

Take-Home Message of Today

- PCA is a dimensionality reduction technique which allows representing high-dimensional data into low-dimensional linear subspace
 - The original space is rotated to make dimensions uncorrelated (i.e., independent)

Take-Home Message of Today

- PCA is a dimensionality reduction technique which allows representing high-dimensional data into low-dimensional linear subspace
 - The original space is rotated to make dimensions uncorrelated (i.e., independent)
- Reduced size of data means faster processing and smaller storage

Take-Home Message of Today

- PCA is a dimensionality reduction technique which allows representing high-dimensional data into low-dimensional linear subspace
 - The original space is rotated to make dimensions uncorrelated (i.e., independent)
- Reduced size of data means faster processing and smaller storage
- PCA can be very expensive for many very large-scale applications

Take-Home Message of Today

- PCA is a dimensionality reduction technique which allows representing high-dimensional data into low-dimensional linear subspace
 - The original space is rotated to make dimensions uncorrelated (i.e., independent)
- Reduced size of data means faster processing and smaller storage
- PCA can be very expensive for many very large-scale applications
- If data do not live on a linear subspace PCA may not work well