

# Big Data Computing

Master's Degree in Computer Science

2022-2023

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)



SAPIENZA  
UNIVERSITÀ DI ROMA

# Recap from Last Lectures

- We presented 2 linear models: linear regression and logistic regression
- Those hypotheses work well whenever there exists a linear relationship between the features (input) and the response (output)
- Model's parameter estimation done either analytically (OLS) or iteratively (Gradient Descent)

# Tree-based Methods

- Suitable for both **regression** and **classification** tasks

# Tree-based Methods

- Suitable for both **regression** and **classification** tasks
- Work by repeatedly splitting the input feature space into a number of regions

# Tree-based Methods

- Suitable for both **regression** and **classification** tasks
- Work by repeatedly splitting the input feature space into a number of regions
- A prediction is computed by taking the **mean** (regression) or the **mode** (classification) of the points in the region which the observation belongs

# Tree-based Methods

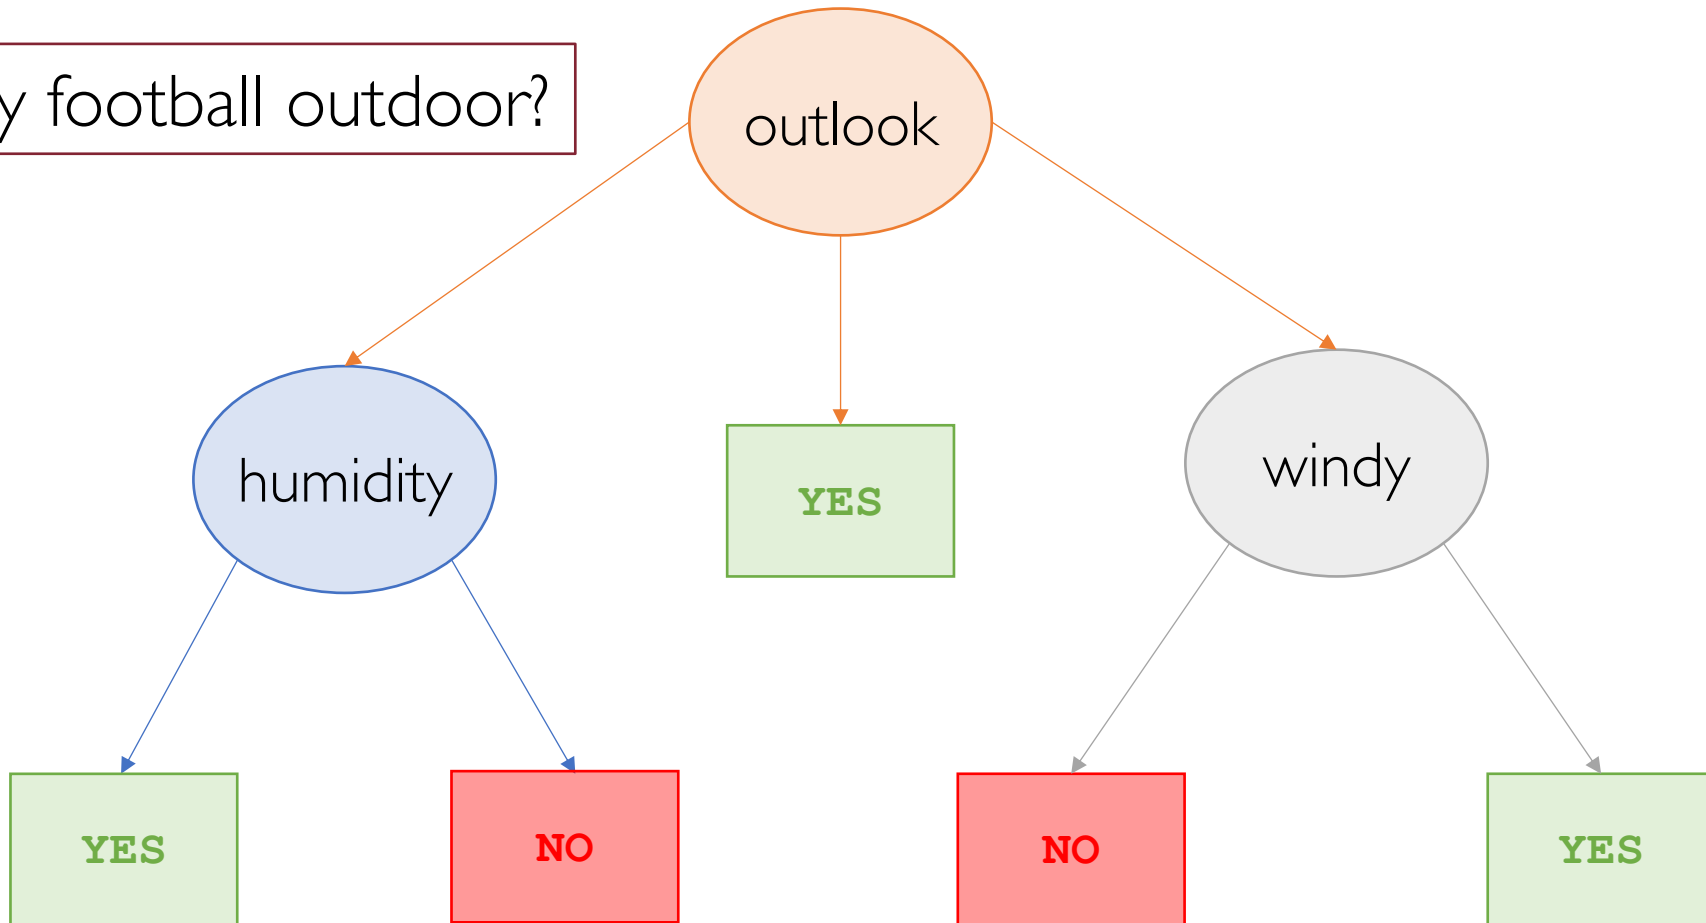
- Suitable for both **regression** and **classification** tasks
- Work by repeatedly splitting the input feature space into a number of regions
- A prediction is computed by taking the **mean** (regression) or the **mode** (classification) of the points in the region which the observation belongs
- The set of splitting rules can be represented as a tree (**decision tree**)

# Tree-based Methods

- Suitable for both **regression** and **classification** tasks
- Work by repeatedly splitting the input feature space into a number of regions
- A prediction is computed by taking the **mean** (regression) or the **mode** (classification) of the points in the region which the observation belongs
- The set of splitting rules can be represented as a tree (**decision tree**)
- Highly human-interpretable models

# Decision Tree: An Example

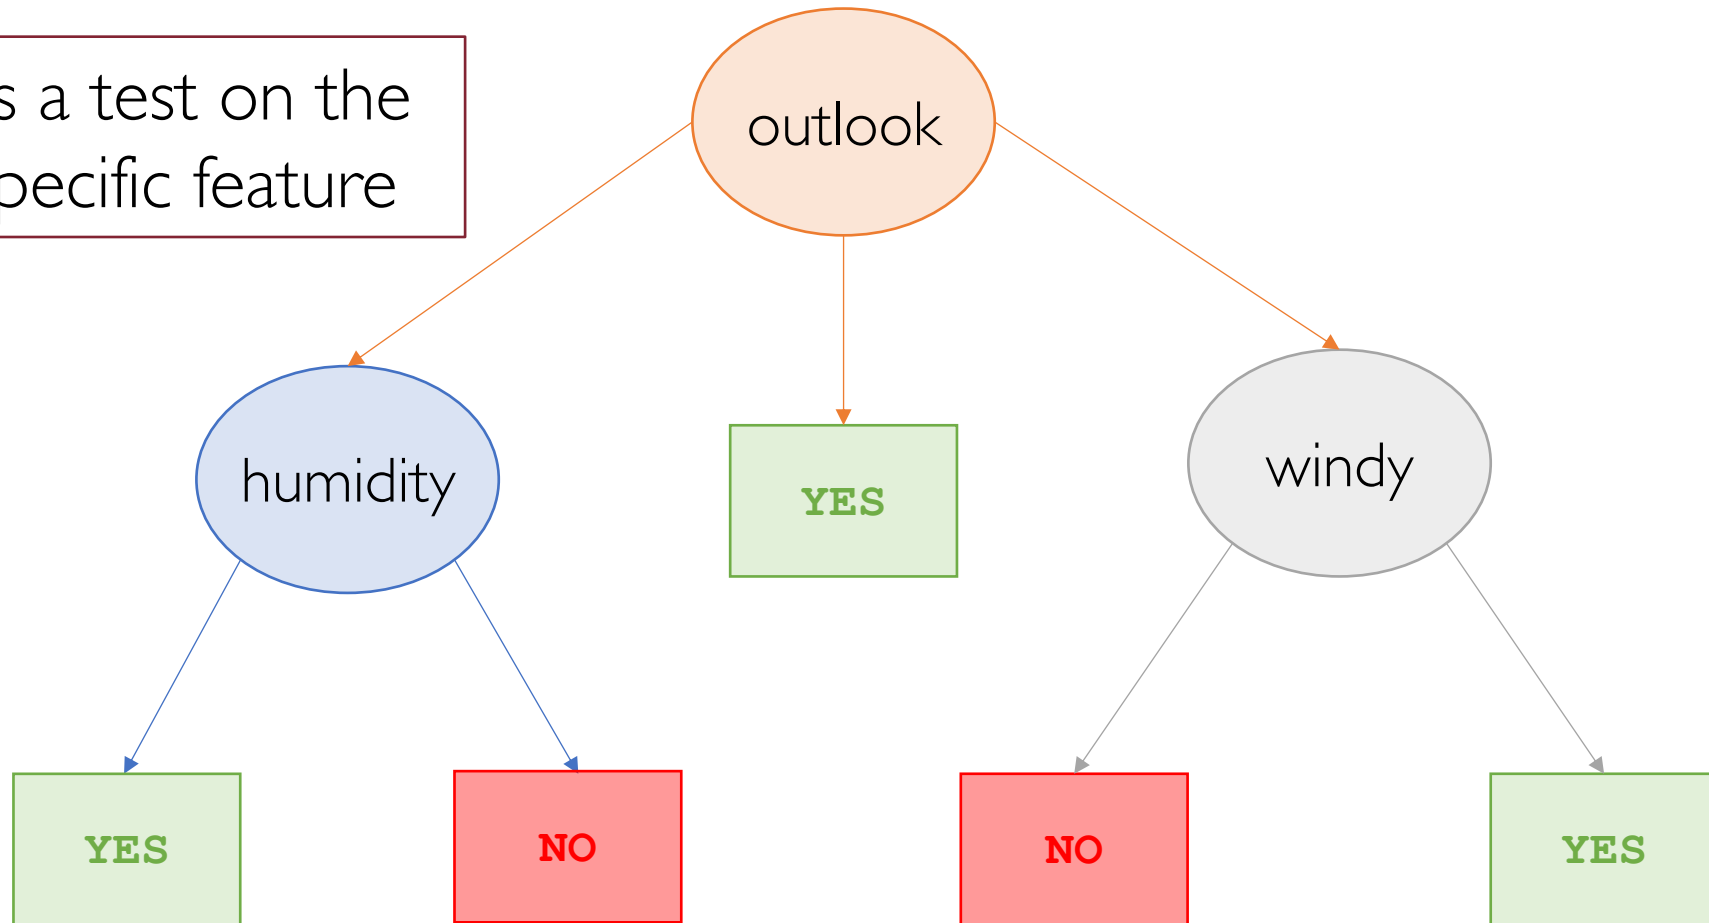
Should we play football outdoor?



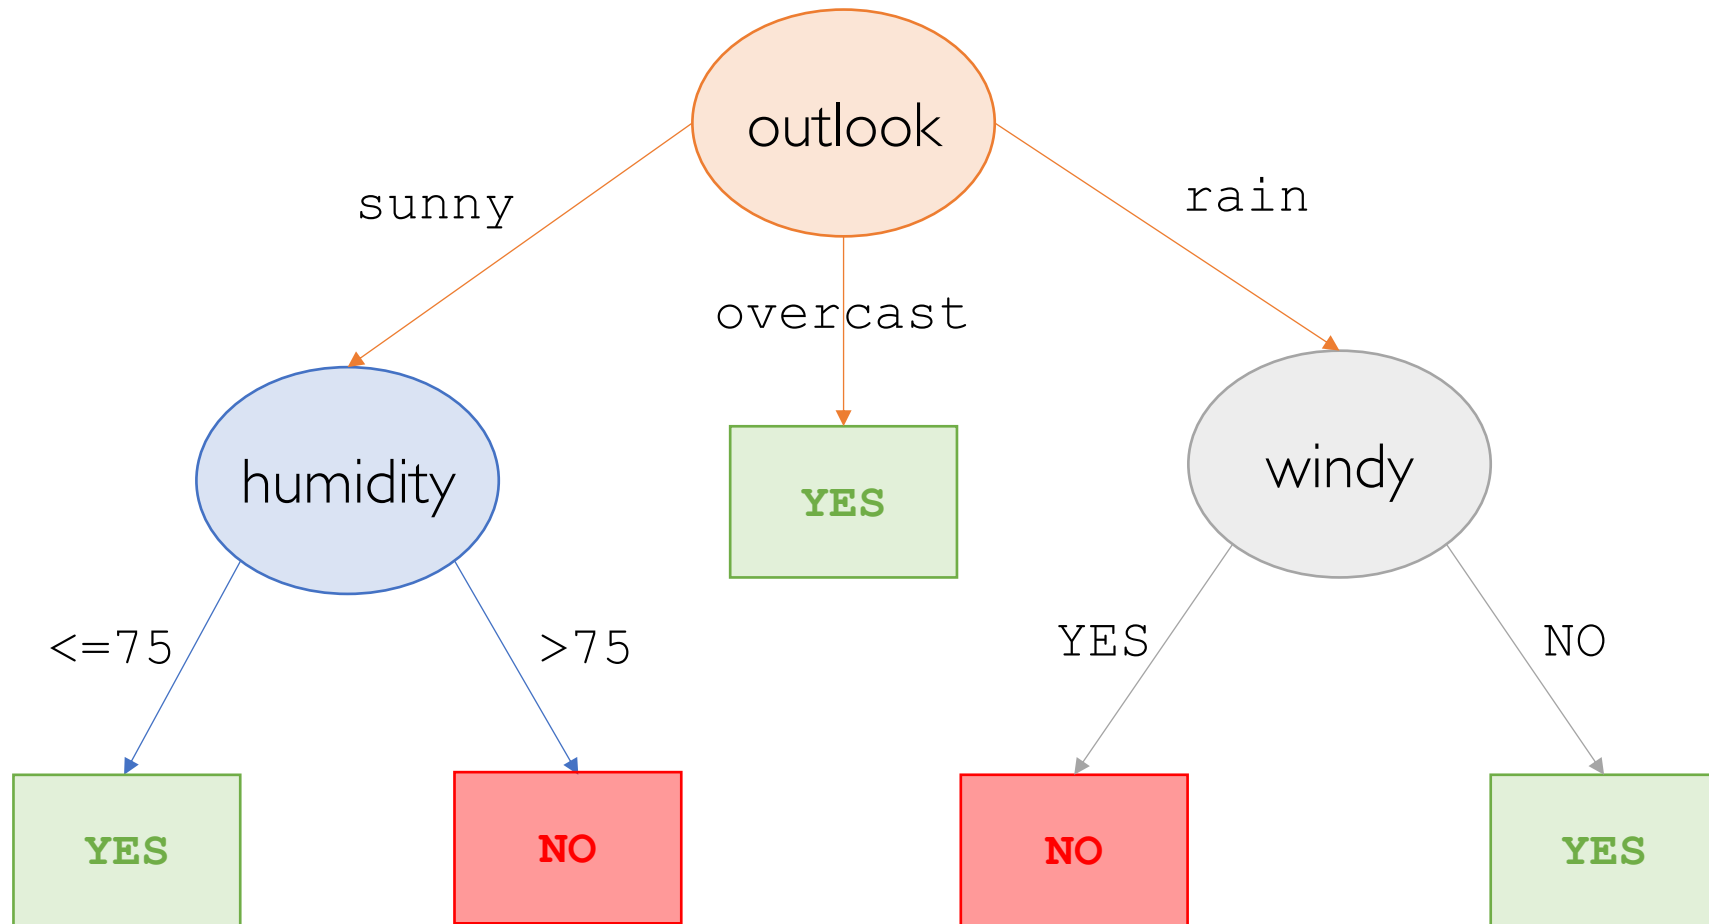


# Decision Tree: An Example

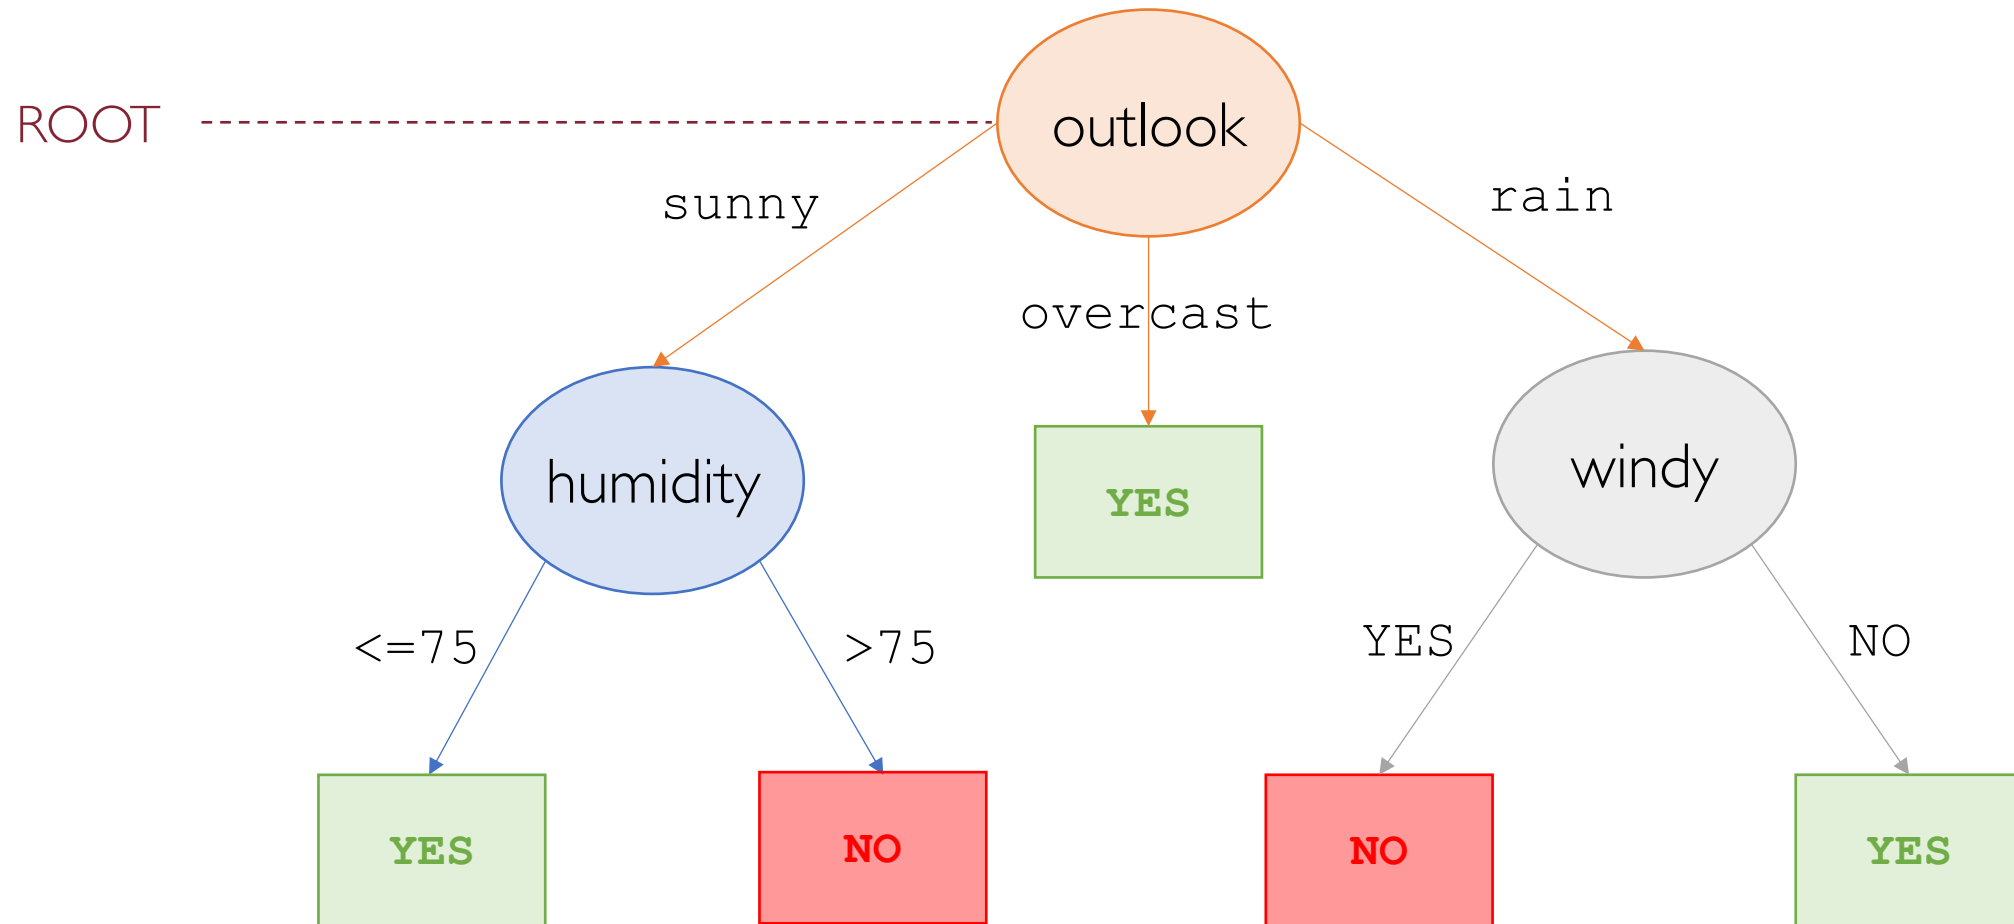
Each node is a test on the value of a specific feature



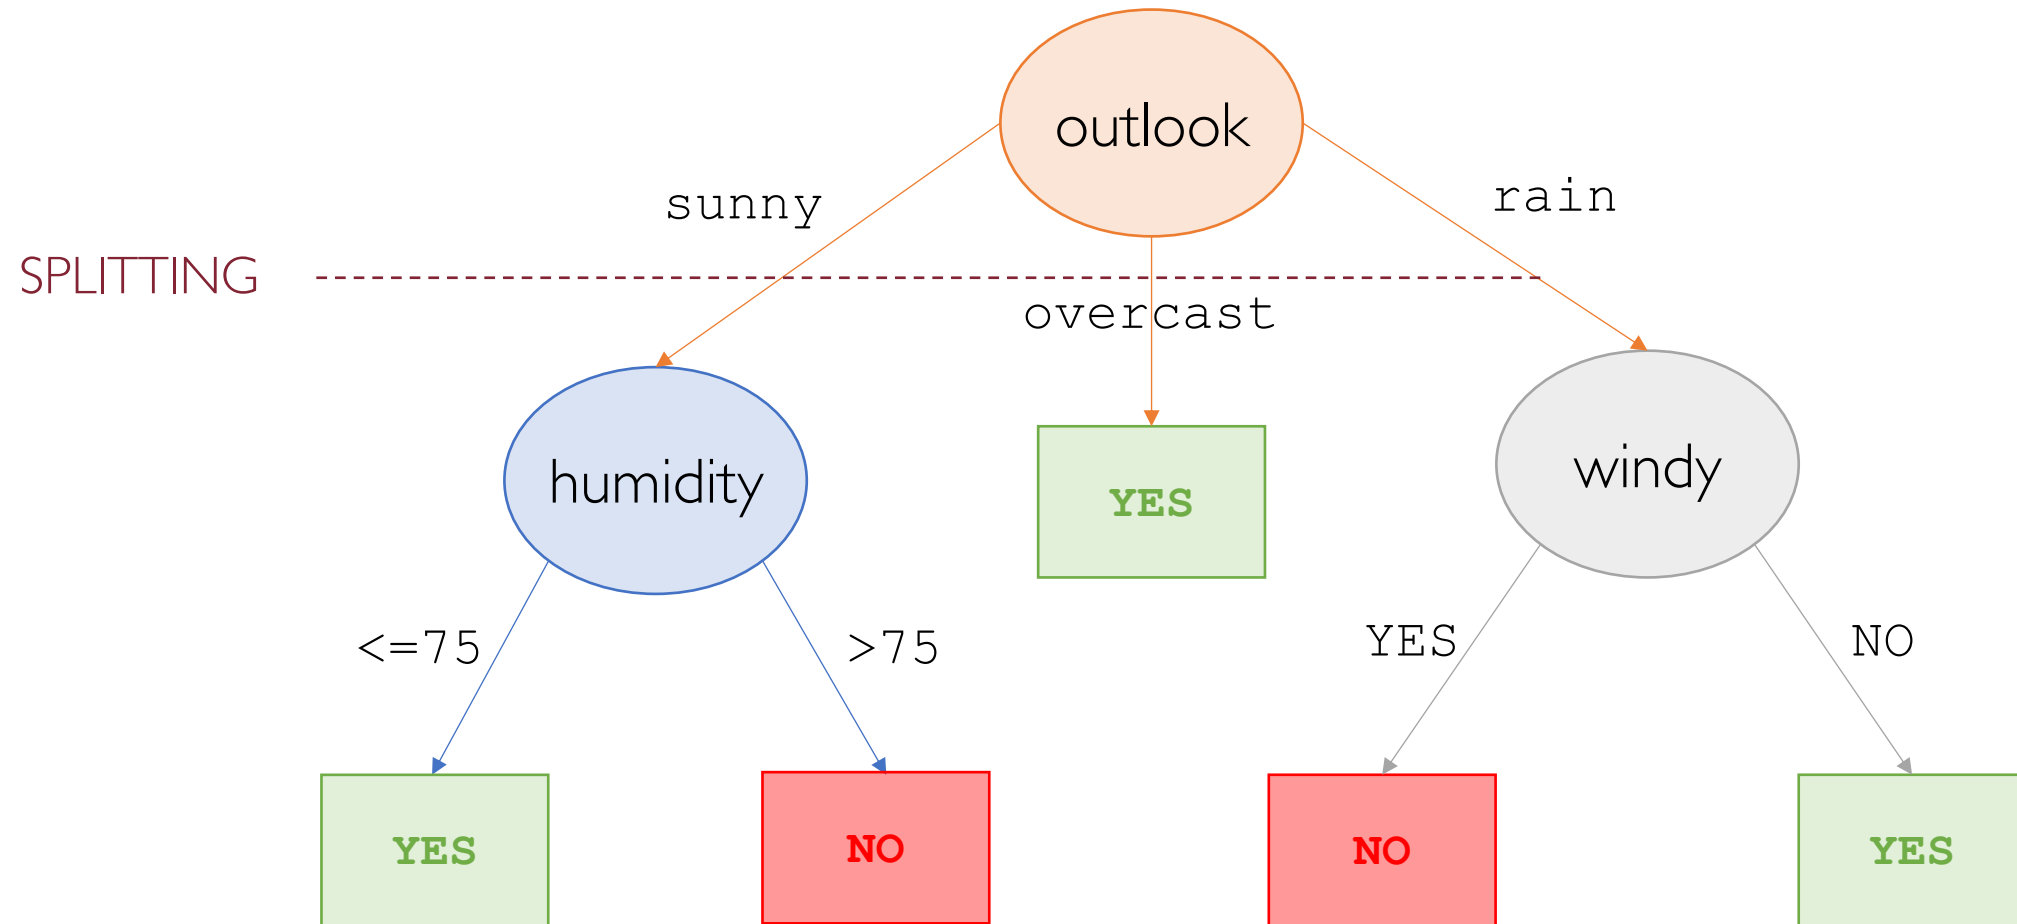
# Decision Tree: An Example



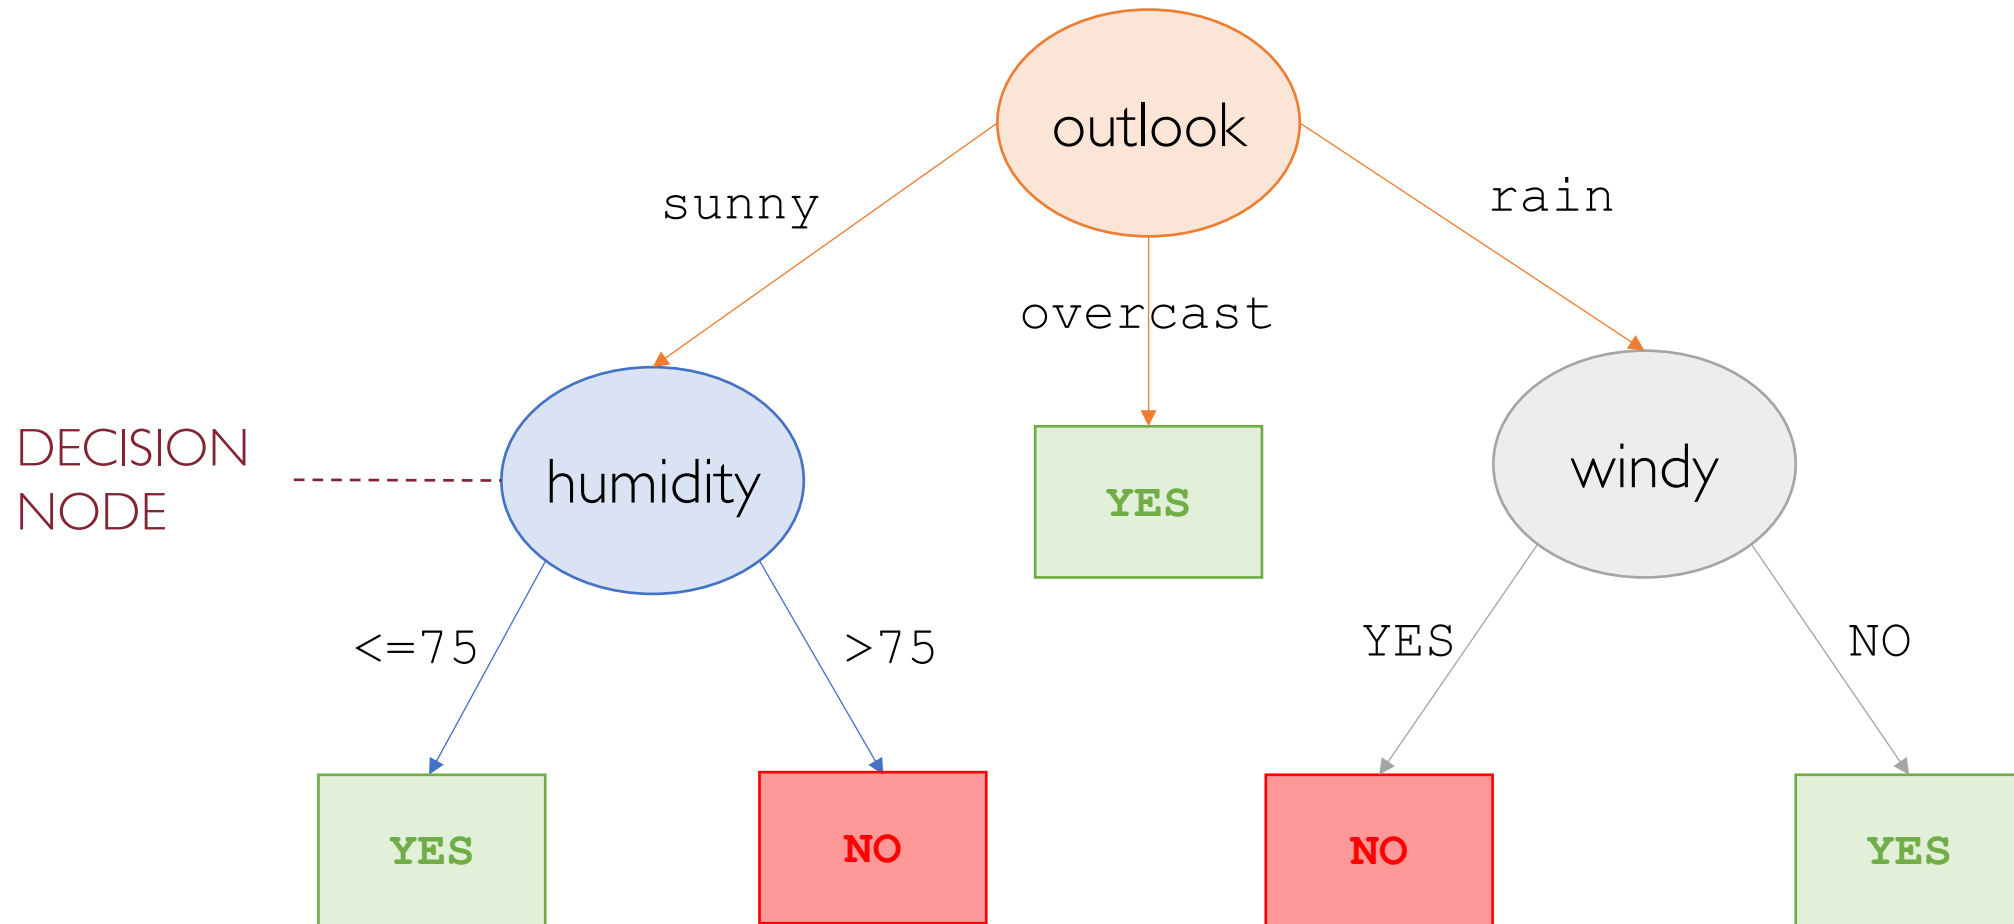
# Decision Tree: An Example



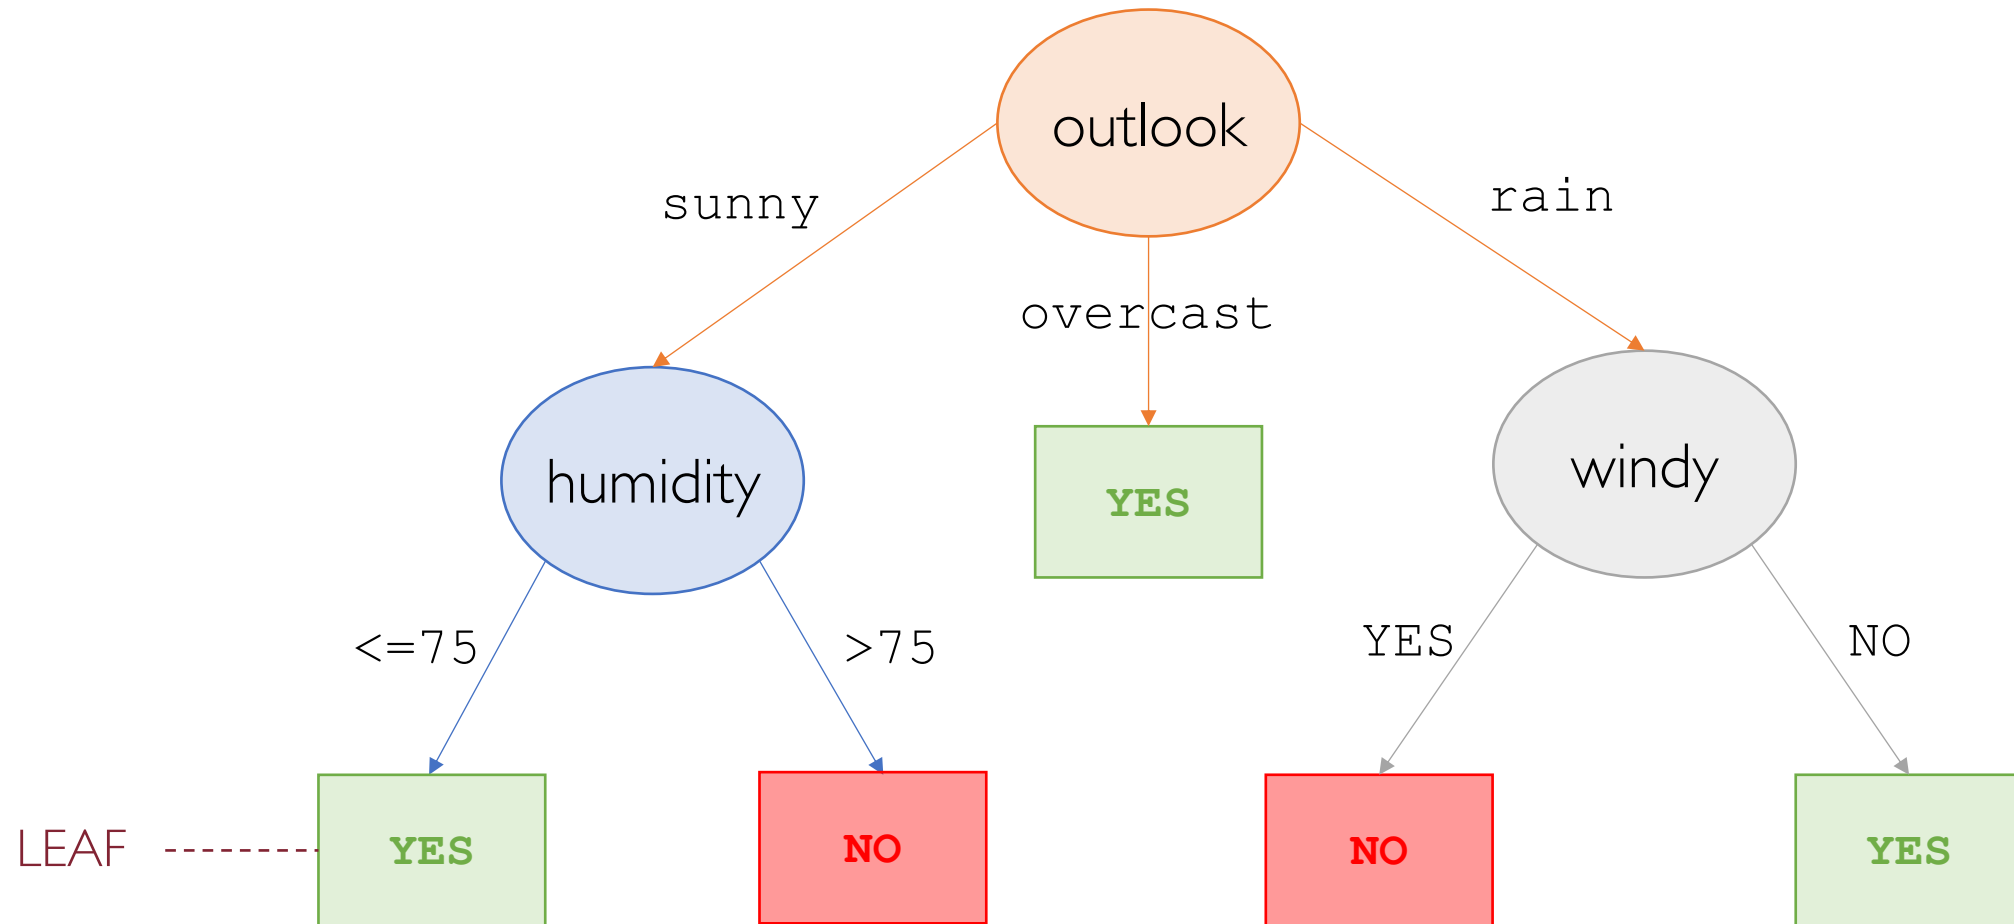
# Decision Tree: An Example



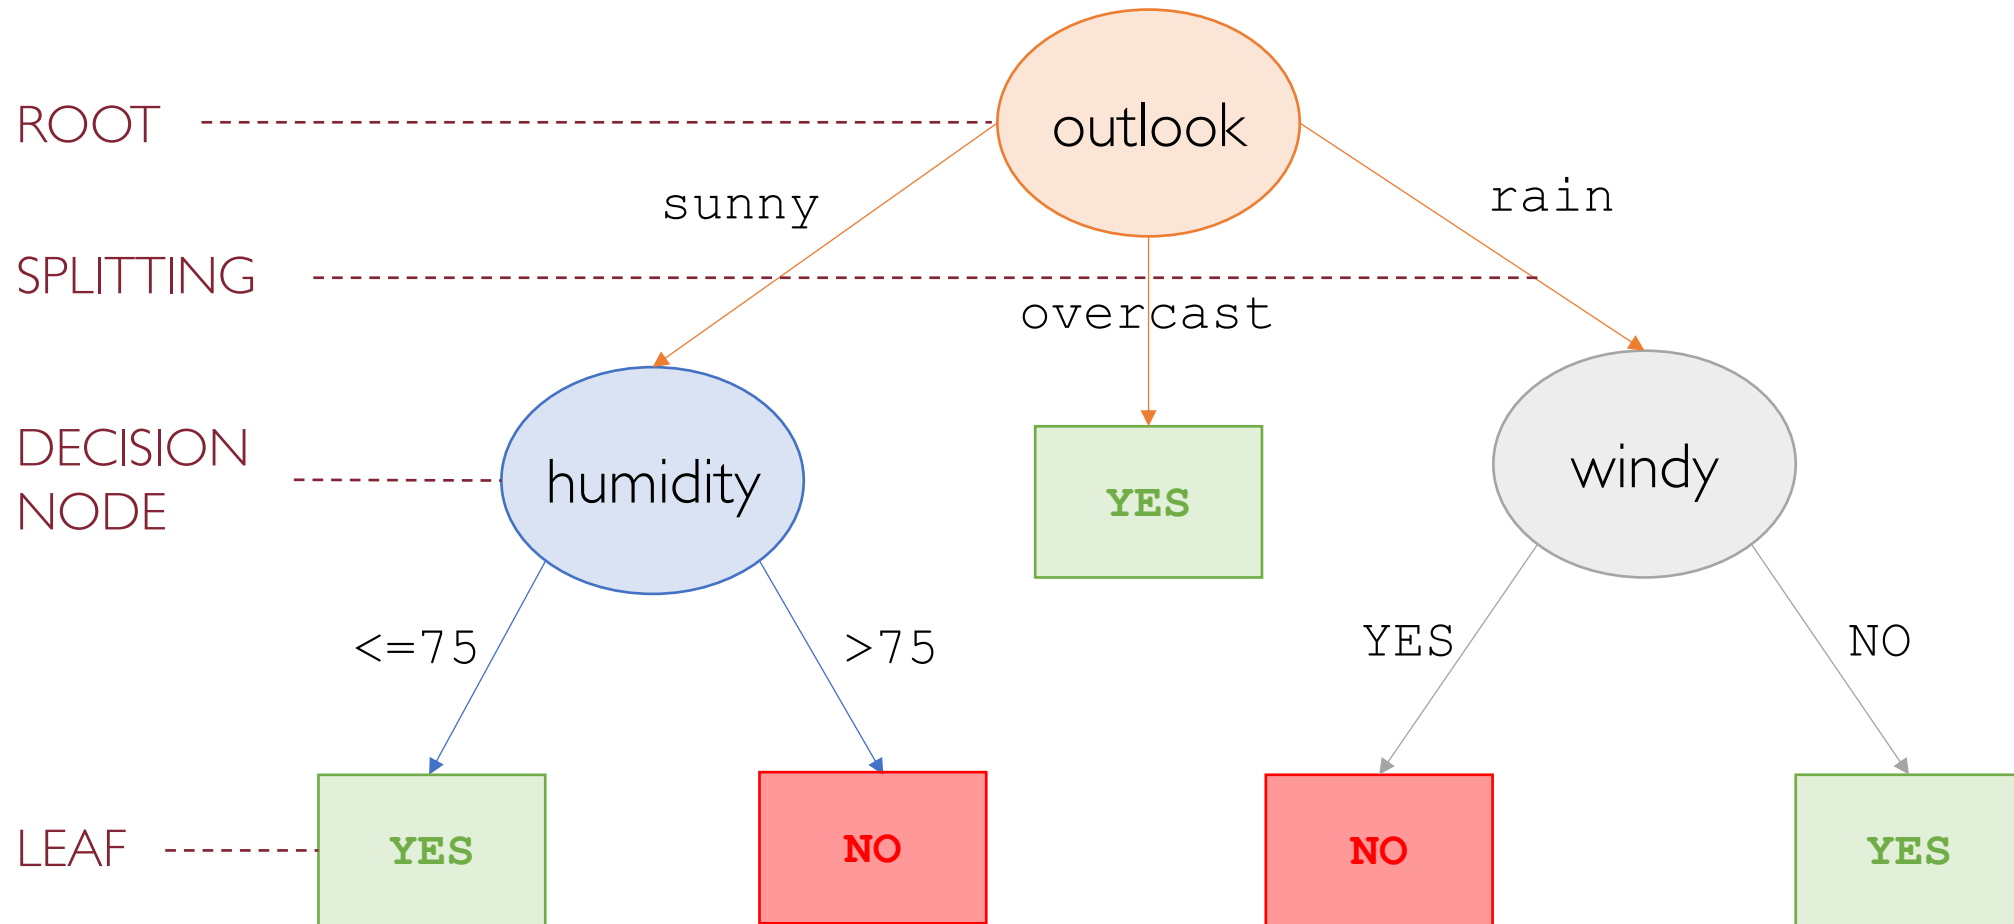
# Decision Tree: An Example



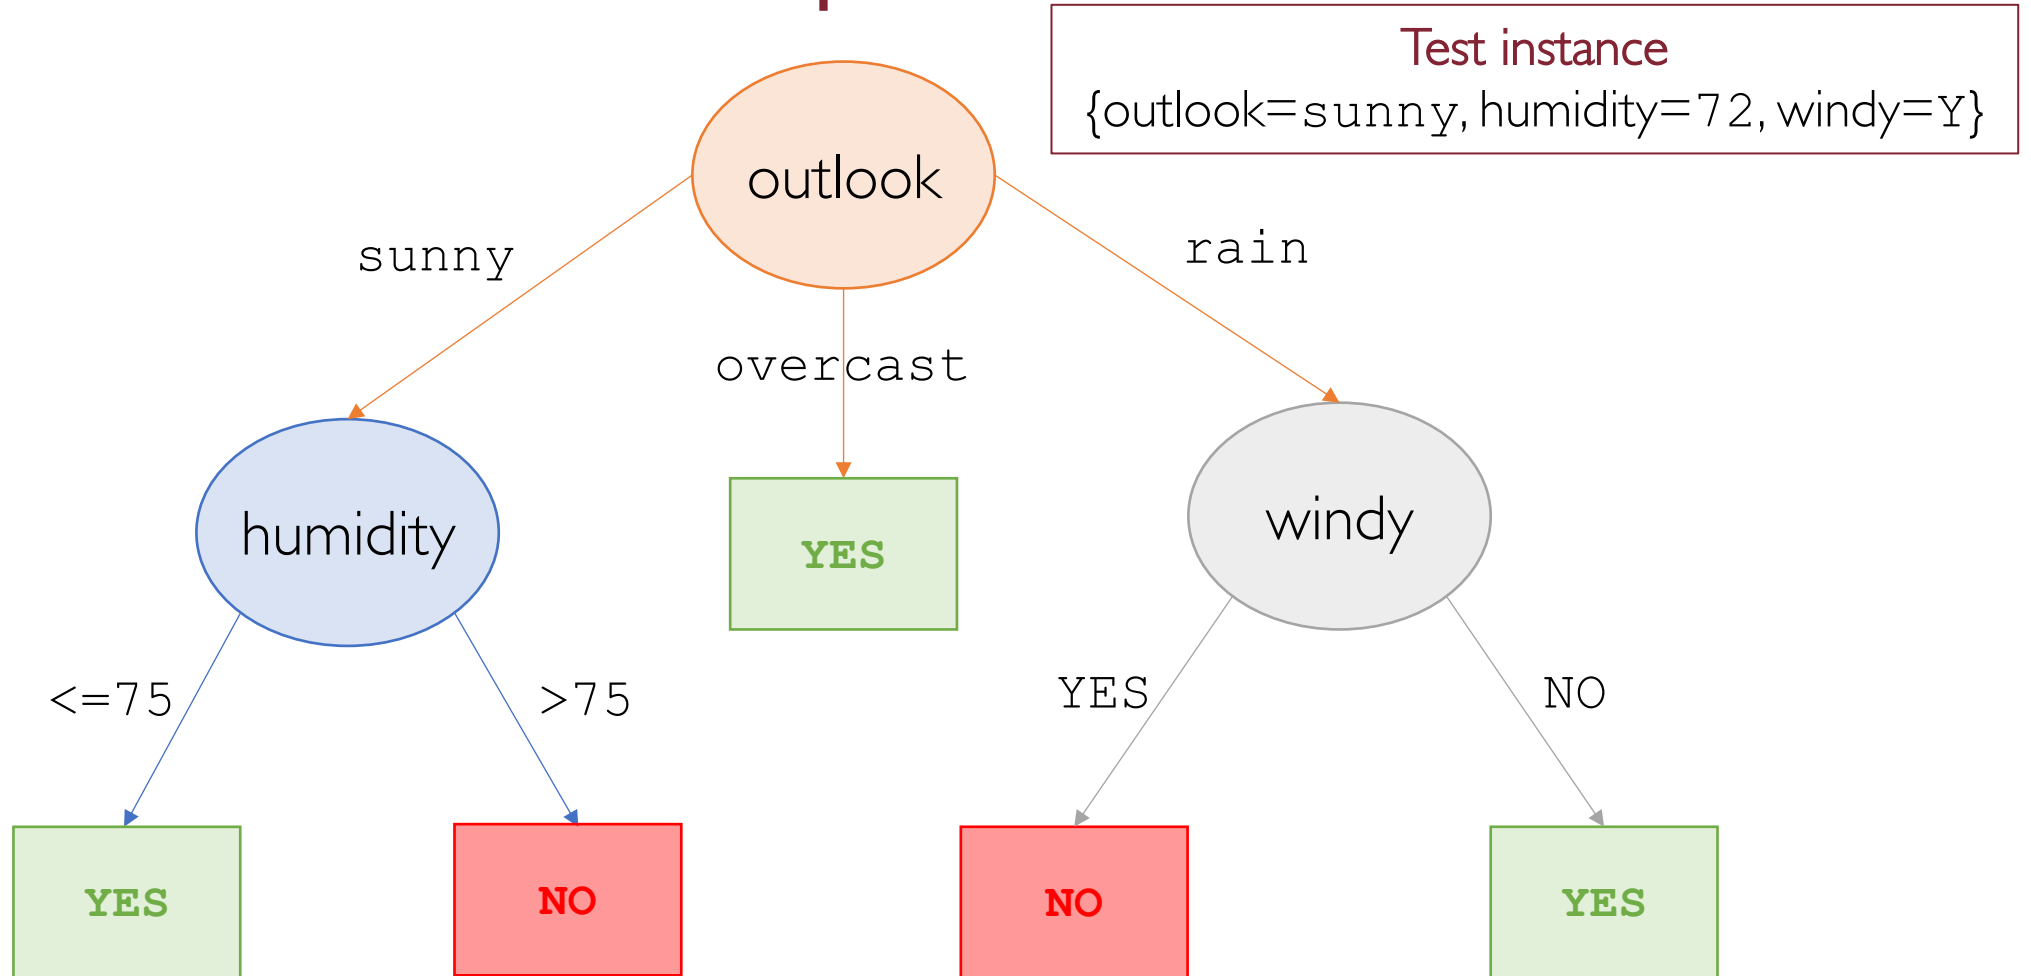
# Decision Tree: An Example



# Decision Tree: An Example

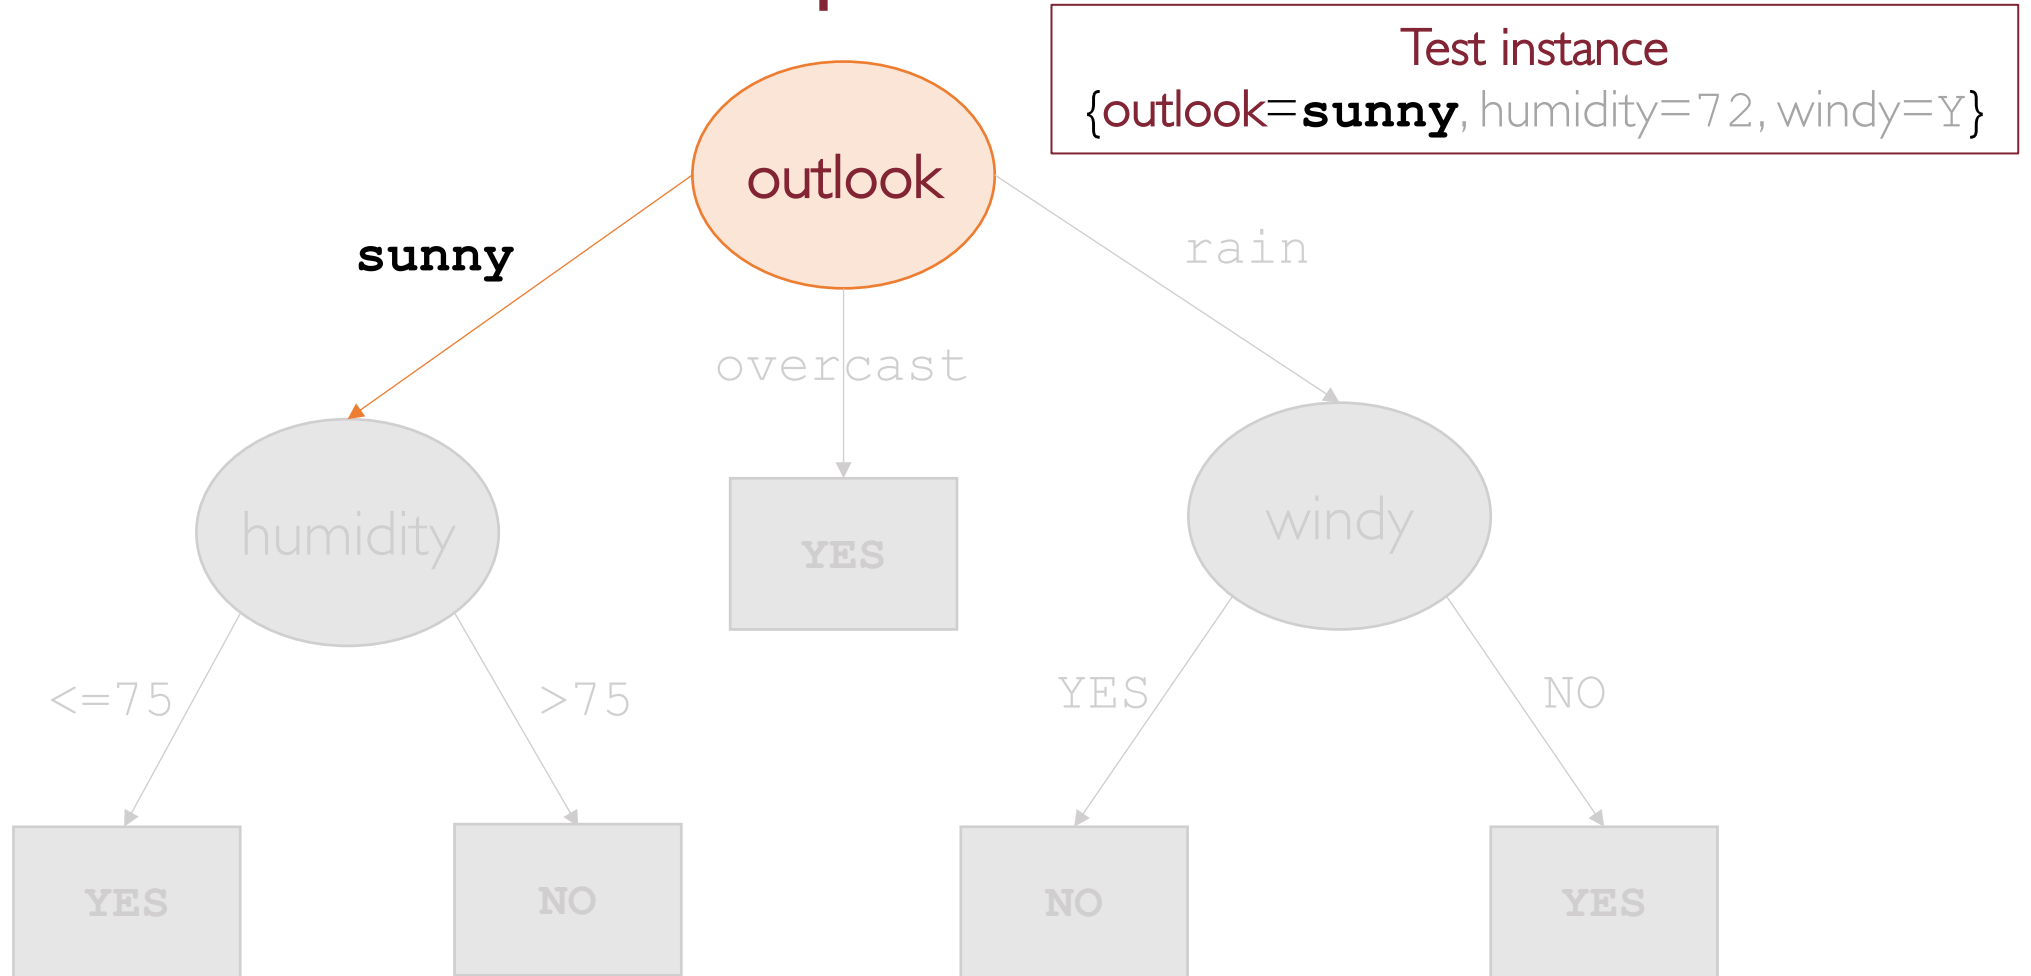


# Decision Tree: An Example

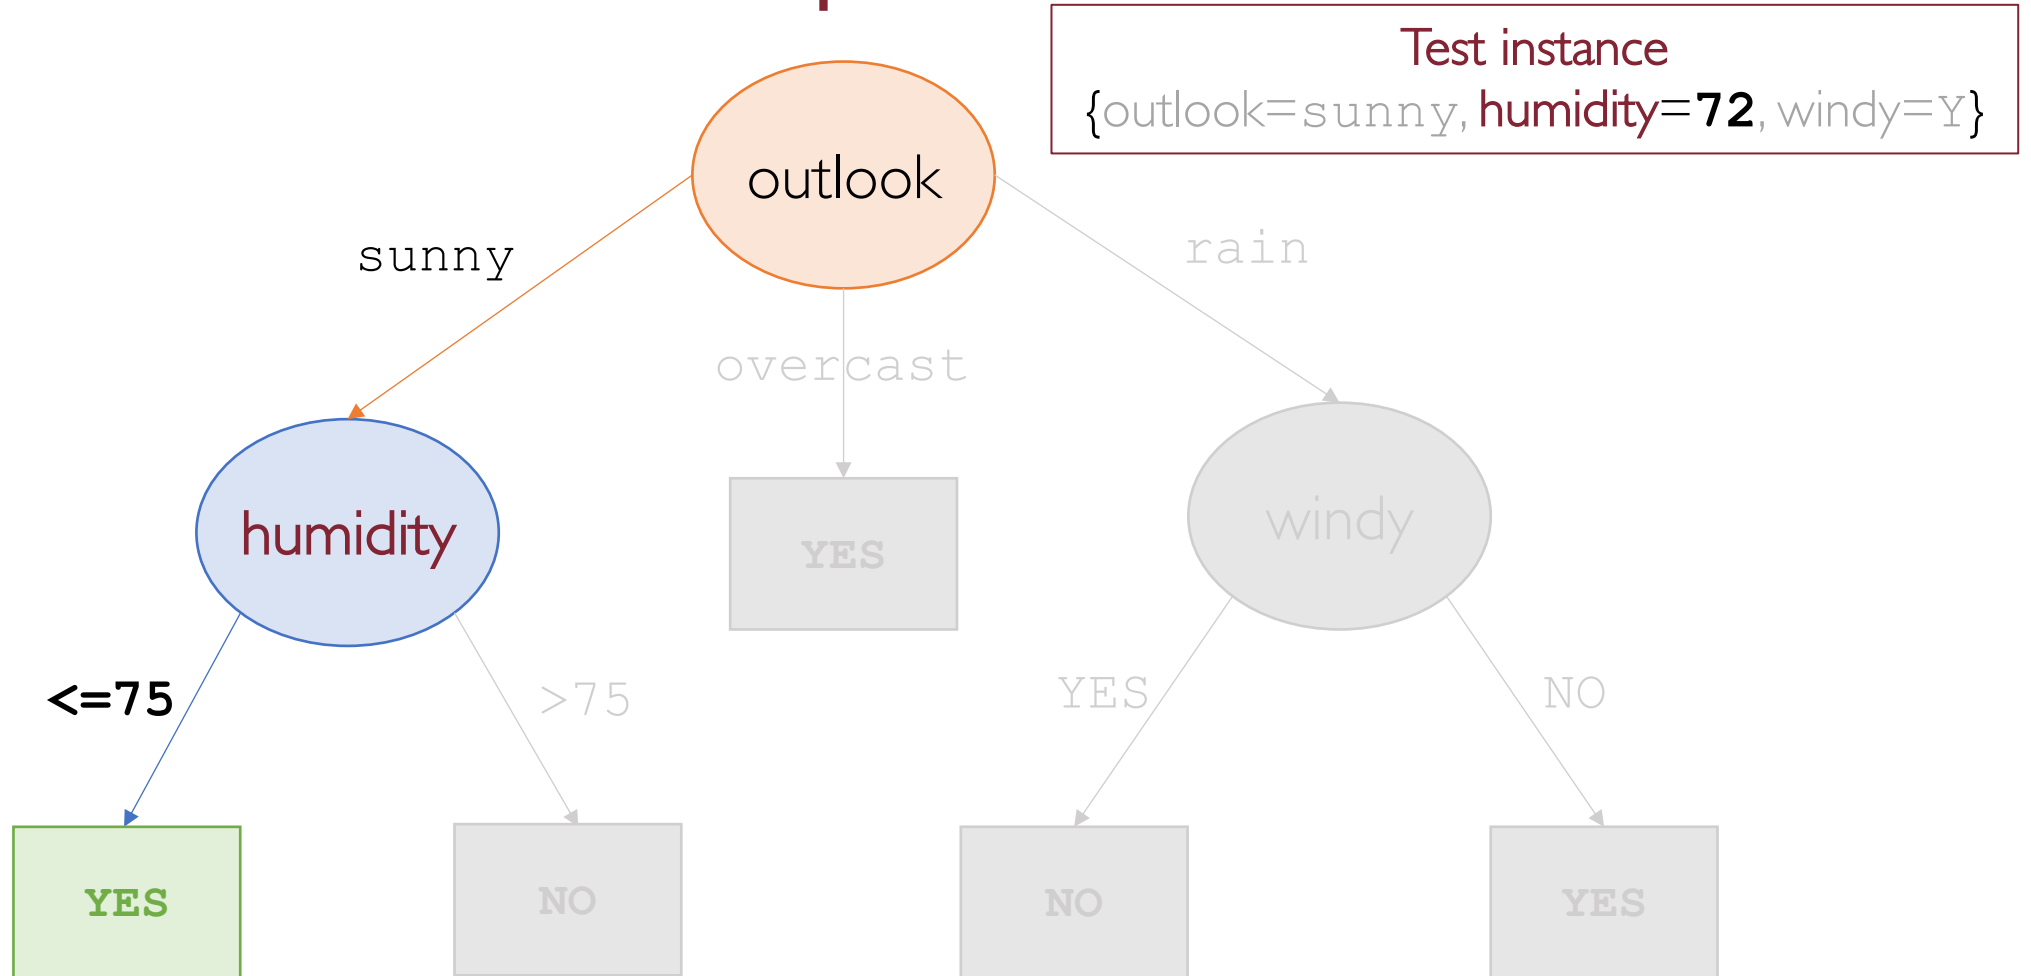




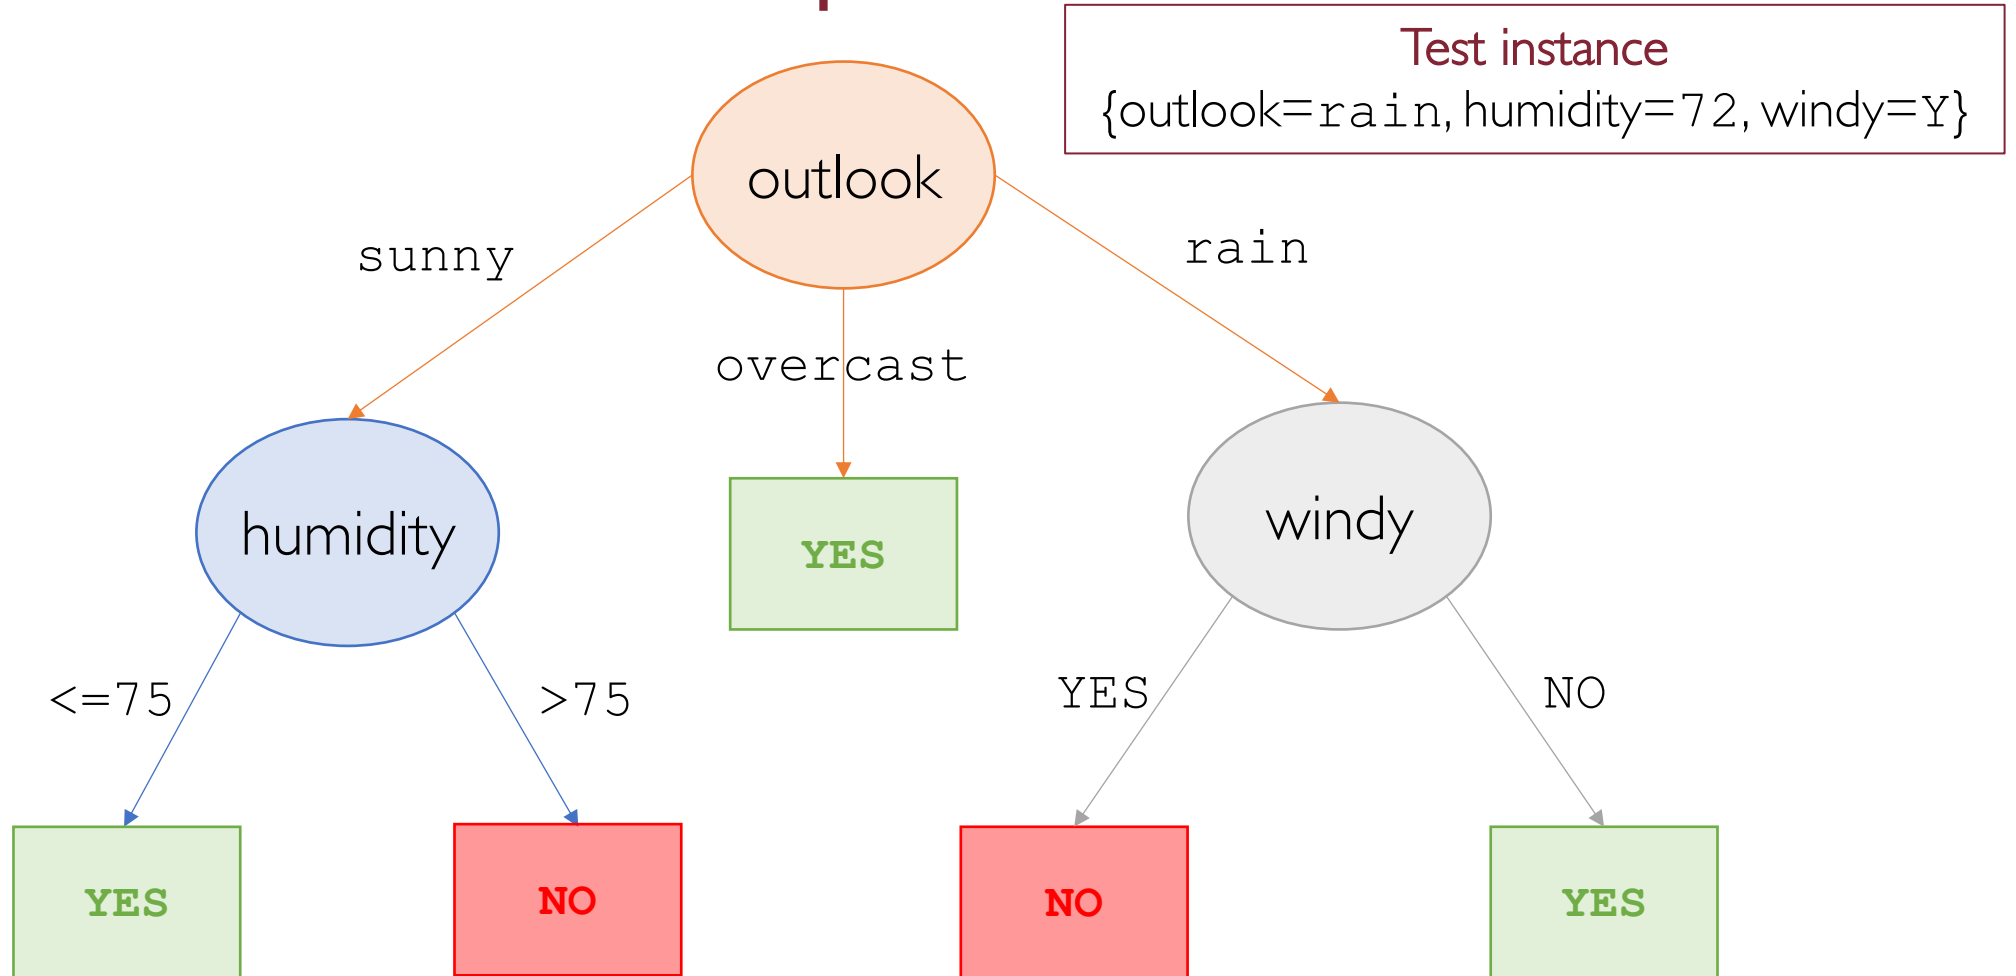
# Decision Tree: An Example



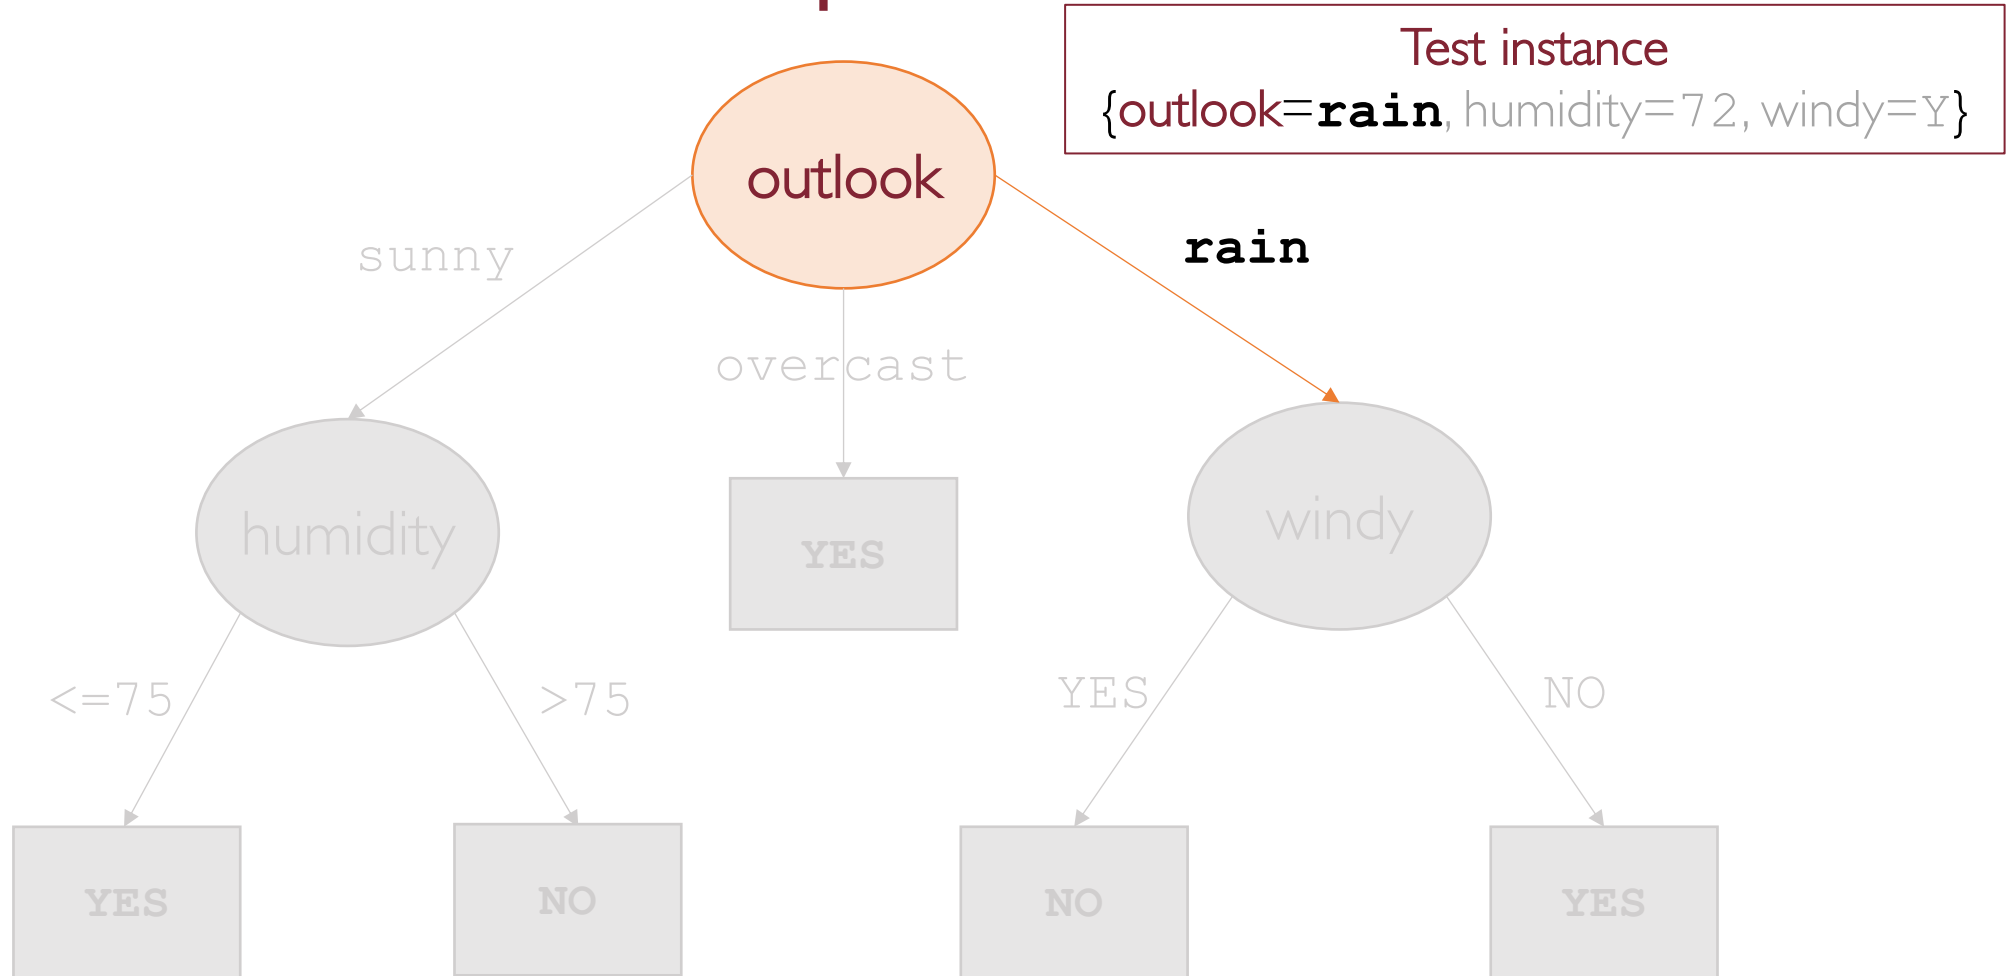
# Decision Tree: An Example



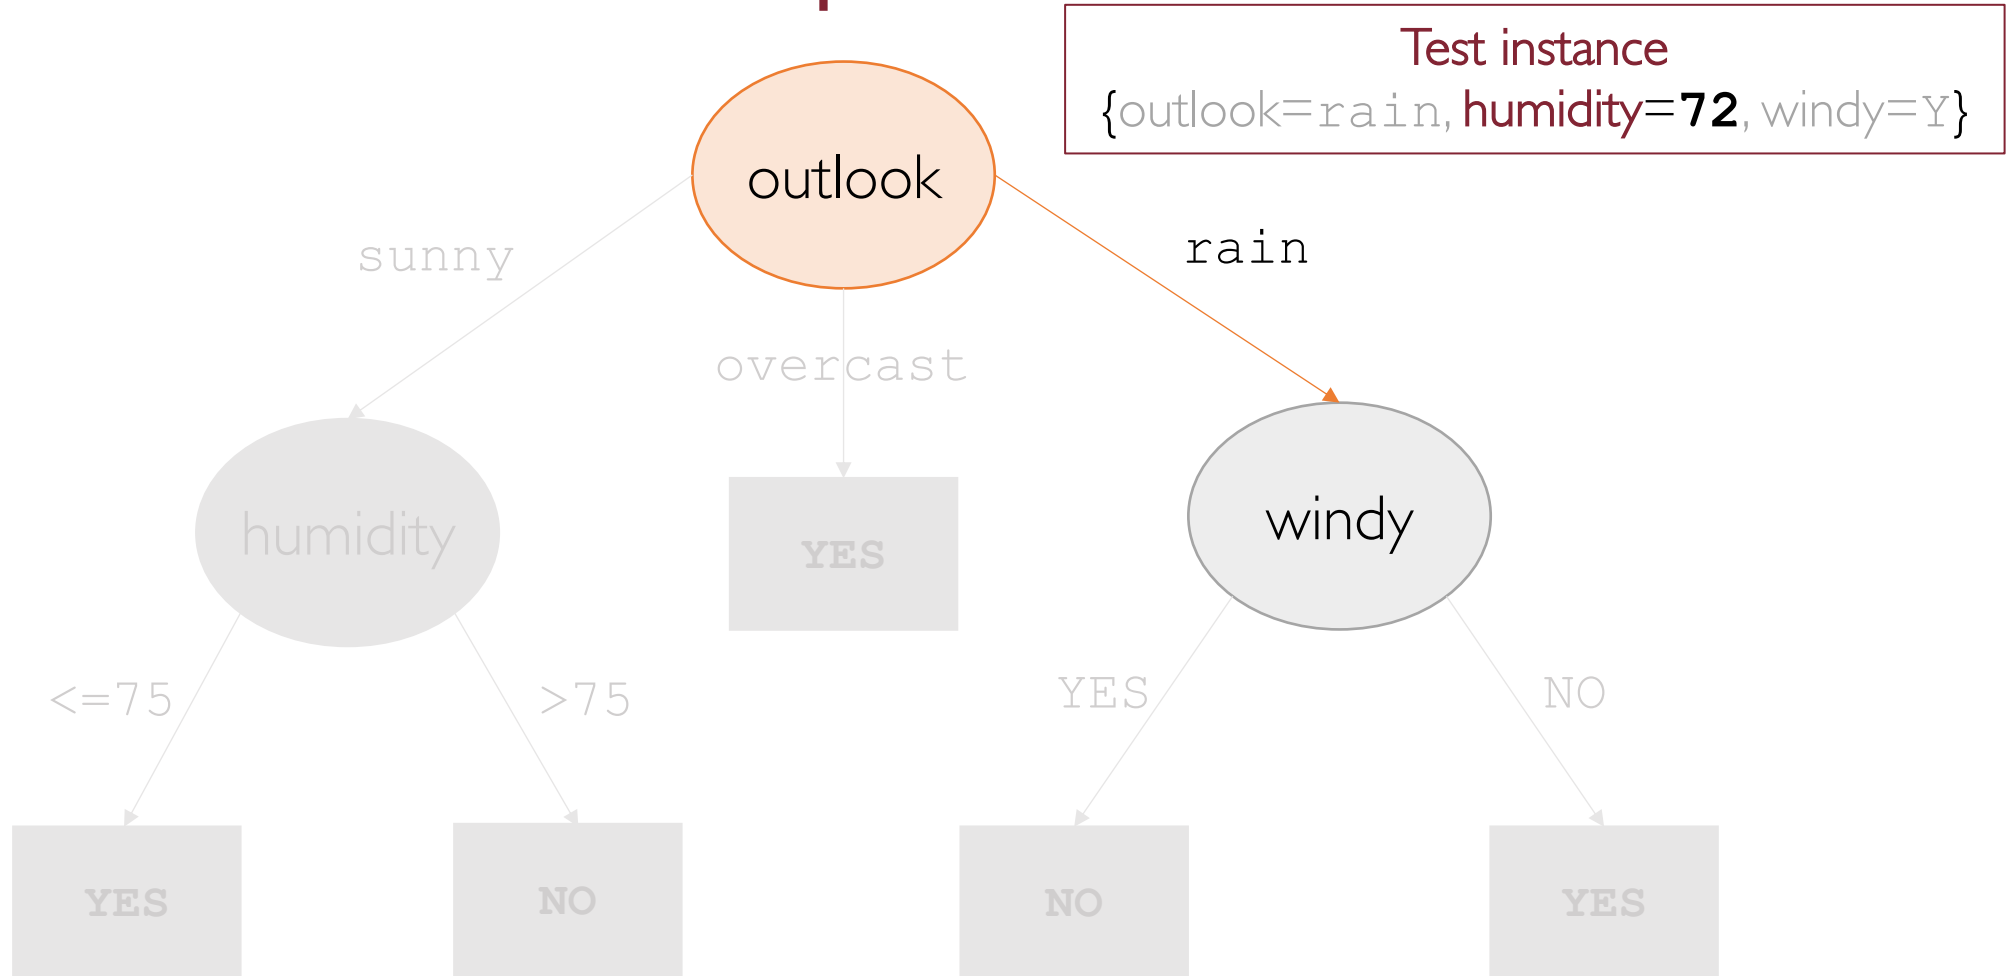
# Decision Tree: An Example



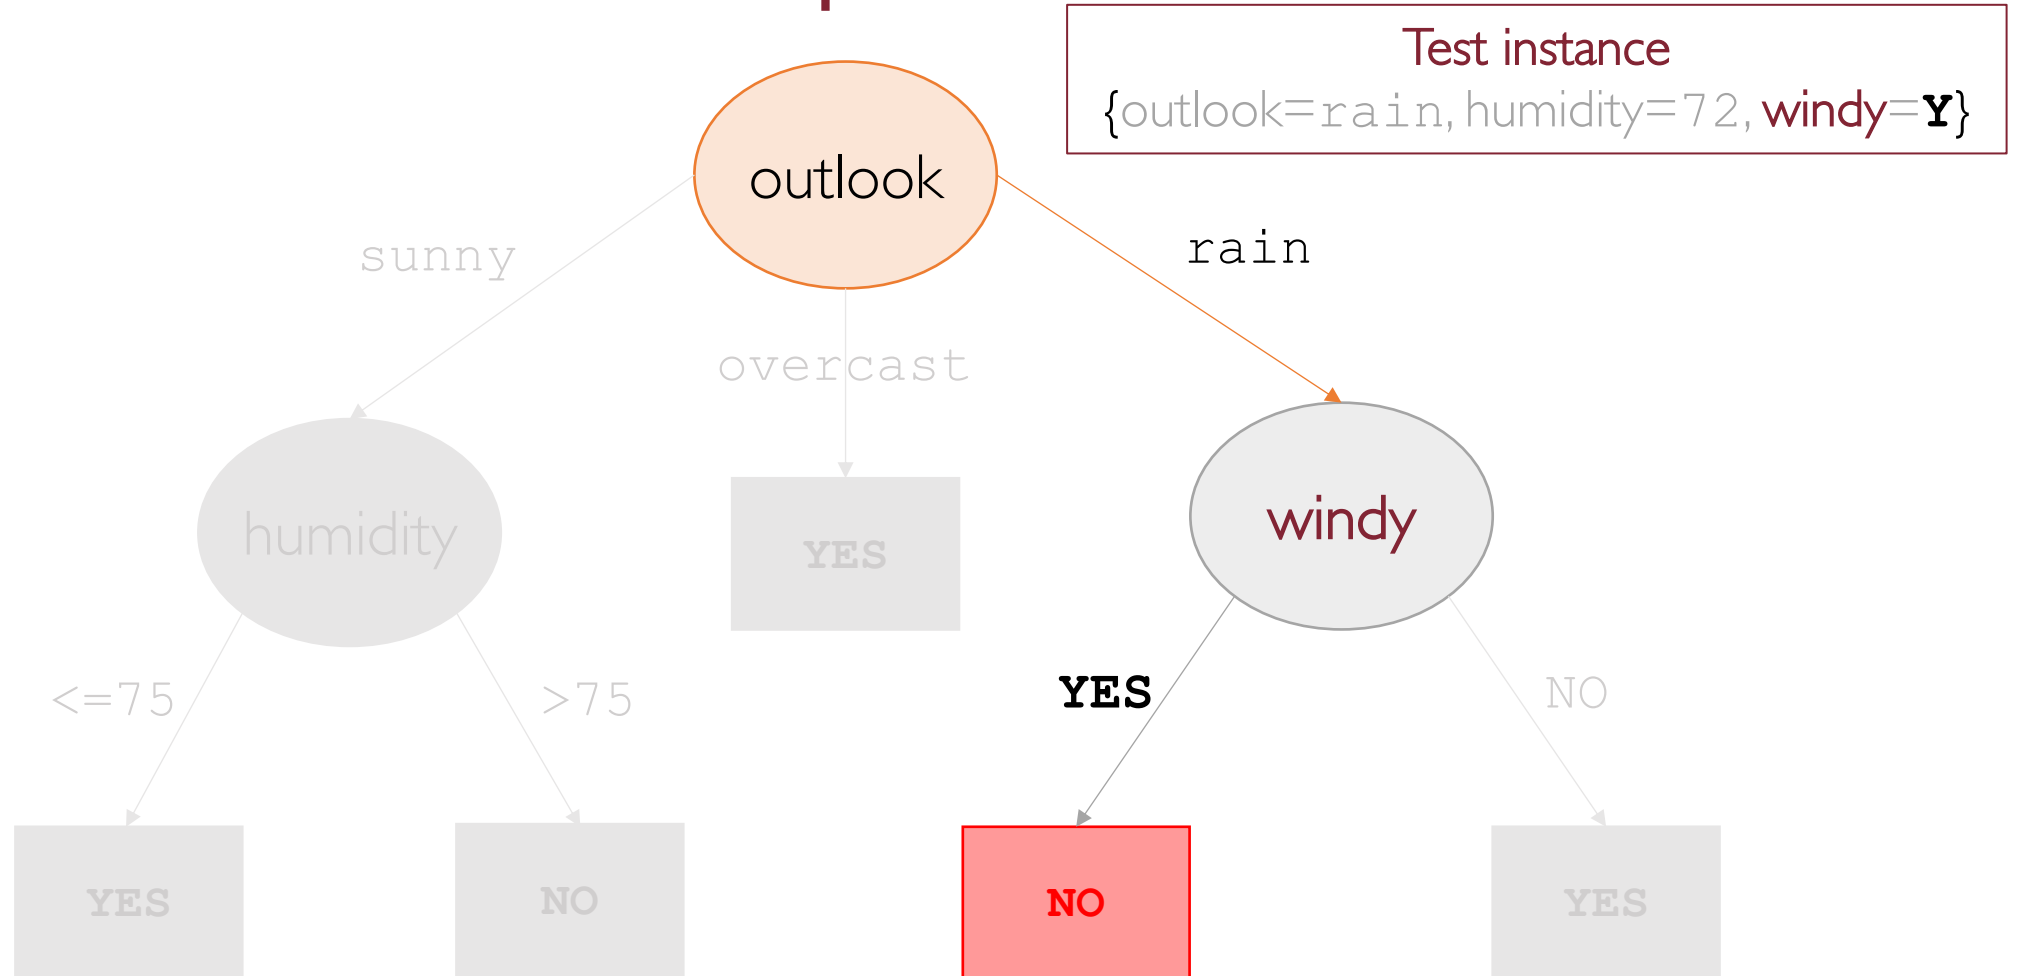
# Decision Tree: An Example



# Decision Tree: An Example



# Decision Tree: An Example



# A Bit of Notation

$$\mathcal{X} \subseteq \mathbb{R}^n$$

input feature space

$$\mathcal{Y}$$

output space

$$\mathcal{Y} \subseteq \mathbb{R}$$

real-value label (**regression**)

$$\mathcal{Y} = \{1, \dots, k\}$$

discrete-value label (**k-ary classification**)

$$(\mathbf{x}_i, y_i)$$

$i$ -th labeled instance

$$\mathbf{x}_i = (x_{i,1}, \dots, x_{i,n}) \in \mathcal{X}$$

$n$ -dimensional feature vector of the  $i$ -th instance

$$y_i \in \mathcal{Y}$$

label of the  $i$ -th instance

$$\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)\}$$

dataset of  $m$  **i.i.d.** labeled instances

# How Do We Build a Decision Tree?

- Split the input feature space (i.e., the set of possible values observed for each feature  $x_i$ ) into a set of non-overlapping regions  $R_1, R_2, \dots, R_j$



# How Do We Build a Decision Tree?

- Split the input feature space (i.e., the set of possible values observed for each feature  $x_i$ ) into a set of non-overlapping regions  $R_1, R_2, \dots, R_j$
- For every observation that falls into the region  $R_j$ , we make the same prediction, i.e., the mean of the response values in  $R_j$

# How Do We Build a Decision Tree?

- Split the input feature space (i.e., the set of possible values observed for each feature  $x_i$ ) into a set of non-overlapping regions  $R_1, R_2, \dots, R_j$
- For every observation that falls into the region  $R_j$ , we make the same prediction, i.e., the mean of the response values in  $R_j$
- Example:
  - Suppose we split the input feature space in 2 regions:  $R_1$  and  $R_2$  and the response mean as computed from  $R_1 = 10$  and  $R_2 = 20$
  - For any  $\mathbf{x}$  belonging to  $R_1$  ( $R_2$ ) will be predicted 10 (20)

# Partitioning the Input Space

- In theory, partitions could have *any* shape

# Partitioning the Input Space

- In theory, partitions could have *any* shape
- We choose to divide the input feature space into **hyper-rectangles** (i.e., high-dimensional rectangles)

# Partitioning the Input Space

- In theory, partitions could have *any* shape
- We choose to divide the input feature space into **hyper-rectangles** (i.e., high-dimensional rectangles)
- This way, the obtained model is easier to interpret

# Partitioning the Input Space

- In theory, partitions could have *any* shape
- We choose to divide the input feature space into **hyper-rectangles** (i.e., high-dimensional rectangles)
- This way, the obtained model is easier to interpret

Minimize the Residual Sum of Squares

$$\text{RSS} = \sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

# Partitioning the Input Space

- In theory, partitions could have *any* shape
- We choose to divide the input feature space into **hyper-rectangles** (i.e., high-dimensional rectangles)
- This way, the obtained model is easier to interpret

Minimize the Residual Sum of Squares

$$\text{RSS} = \sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

The mean computed from observations in  $R_j$

# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value

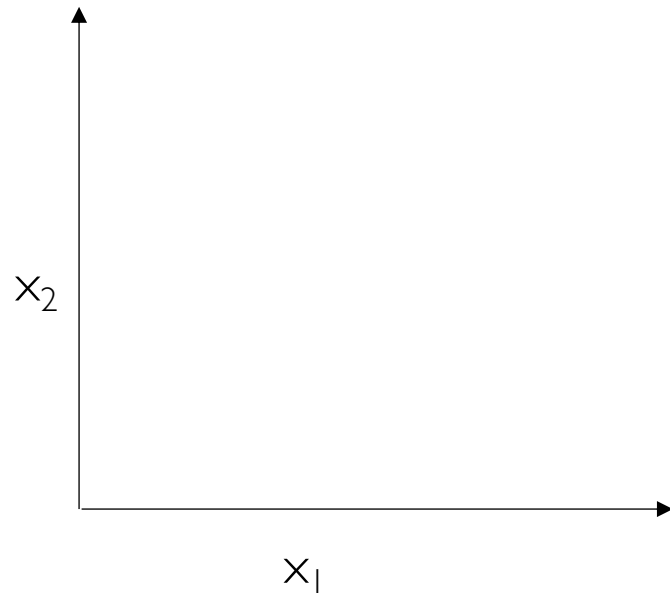


# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value

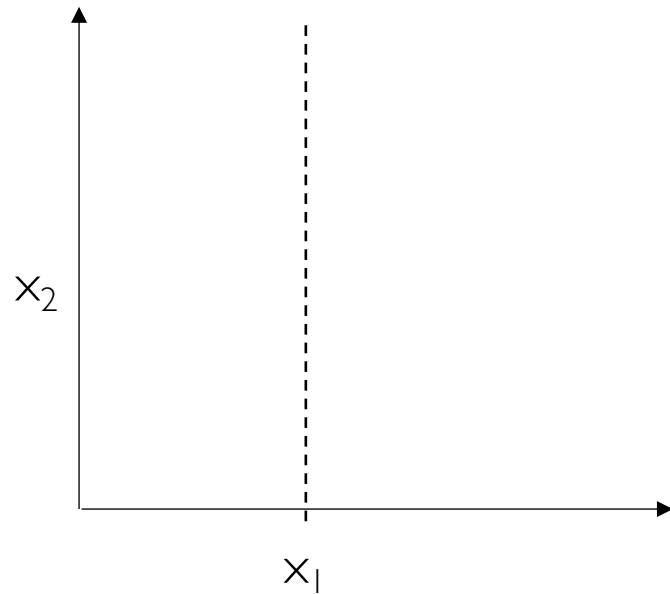


# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value

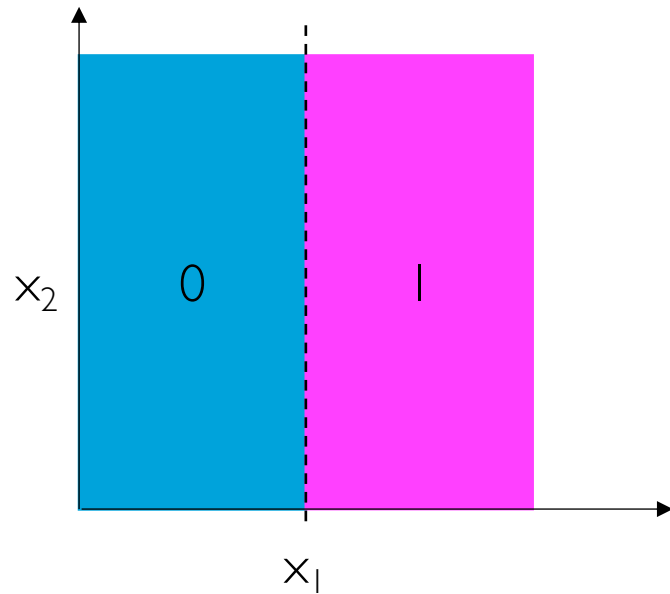


# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value

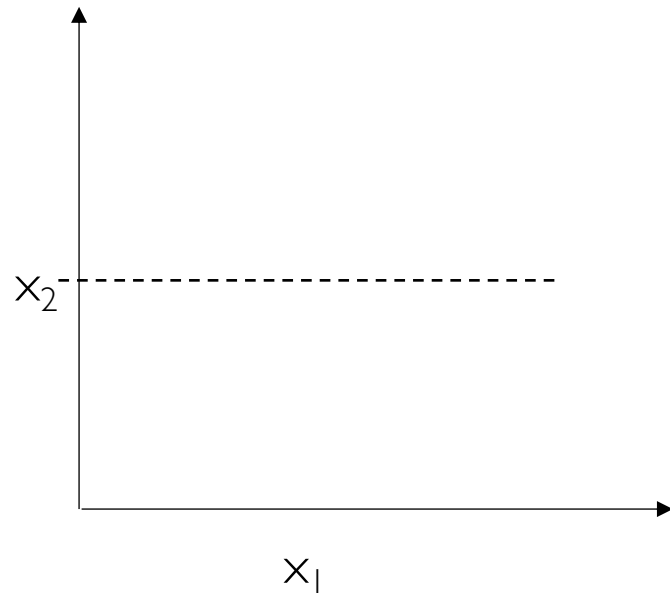


# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value

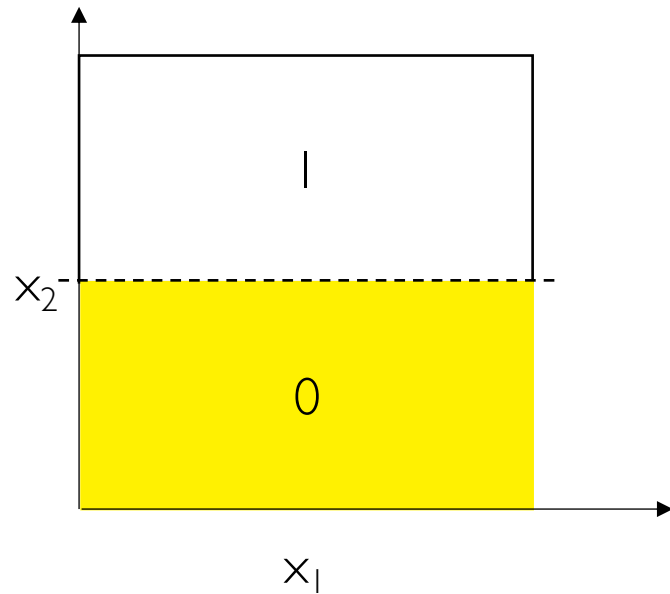


# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value




# Discrete vs. Continuous Inputs

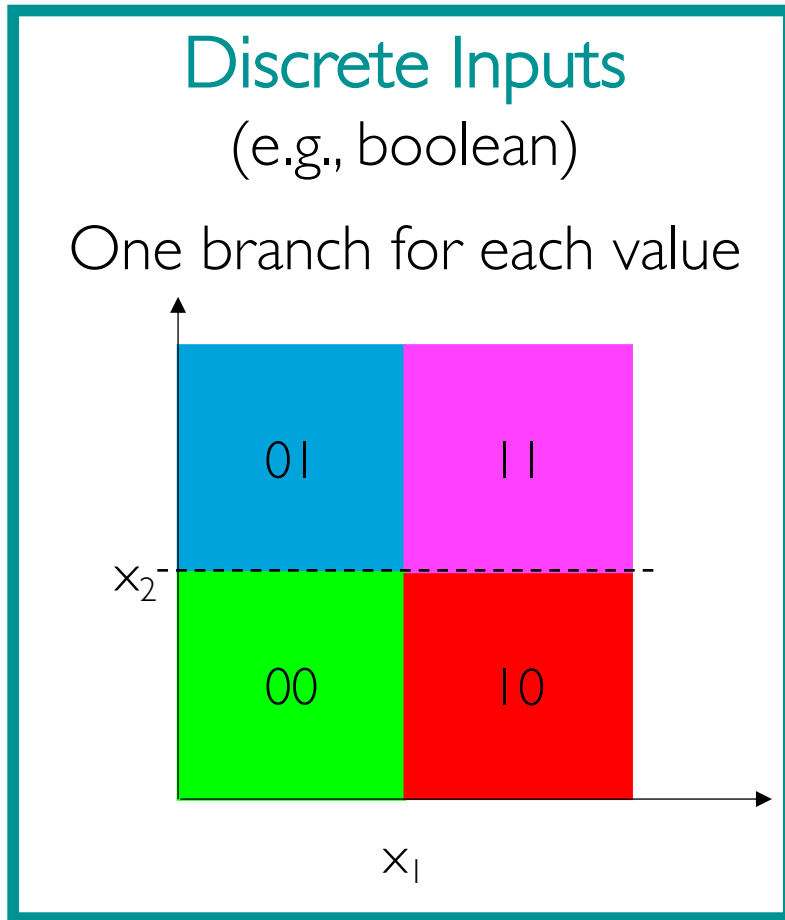
## Discrete Inputs

(e.g., boolean)

One branch for each value



# Discrete vs. Continuous Inputs

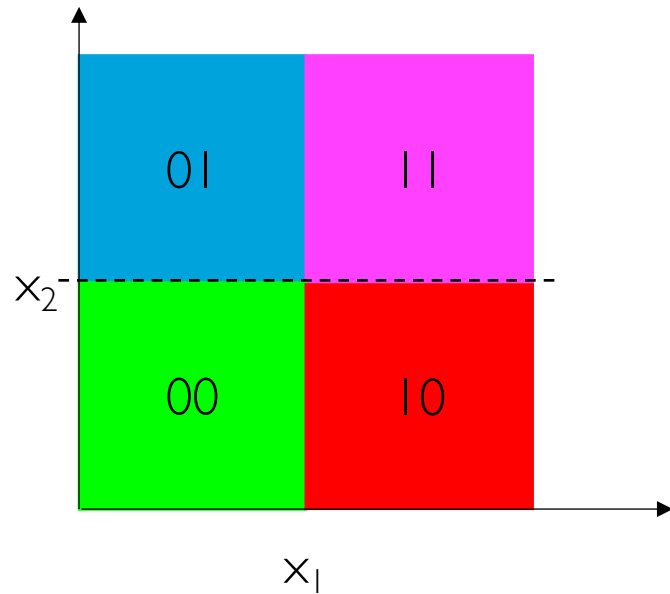


# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value




## Continuous Inputs

For each attribute, find a splitting point  $s$

Test  $x_j \leq s$  and create 2 branches

The same attribute may be further split in each subtree






# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value






## Continuous Inputs

For each attribute, find a splitting point  $s$

Test  $x_j \leq s$  and create 2 branches

The same attribute may be further split in each subtree




# Discrete vs. Continuous Inputs

## Discrete Inputs

(e.g., boolean)

One branch for each value






## Continuous Inputs

For each attribute, find a splitting point  $s$

Test  $x_j \leq s$  and create 2 branches

The same attribute may be further split in each subtree



# Expressiveness

Discrete Inputs/Discrete Outputs

# Expressiveness

## Discrete Inputs/Discrete Outputs

Decision Trees can express any function of the inputs


# Expressiveness

## Discrete Inputs/Discrete Outputs

Decision Trees can express any function of the inputs

Example: Boolean Functions

A	B	A xor B
F	F	F
F	T	T
T	F	T
T	T	F



# Expressiveness


## Discrete Inputs/Discrete Outputs

Decision Trees can express any function of the inputs

Example: Boolean Functions

Truth table

A	B	A xor B
F	F	F
F	T	T
T	F	T
T	T	F




# Expressiveness

## Discrete Inputs/Discrete Outputs

Decision Trees can express any function of the inputs

Example: Boolean Functions

A	B	A xor B
F	F	F
F	T	T
T	F	T
T	T	F



Each row of the truth table maps to a root-to-leaf path on the tree

# Expressiveness

Continuous Inputs/Continuous Outputs



# Expressiveness

## Continuous Inputs/Continuous Outputs

Decision Trees can *approximate* any function arbitrarily closely

# Expressiveness

## Continuous Inputs/Continuous Outputs

Decision Trees can *approximate* any function arbitrarily closely

Trivially, there exists a consistent decision tree for any training set

# Expressiveness

## Continuous Inputs/Continuous Outputs

Decision Trees can *approximate* any function arbitrarily closely

Trivially, there exists a consistent decision tree for any training set

Such a tree will have one dedicated root-to-leaf path for each training instance

# Expressiveness

## Continuous Inputs/Continuous Outputs

Decision Trees can *approximate* any function arbitrarily closely

Trivially, there exists a consistent decision tree for any training set

Such a tree will have one dedicated root-to-leaf path for each training instance

Of course, this tree clearly overfits the training data and it will not generalize to unseen examples (needs regularization)

# Hypothesis Space

How many distinct decision trees can be built out of  $n$  boolean attributes?

# Hypothesis Space

How many distinct decision trees can be built out of  $n$  boolean attributes?

This is equivalent to say how many boolean functions can be defined

# Hypothesis Space

How many distinct decision trees can be built out of  $n$  boolean attributes?

This is equivalent to say how many boolean functions can be defined

Each boolean function of  $n$  boolean inputs is represented by a truth table

$x_1$	$x_1$	...	$x_n$	$y$
0	0	...	0	
1	0	...	0	
1	1	...	0	
...	...	...	...	
...	...	...	...	
...	...	...	...	
1	1	1	1	

# Hypothesis Space

How many distinct decision trees can be built out of  $n$  boolean attributes?

This is equivalent to say how many boolean functions can be defined

Each boolean function of  $n$  boolean inputs is represented by a truth table

$x_1$	$x_1$	...	$x_n$	$y$
0	0	...	0	
1	0	...	0	
1	1	...	0	
...	...	...	...	
...	...	...	...	
...	...	...	...	
1	1	1	1	

$2^n$  rows



# Hypothesis Space

How many distinct decision trees can be built out of  $n$  boolean attributes?

This is equivalent to say how many boolean functions can be defined

Each boolean function of  $n$  boolean inputs is represented by a truth table

$2^n$  rows

$x_1$	$x_1$	...	$x_n$	$y$
0	0	...	0	
1	0	...	0	
1	1	...	0	
...	...	...	...	
...	...	...	...	
...	...	...	...	
1	1	1	1	

For each input  $y = 0$  or  $1$

# Hypothesis Space

A possible boolean function is the one which will output all 0s

2 <sup>n</sup> rows	$x_1$	$x_1$	...	$x_n$	$y$
	0	0	...	0	0
	1	0	...	0	0
	1	1	...	0	0
	...	...	...	...	0
	...	...	...	...	0
	...	...	...	...	0
	1	1	1	1	0

# Hypothesis Space

Another possible boolean function is the one which will output all 1s

$x_1$	$x_1$	...	$x_n$	$y$
0	0	...	0	1
1	0	...	0	1
1	1	...	0	1
...	...	...	...	1
...	...	...	...	1
...	...	...	...	1
1	1	1	1	1

$2^n$  rows

# Hypothesis Space

Overall, there are  $2^{(2^n)}$  possible boolean functions, therefore decision trees

# Hypothesis Space

Overall, there are  $2^{(2^n)}$  possible boolean functions, therefore decision trees

e.g., with just 6 boolean attributes, there are 18,446,744,073,709,551,616 trees

# Hypothesis Space

Overall, there are  $2^{(2^n)}$  possible boolean functions, therefore decision trees

e.g., with just 6 boolean attributes, there are 18,446,744,073,709,551,616 trees

The hypothesis space defined by decisions tree is **very expressive** because there are a lot of different functions it can represent

# Hypothesis Space

Overall, there are  $2^{(2^n)}$  possible boolean functions, therefore decision trees

e.g., with just 6 boolean attributes, there are 18,446,744,073,709,551,616 trees

The hypothesis space defined by decisions tree is **very expressive** because there are a lot of different functions it can represent

Larger hypothesis space means also it is generally harder for the learning algorithm to find the best hypothesis (larger space to explore)

# Partitioning the Input Space: NP-Complete

- Unfortunately, minimizing RSS would require exploring all the possible partitions of the feature space into  $J$  boxes



# Partitioning the Input Space: NP-Complete

- Unfortunately, minimizing RSS would require exploring all the possible partitions of the feature space into  $J$  boxes
- The problem is known to be [NP-Complete](#), therefore computationally unfeasible

# Partitioning the Input Space: NP-Complete

- Unfortunately, minimizing RSS would require exploring all the possible partitions of the feature space into  $J$  boxes
- The problem is known to be [NP-Complete](#), therefore computationally unfeasible



## Solution

Top-Down greedy heuristic Recursive Binary Splitting

# Recursive Binary Splitting (RBS)

- Initially, all the observations belong to the **root** of the tree

# Recursive Binary Splitting (RBS)

- Initially, all the observations belong to the **root** of the tree
- Split the input feature space into 2 subtrees (i.e., **left** and **right**)

# Recursive Binary Splitting (RBS)

- Initially, all the observations belong to the **root** of the tree
- Split the input feature space into 2 subtrees (i.e., **left** and **right**)
- Recursively repeat the step above on both subtrees

# Recursive Binary Splitting (RBS)

- Initially, all the observations belong to the **root** of the tree
- Split the input feature space into 2 subtrees (i.e., **left** and **right**)
- Recursively repeat the step above on both subtrees

top-down  
↓

# Recursive Binary Splitting (RBS)

- Initially, all the observations belong to the **root** of the tree
- Split the input feature space into 2 subtrees (i.e., **left** and **right**)
- Recursively repeat the step above on both subtrees
- Greedy strategy:
  - At each step, the best "local" split is made
  - Looking ahead might result in a different split, which leads to a better tree

top-down  
↓

# How to Choose the Split?

- Intuitively, we need to select the feature  $f$  and a splitting point  $s$  (of  $f$ )



# How to Choose the Split?

- Intuitively, we need to select the feature  $f$  and a splitting point  $s$  (of  $f$ )
- Such a splitting will lead to the partitioning of the feature space into the following regions:

$$R_{\text{left}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} \leq s\} \quad R_{\text{right}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} > s\}$$

# How to Choose the Split?

- Intuitively, we need to select the feature  $f$  and a splitting point  $s$  (of  $f$ )
- Such a splitting will lead to the partitioning of the feature space into the following regions:

$$R_{\text{left}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} \leq s\} \quad R_{\text{right}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} > s\}$$

- Among all the possible splittings, we select the one which reduces RSS **the most**

# How to Choose the Split?

- Intuitively, we need to select the feature  $f$  and a splitting point  $s$  (of  $f$ )
- Such a splitting will lead to the partitioning of the feature space into the following regions:

$$R_{\text{left}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} \leq s\} \quad R_{\text{right}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} > s\}$$

- Among all the possible splittings, we select the one which reduces RSS **the most**

$$\{x_{i,f} \leq s\}$$

is the region of the feature space in which the  $f$ -th feature takes on values less than or equal to  $s$

# How to Choose the Split?

- For each feature  $f = 1, \dots, d$  we look at every possible splitting value  $s$  (i.e., all the values we observed in the training set for the feature  $f$ )

# How to Choose the Split?

- For each feature  $f = 1, \dots, d$  we look at every possible splitting value  $s$  (i.e., all the values we observed in the training set for the feature  $f$ )
- For each pair of feature and splitting value  $(f, s)$ , we obtain 2 regions:

$$R_{\text{left}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} \leq s\} \quad R_{\text{right}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} > s\}$$

# How to Choose the Split?

- For each feature  $f = 1, \dots, d$  we look at every possible splitting value  $s$  (i.e., all the values we observed in the training set for the feature  $f$ )
- For each pair of feature and splitting value  $(f, s)$ , we obtain 2 regions:

$$R_{\text{left}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} \leq s\} \quad R_{\text{right}}(f, s) = \{\mathbf{x}_i \mid x_{i,f} > s\}$$

**Goal:** find the pair  $(f, s)$  which minimizes the following

$$\sum_{i: \mathbf{x}_i \in R_{\text{left}}(f, s)} (y_i - \hat{y}_{R_{\text{left}}})^2 + \sum_{i: \mathbf{x}_i \in R_{\text{right}}(f, s)} (y_i - \hat{y}_{R_{\text{right}}})^2$$

# Growing the Tree

- Finding the pair  $(f, s)$  which minimizes the quantity below can be done "easily", especially when the number of features  $d$  is not too large

$$\sum_{i: \mathbf{x}_i \in R_{\text{left}}(f, s)} (y_i - \hat{y}_{R_{\text{left}}})^2 + \sum_{i: \mathbf{x}_i \in R_{\text{right}}(f, s)} (y_i - \hat{y}_{R_{\text{right}}})^2$$

# Growing the Tree

- Finding the pair  $(f, s)$  which minimizes the quantity below can be done "easily", especially when the number of features  $d$  is not too large

$$\sum_{i: \mathbf{x}_i \in R_{\text{left}}(f, s)} (y_i - \hat{y}_{R_{\text{left}}})^2 + \sum_{i: \mathbf{x}_i \in R_{\text{right}}(f, s)} (y_i - \hat{y}_{R_{\text{right}}})^2$$

- At the next step, the same splitting strategy applies yet considering only the data falling into the previously identified regions



# Growing the Tree

- Finding the pair  $(f, s)$  which minimizes the quantity below can be done "easily", especially when the number of features  $d$  is not too large

$$\sum_{i: \mathbf{x}_i \in R_{\text{left}}(f, s)} (y_i - \hat{y}_{R_{\text{left}}})^2 + \sum_{i: \mathbf{x}_i \in R_{\text{right}}(f, s)} (y_i - \hat{y}_{R_{\text{right}}})^2$$

- At the next step, the same splitting strategy applies yet considering only the data falling into the previously identified regions
- Each time, we reduce the RSS

# When Do We Stop?

- We need a stopping criterion otherwise the tree will grow until each training instance falls into a leaf node

# When Do We Stop?

- We need a stopping criterion otherwise the tree will grow until each training instance falls into a leaf node
- Clearly, when that happens the RSS is **minimum** (in fact, it is 0!)
  - That would correspond to an overfitted tree

# When Do We Stop?

- We need a stopping criterion otherwise the tree will grow until each training instance falls into a leaf node
- Clearly, when that happens the RSS is **minimum** (in fact, it is 0!)
  - That would correspond to an overfitted tree
- Possible **stopping criteria** (tree grows until):
  - no region contains more than  $N$  observations
  - max depth of the tree is  $D$
  - RSS is reduced by at least a threshold value  $t$

# Test Time Predictions

- Once the tree is built using training instances, we will have  $J$  regions

# Test Time Predictions

- Once the tree is built using training instances, we will have  $J$  regions
- In each of them, we have a fraction of the initial training examples

# Test Time Predictions

- Once the tree is built using training instances, we will have  $J$  regions
- In each of them, we have a fraction of the initial training examples
- We can compute the mean of each region as the mean of the responses of the instances falling in that region

# Test Time Predictions

- Once the tree is built using training instances, we will have  $J$  regions
- In each of them, we have a fraction of the initial training examples
- We can compute the mean of each region as the mean of the responses of the instances falling in that region
- At test time, an unseen instance follows a root-to-leaf path on the tree and ends up into a region  $R_j$



# Test Time Predictions

- Once the tree is built using training instances, we will have  $J$  regions
- In each of them, we have a fraction of the initial training examples
- We can compute the **mean** of each region as the mean of the responses of the instances falling in that region
- At test time, an unseen instance follows a root-to-leaf path on the tree and ends up into a region  $R_j$
- The prediction for that test instance will be the **mean** of the region  $R_j$

# Take-Home Message of Today

- Decision Trees (DTs) highly **expressive** yet **interpretable** models both for regression and classification

# Take-Home Message of Today

- Decision Trees (DTs) highly **expressive** yet **interpretable** models both for regression and classification
- Learning the optimal DT is NP-Complete: **Recursive Binary Splitting** algorithm is an effective greedy training heuristic

# Take-Home Message of Today

- Decision Trees (DTs) highly **expressive** yet **interpretable** models both for regression and classification
- Learning the optimal DT is NP-Complete: **Recursive Binary Splitting** algorithm is an effective greedy training heuristic
- Regression Trees:
  - Use Residual Sum of Squares (RSS) as splitting criterion
  - At inference time, predictions are the mean of the leaf observations