

# Big Data Computing

Master's Degree in Computer Science  
2024-2025



**SAPIENZA**  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

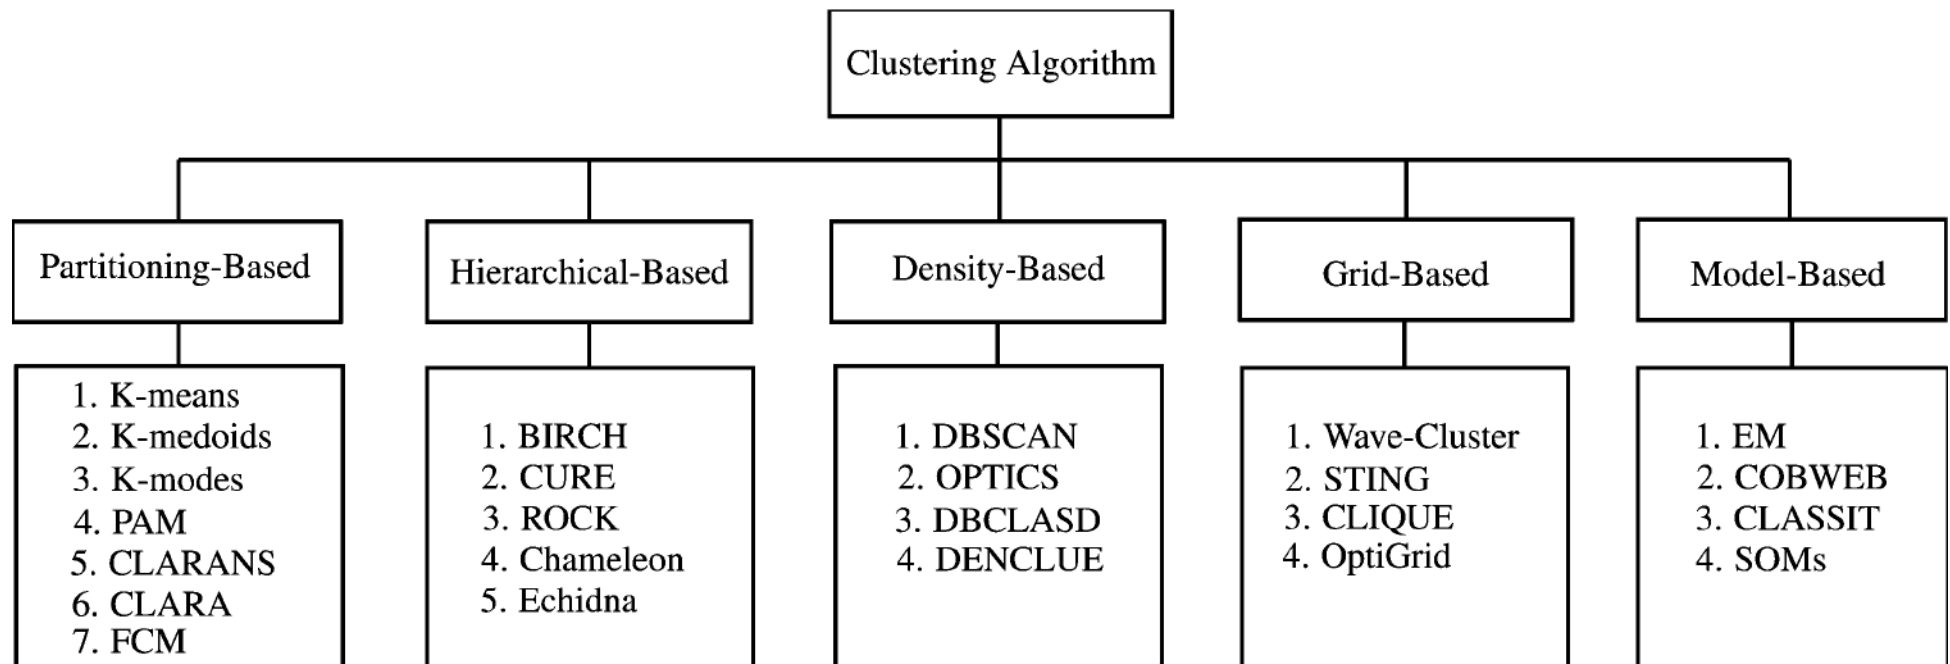
[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Recap from Last Lecture(s)

- Clustering is an unsupervised learning technique to group "similar" data objects together
- Depends on:
  - object representation
  - similarity measure
- Harder when data dimensionality gets large (**curse of dimensionality**)
- Number of output clusters is part of the problem itself!


# Clustering Algorithms

# Clustering Algorithms: Taxonomy



source: <https://www.computer.org/csdl/journal/ec/2014/03/06832486/13rRUEgs2xB>

# Clustering Algorithms: Taxonomy



source: <https://www.computer.org/csdl/journal/ec/2014/03/06832486/13rRUEgs2xB>

# Partitioning: Hard Clustering

- **Input:** A set of  $N$  data points and a number  $K$  ( $K < N$ )

# Partitioning: Hard Clustering

- **Input:** A set of  $N$  data points and a number  $K$  ( $K < N$ )
- **Output:** A partition of the  $N$  data points into  $K$  clusters

# Partitioning: Hard Clustering

- **Input:** A set of  $N$  data points and a number  $K$  ( $K < N$ )
- **Output:** A partition of the  $N$  data points into  $K$  clusters
- **Goal:** Find the partition which optimizes a certain criterion



# Partitioning: Intuition

Let  $K=2$  for simplicity

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either**  
one of the  $K=2$  clusters

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either** one of the  $K=2$  clusters

We use an  $N$ -dimensional assignment vector  $C$ ,  
where  $C[i] = \{0, 1\}$

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either** one of the  $K=2$  clusters

We use an  $N$ -dimensional assignment vector  $C$ ,  
where  $C[i] = \{0, 1\}$

Here is a possible assignment (i.e., clustering output):

	0	1	2			...				N-1
C	0	1	1	...	0	0	1	...	0	1

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either** one of the  $K=2$  clusters

We use an  $N$ -dimensional assignment vector  $C$ ,  
where  $C[i] = \{0, 1\}$

Here is another one:

	0	1	2			...			N-1
C	1	1	0	...	0	0	1	...	1

# Partitioning: Intuition

Let  $K=2$  for simplicity

We can assign each of the  $N$  data points to **either** one of the  $K=2$  clusters

We use an  $N$ -dimensional assignment vector  $C$ ,  
where  $C[i] = \{0, 1\}$

... And another one:

	0	1	2			...				N-1
C	1	1	0	...	0	1	1	...	1	0

# Partitioning: Intuition

So, how many possible clustering outputs?

# Partitioning: Intuition

So, how many possible clustering outputs?

Roughly,  $2^N$ ; More generally,  $K^N$

0	1	2		...					N-1
{0..K-1}	{0..K-1}	{0..K-1}	...	{0..K-1}	{0..K-1}	{0..K-1}	...	{0..K-1}	{0..K-1}



# Partitioning: Intuition

So, how many possible clustering outputs?

Roughly,  $2^N$ ; More generally,  $K^N$

0	1	2		...					N-1
{0..K-1}	{0..K-1}	{0..K-1}	...	{0..K-1}	{0..K-1}	{0..K-1}	...	{0..K-1}	{0..K-1}

Actually, **slightly less** than that because we want each of the  $K$  clusters to contain at least one data point!

# Partitioning: Intuition

So, how many possible clustering outputs?

Roughly,  $2^N$ ; More generally,  $K^N$

0	1	2	...	...	...	...	...	N-1
{0..K-1}	{0..K-1}	{0..K-1}	...	{0..K-1}	{0..K-1}	{0..K-1}	...	{0..K-1}

Actually, **slightly less** than that because we want each of the  $K$  clusters to contain at least one data point!


In the previous example ( $K=2$ ), the following is **not** allowed

0	1	2	...				N-1	0	1	2	...				N-1				
0	0	0	...	0	0	0	...	0	0	1	1	1	...	1	1	1	...	1	1

# Partitioning: Intuition


So, how many possible clustering outputs?

Roughly,  $2^N$ ; More generally,  $K^N$



Actually, **slightly less** than that because we want each of the K clusters to contain at least one data point!

In the previous example ( $K=2$ ), the following is **not** allowed



# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$

$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$

...

$\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$	$\{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$
$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$	$\{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$
$\vdots$	$\vdots$
$\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$	$\{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$	$\{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$	$\{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$
$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$	$\{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$	$\{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$
$\vdots$	$\vdots$	$\vdots$
$\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$	$\{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$	$\{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$



# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$	$\{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$	$\{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$
$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$	$\{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$	$\{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$
$\vdots$	$\vdots$	$\vdots$
$\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$	$\{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$	$\{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$

Some combinations are counted twice!

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$	$\{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$	$\{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$
$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$	$\{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$	$\{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$
$\vdots$	$\vdots$	$\vdots$
$\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$	$\{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$	$\{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$

Some combinations are counted twice!

# Partitioning: Intuition

Let  $D = \{1, 2, 3, 4, 5\}$  a set of  $N=5$  data points

Let  $K=2$  be the desired number of output clusters

$\{\{1\}, \{2,3,4,5\}\} \rightarrow (0,1,1,1,1)$	$\{\{1,2\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$	$\{\{2,1\}, \{3,4,5\}\} \rightarrow (0,0,1,1,1)$
$\{\{2\}, \{1,3,4,5\}\} \rightarrow (1,0,1,1,1)$	$\{\{1,3\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$	$\{\{3,1\}, \{2,4,5\}\} \rightarrow (0,1,0,1,1)$
$\vdots$	$\vdots$	$\vdots$
$\{\{5\}, \{1,2,3,4\}\} \rightarrow (1,1,1,1,0)$	$\{\{1,5\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$	$\{\{5,1\}, \{2,3,4\}\} \rightarrow (0,1,1,1,0)$

Some combinations are counted twice!

# Partitioning: NP-Hardness

- Stirling Partition Number  $\rightarrow$  K-way non-empty partitions of N elements

# Partitioning: NP-Hardness

- Stirling Partition Number  $\rightarrow$  K-way non-empty partitions of N elements

$$S(K, N) \sim K^N / K! = O(K^N)$$

# Partitioning: NP-Hardness

- Stirling Partition Number  $\rightarrow$  K-way non-empty partitions of N elements

$$S(K, N) \sim K^N / K! = O(K^N)$$

- Finding the global optimum  $\rightarrow$  Intractable for many objective function (enumerate all the possible partitions)\*

\*Kleinberg, J., "An Impossibility Theorem for Clustering" (NIPS 2002)

# Partitioning: NP-Hardness

- Stirling Partition Number  $\rightarrow$  K-way non-empty partitions of N elements

$$S(K, N) \sim K^N / K! = O(K^N)$$

- Finding the global optimum  $\rightarrow$  Intractable for many objective function (enumerate all the possible partitions)\*
- Effective heuristics  $\rightarrow$  K-means, K-medoids, K-means++, etc.

\*Kleinberg, J., "An Impossibility Theorem for Clustering" (NIPS 2002)

# Flat Hard Clustering: General Framework

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  the set of  $N$  input data points  
 $\{C_1, \dots, C_K\}$  the set of  $K$  output clusters  
 $C_k$  the generic  $k$ -th cluster  
 $\boldsymbol{\theta}_k$  is the *representative* of cluster  $C_k$



# Flat Hard Clustering: General Framework

$\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  the set of  $N$  input data points  
 $\{C_1, \dots, C_K\}$  the set of  $K$  output clusters  
 $C_k$  the generic  $k$ -th cluster  
 $\theta_k$  is the *representative* of cluster  $C_k$

## Note:

At this stage we haven't yet specified what a cluster representative actually is

# Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$$

where:

- $A$  is an  $N \times K$  matrix s.t.  $\alpha_{n,k} = 1$  iff  $\mathbf{x}_n$  is assigned to cluster  $C_k$ , 0 otherwise
- $\Theta = \{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K\}$  are the cluster representatives
- $\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$  is a function measuring the distance between  $\mathbf{x}_n$  and  $\boldsymbol{\theta}_k$

# Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$$

$\forall n \exists! k$  such that  $\alpha_{n,k} = 1 \wedge \alpha_{n,k'} = 0 \forall k' \neq k$

hard clustering

where:

- $A$  is an  $N \times K$  matrix s.t.  $\alpha_{n,k} = 1$  iff  $\mathbf{x}_n$  is assigned to cluster  $C_k$ , 0 otherwise
- $\Theta = \{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K\}$  are the cluster representatives
- $\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$  is a function measuring the distance between  $\mathbf{x}_n$  and  $\boldsymbol{\theta}_k$

# Objective Function

$$L(A, \Theta) = \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$$

$\forall n \exists! k$  such that  $\alpha_{n,k} = 1 \wedge \alpha_{n,k'} = 0 \forall k' \neq k$

hard clustering

where:

- $A$  is an  $N \times K$  matrix s.t.  $\alpha_{n,k} = 1$  iff  $\mathbf{x}_n$  is assigned to cluster  $C_k$ , 0 otherwise
- $\Theta = \{\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K\}$  are the cluster representatives
- $\delta(\mathbf{x}_n, \boldsymbol{\theta}_k)$  is a function measuring the distance between  $\mathbf{x}_n$  and  $\boldsymbol{\theta}_k$

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

# Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

# Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

exact solution must explore  
exponential search space  
 $S(K, N) \sim O(K^N)$



NP-hard

# Objective Function

$$A^*, \Theta^* = \operatorname{argmin}_{A, \Theta} \underbrace{\sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k)}_{L(A, \Theta)}$$

exact solution must explore  
exponential search space  
 $S(K, N) \sim O(K^N)$



NP-hard

non-convex due to the  
discrete assignment matrix  $A$



multiple local  
minima

# Iterative Solution: Lloyd-Forgy Algorithm

- **NP-hardness** doesn't allow us to compute the exact solution (i.e., global optimum)



# Iterative Solution: Lloyd-Forgy Algorithm

- **NP-hardness** doesn't allow us to compute the exact solution (i.e., global optimum)
- **Non-convexity** doesn't allow us to rely on nice property of convex optimization (unique global optimum)

# Iterative Solution: Lloyd-Forgy Algorithm

- **NP-hardness** doesn't allow us to compute the exact solution (i.e., global optimum)
- **Non-convexity** doesn't allow us to rely on nice property of convex optimization (unique global optimum)
- A convex objective can be (approximately) solved with numerical methods to find the global optimum

# Iterative Solution: Lloyd-Forgy Algorithm

- **Lloyd-Forgy Algorithm:** 2-step **iterative** approximated solution

# Iterative Solution: Lloyd-Forgy Algorithm

- Lloyd-Forgy Algorithm: 2-step **iterative** approximated solution
- Assignment step
- Update step

# Iterative Solution: Lloyd-Forgy Algorithm

- Lloyd-Forgy Algorithm: 2-step **iterative** approximated solution
- Assignment step
- Update step

Does not guarantee to find the global optimum as it may stuck to a local optimum or a saddle point

## 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

## 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

Note:

Can't take the gradient of  $L$  w.r.t.  $A$   
since  $A$  is discrete!

## 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

Intuitively, given a set of fixed representatives,  $L$  is minimized if each data point is assigned to the closest cluster representative according to  $\delta$

( $L$  is the sum of all the distances from each data point to its representative)



## 2-Step Optimization: Assignment Step

Minimize  $L$  w.r.t.  $A$  by fixing  $\Theta$

$L(A|\Theta) = L(A; \Theta) = L$  is a function of  $A$  parametrized by  $\Theta$

Intuitively, given a set of fixed representatives,  $L$  is minimized if each data point is assigned to the closest cluster representative according to  $\delta$

( $L$  is the sum of all the distances from each data point to its representative)

$$\alpha_{n,k} = \begin{cases} 1 & \text{if } \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) = \min_{1 \leq j \leq K} \{\delta(\mathbf{x}_n, \boldsymbol{\theta}_j)\} \\ 0 & \text{otherwise} \end{cases}$$

## 2-Step Optimization: Update Step

Minimize  $L$  w.r.t.  $\Theta$  by fixing  $A$

$L(\Theta|A) = L(\Theta; A) = L$  is a function of  $\Theta$  parametrized by  $A$

## 2-Step Optimization: Update Step

Minimize  $L$  w.r.t.  $\Theta$  by fixing  $A$

$L(\Theta|A) = L(\Theta; A) = L$  is a function of  $\Theta$  parametrized by  $A$

We can minimize  $L$  by taking the **gradient** of  $L$  w.r.t  $\Theta$  (i.e., the vector of partial derivatives), set it to 0 and solve it for  $\Theta$

## 2-Step Optimization: Update Step

$$\nabla L(\Theta; A) = \left( \frac{\partial L(\Theta; A)}{\partial \theta_1}, \dots, \frac{\partial L(\Theta; A)}{\partial \theta_K} \right)$$

## 2-Step Optimization: Update Step

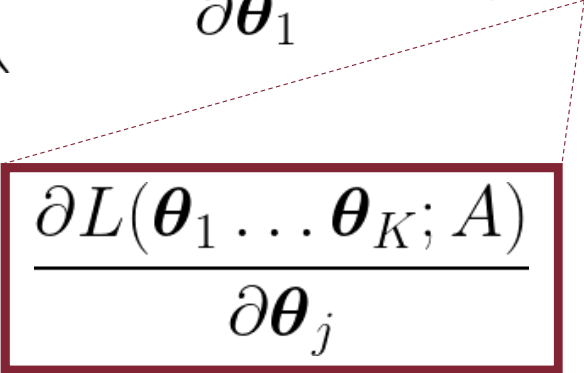
$$\nabla L(\Theta; A) = \left( \frac{\partial L(\Theta; A)}{\partial \theta_1}, \dots, \frac{\partial L(\Theta; A)}{\partial \theta_K} \right)$$

$$\nabla L(\Theta; A) = \left( \frac{\partial L(\theta_1 \dots \theta_K; A)}{\partial \theta_1}, \dots, \frac{\partial L(\theta_1 \dots \theta_K; A)}{\partial \theta_K} \right)$$

## 2-Step Optimization: Update Step

$$\nabla L(\Theta; A) = \left( \frac{\partial L(\Theta; A)}{\partial \theta_1}, \dots, \frac{\partial L(\Theta; A)}{\partial \theta_K} \right)$$

$$\nabla L(\Theta; A) = \left( \frac{\partial L(\theta_1 \dots \theta_K; A)}{\partial \theta_1}, \dots, \frac{\partial L(\theta_1 \dots \theta_K; A)}{\partial \theta_K} \right)$$


$$\frac{\partial L(\theta_1 \dots \theta_K; A)}{\partial \theta_j}$$

The general j-th partial derivative

## 2-Step Optimization: Update Step

$$\nabla L(\Theta; A) = \mathbf{0} \Leftrightarrow \frac{\partial L(\theta_1, \dots, \theta_K; A)}{\partial \theta_j} = 0 \quad \forall j \in \{1, \dots, K\}$$

## 2-Step Optimization: Update Step

$$\nabla L(\boldsymbol{\Theta}; A) = \mathbf{0} \Leftrightarrow \frac{\partial L(\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j} = 0 \quad \forall j \in \{1, \dots, K\}$$

$$\frac{\partial L(\boldsymbol{\theta}_1, \dots, \boldsymbol{\theta}_K; A)}{\partial \boldsymbol{\theta}_j} = \frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right]$$



## 2-Step Optimization: Update Step

$$\nabla L(\Theta; A) = \mathbf{0} \Leftrightarrow \frac{\partial L(\theta_1, \dots, \theta_K; A)}{\partial \theta_j} = 0 \quad \forall j \in \{1, \dots, K\}$$

$$\frac{\partial L(\theta_1, \dots, \theta_K; A)}{\partial \theta_j} = \frac{\partial}{\partial \theta_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \theta_k) \right]$$

$$\downarrow$$
$$\frac{\partial L}{\partial \theta_j}$$

To make the notation easier!

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_j} = \frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_j} = \frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

When computing the partial derivative w.r.t.  $\boldsymbol{\theta}_j$  any other term  $\boldsymbol{\theta}_k$  of the inner summation is treated as constant!

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_j} = \frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

$$= \frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \alpha_{n,j} \delta(\mathbf{x}_n, \boldsymbol{\theta}_j) \right] = 0$$

## 2-Step Optimization: Update Step

$$\frac{\partial L}{\partial \boldsymbol{\theta}_j} = \frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \sum_{k=1}^K \alpha_{n,k} \delta(\mathbf{x}_n, \boldsymbol{\theta}_k) \right] = 0$$

$$= \frac{\partial}{\partial \boldsymbol{\theta}_j} \left[ \sum_{n=1}^N \alpha_{n,j} \delta(\mathbf{x}_n, \boldsymbol{\theta}_j) \right] = 0$$

Solve for each  $\boldsymbol{\theta}_j$   
independently

Depends on the distance  
function  $\delta$

# Take-Home Message of Today

- Focus on hard partitioning clustering

# Take-Home Message of Today

- Focus on hard partitioning clustering
- Formulate hard partitioning clustering as a (**non-convex**) optimization problem
  - Minimizing “some” aggregated internal cluster distance

# Take-Home Message of Today

- Focus on hard partitioning clustering
- Formulate hard partitioning clustering as a (**non-convex**) optimization problem
  - Minimizing “some” aggregated internal cluster distance
- Computing exact solution is **NP-hard** due to exponential search space



# Take-Home Message of Today

- Focus on hard partitioning clustering
- Formulate hard partitioning clustering as a (**non-convex**) optimization problem
  - Minimizing “some” aggregated internal cluster distance
- Computing exact solution is **NP-hard** due to exponential search space
- Use an iterative (approximate) solution