

# Big Data Computing

Master's Degree in Computer Science  
2023-2024

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)



SAPIENZA  
UNIVERSITÀ DI ROMA

# Who Am I?



# Who Am I?



UniPI  
(1999-2005)



# Who Am I?



UniPI  
(1999-2005)



UniVE  
(2008-2013)

# Who Am I?



UniPI  
(1999-2005)



UniVE  
(2008-2013)



Yahoo! Labs  
(2014-2017)

11/21/23



# Who Am I?



UniPI  
(1999-2005)



UniVE  
(2008-2013)



11/21/23  
Yahoo! Labs  
(2014-2017)



UniPD  
(2017-2019)

# Who Am I?



UniPI  
(1999-2005)



UniVE  
(2008-2013)



Yahoo! Labs  
(2014-2017)  
11/21/23




UniPD  
(2017-2019)



Sapienza  
(2019-)

# My Research Interests




HERCOLE Lab



# My Research Interests

Human-Explainable




HERCOLE Lab

# My Research Interests

Robust

Human-Explainable




HERCOLÉ Lab

# My Research Interests

Robust

Human-Explainable

COLlaborative




HERCOL Lab

# My Research Interests

Robust

Human-Explainable

COllaborative

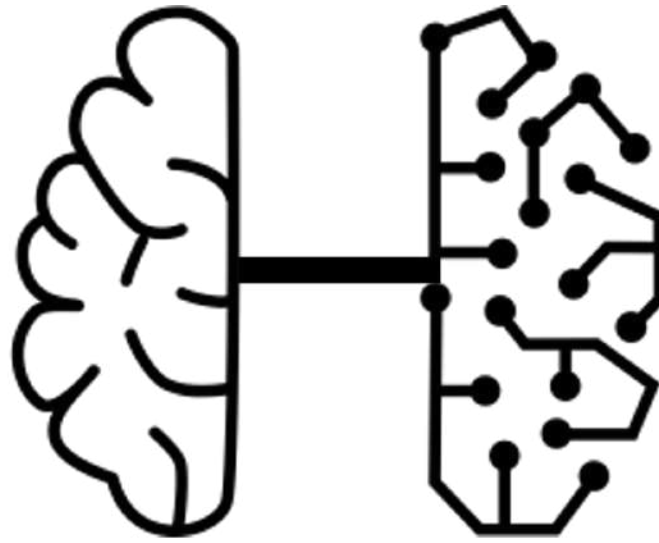


HERCOLE Lab

LEarning

# My Research Interests

Sounds cool?

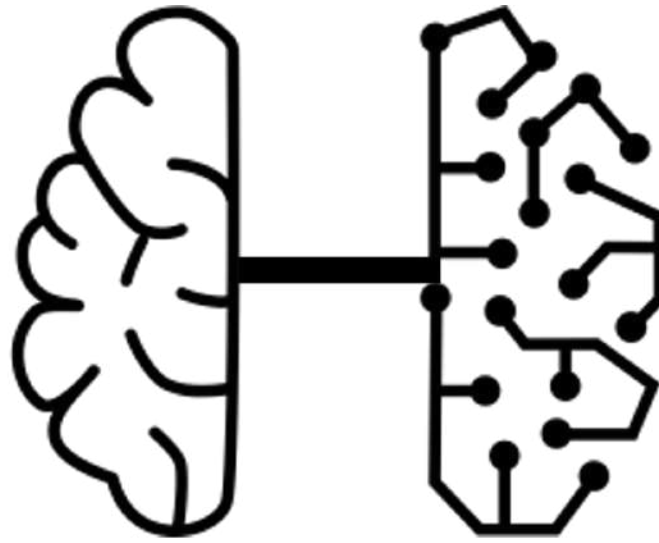


HERCOLE Lab



# My Research Interests

Sounds cool?




HERCOLE Lab

Check out the  
lab's [home](#)  
[page](#)  
(still under  
construction, sic!)



# My Research Interests

Sounds cool?



HERCOLE Lab

Meanwhile you  
can follow us  
on Twitter  
[@HercoleLab](https://twitter.com/HercoleLab)

# Administrivia

- Class schedule:
  - Tuesday from 2:00 p.m. to 4:00 p.m.
  - Wednesday from 10:00 a.m. to 1:00 p.m.

# Administrivia

- Class schedule:

- Tuesday from 2:00 p.m. to 4:00 p.m.
- Wednesday from 10:00 a.m. to 1:00 p.m.

Aula Magna @  
Viale Regina Elena, 295

# Administrivia

- Class schedule:

- Tuesday from 2:00 p.m. to 4:00 p.m.
- Wednesday from 10:00 a.m. to 1:00 p.m.

Aula Magna @  
Viale Regina Elena, 295

- Office hours:

- Drop me a message to ask for a meeting **online** (Google Meet or Zoom) or in-person at my office (Room 106 @ Viale Regina Elena, 295 – 1st Floor, Building E)



# Administrivia

- Contacts:

- Personal homepage: <https://www.di.uniroma1.it/~tolomei>
- Email: [tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Administrivia

- Resources:

- Course's website: <https://github.com/gtolomei/big-data-computing>
- Moodle's web page: <https://elearning.uniroma1.it/course/view.php?id=16942>

# Administrivia

- Resources:
  - Course's website: <https://github.com/gtolomei/big-data-computing>
  - Moodle's web page: <https://elearning.uniroma1.it/course/view.php?id=16942>
- Class material will be published on the course's website **only**
  - Along with other resources (e.g., suggested readings/books) if needed

# Administrivia

- Prerequisites:
  - Familiarity with basics of Data Science and Machine Learning
  - Solid knowledge of Calculus, Linear Algebra, and Probability&Statistics
  - (Python) Programming skills desirable yet not mandatory!

# Administrivia

- Prerequisites:
  - Familiarity with basics of Data Science and Machine Learning
  - Solid knowledge of Calculus, Linear Algebra, and Probability&Statistics
  - (Python) Programming skills desirable yet not mandatory!

No worries!

Many subjects will be anyway revisited during class lectures



# Administrivia

- Exam:
  - A seminar on a research paper

# Administrivia

- Exam:
  - A seminar on a research paper
  - The topic of the seminar must align with the spirit of this unit, e.g.:
    - Dealing with high-dimensional data
    - Search/Retrieve/Filter relevant data from large collections of items
    - Optimization/Learning

# Administrivia

- Exam:
  - A seminar on a research paper
  - The topic of the seminar must align with the spirit of this unit, e.g.:
    - Dealing with high-dimensional data
    - Search/Retrieve/Filter relevant data from large collections of items
    - Optimization/Learning
  - We will give you further details in the upcoming weeks!

Questions?

# Outline of the Course

Big Data Infrastructure



# Outline of the Course

Big Data Infrastructure

HDFS

Hadoop

# Outline of the Course


## Big Data Infrastructure

HDFS

Hadoop

Spark

# Outline of the Course



# Outline of the Course

Big Data Problems and Applications

Big Data Infrastructure

HDFS

Hadoop

Spark

...

# Outline of the Course

Big Data Problems and Applications

Similarity in High-Dimensional Spaces

Big Data Infrastructure

HDFS

Hadoop

Spark

...

# Outline of the Course

## Big Data Problems and Applications

Clustering

k-NN

...

Similarity in High-Dimensional Spaces

## Big Data Infrastructure

HDFS

Hadoop

Spark

...

# Outline of the Course

## Big Data Problems and Applications

Clustering

k-NN

...

Similarity in High-Dimensional Spaces

Dimensionality Reduction

## Big Data Infrastructure

HDFS

Hadoop

Spark

...

# Outline of the Course

## Big Data Problems and Applications

Clustering

k-NN

...

PCA

SVD

...

Similarity in High-Dimensional Spaces

Dimensionality Reduction

## Big Data Infrastructure

HDFS

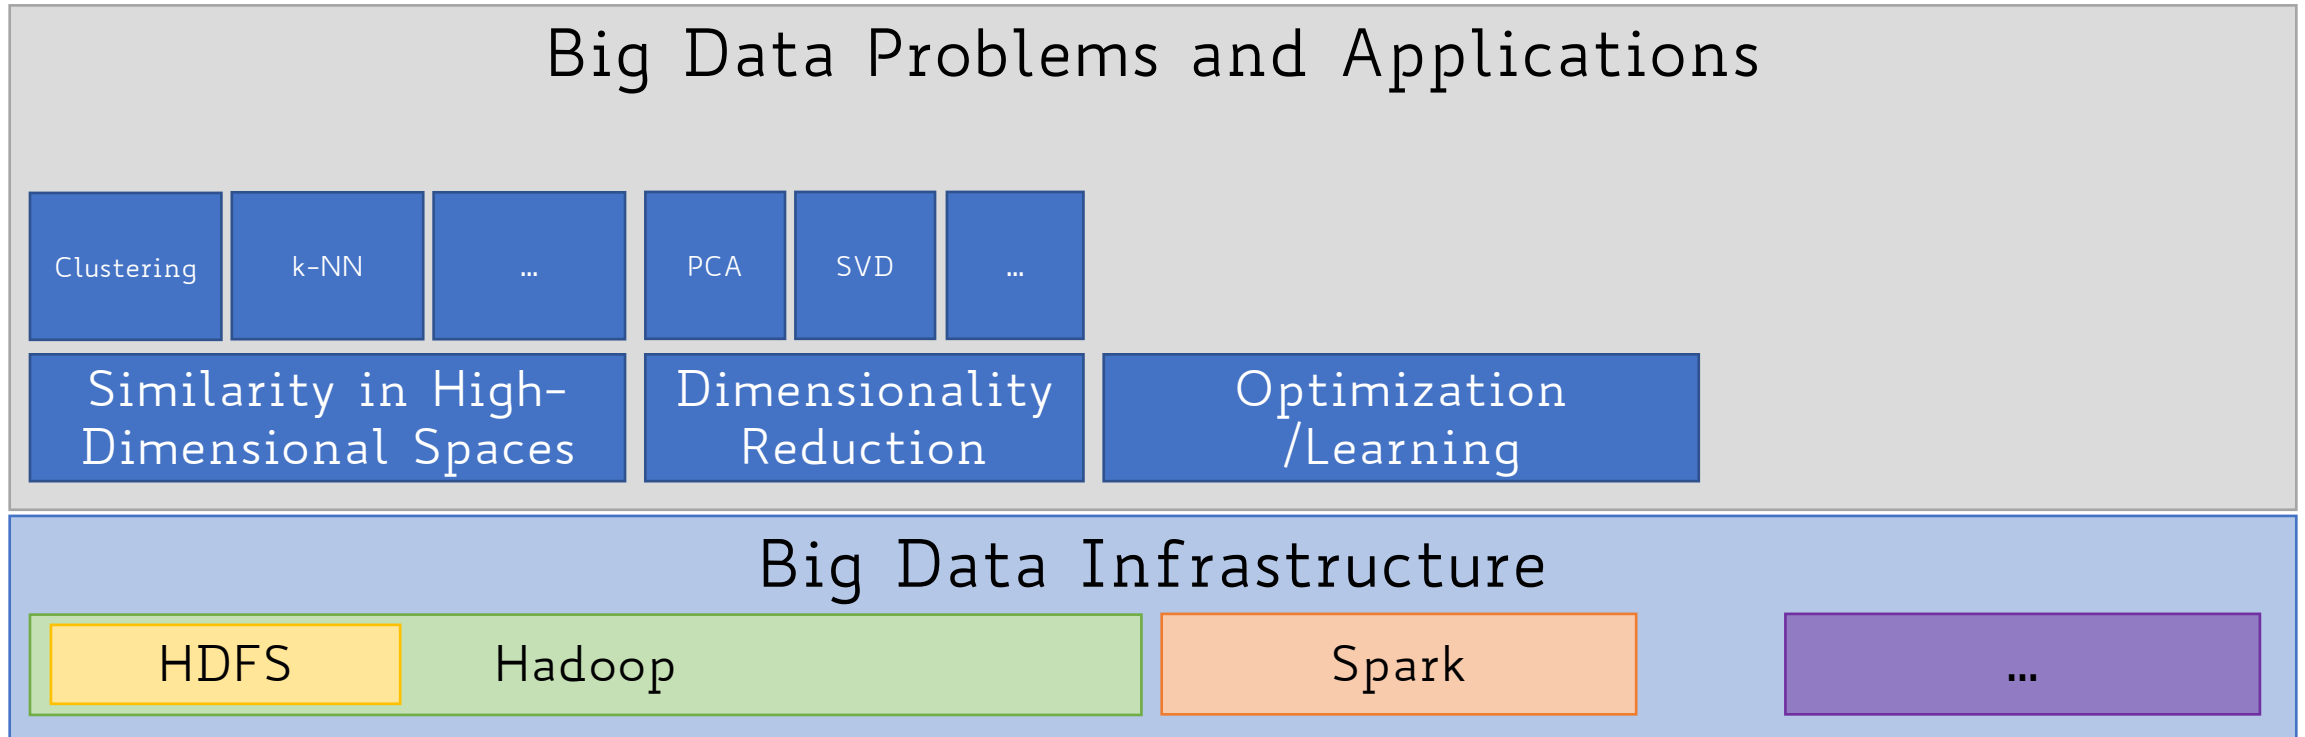
Hadoop

Spark

...



# Outline of the Course



# Outline of the Course

## Big Data Problems and Applications

Clustering

k-NN

...

PCA

SVD

...

SGD

Matrix  
Factorization

...

Similarity in High-  
Dimensional Spaces

Dimensionality  
Reduction

Optimization  
/Learning

## Big Data Infrastructure

HDFS

Hadoop

Spark

...

# Outline of the Course

## Big Data Problems and Applications

Clustering

k-NN

...

PCA

SVD

...

SGD

Matrix  
Factorization

...

Similarity in High-  
Dimensional Spaces

Dimensionality  
Reduction

Optimization  
/Learning

Graph  
Analysis

## Big Data Infrastructure

HDFS

Hadoop

Spark

...

# Outline of the Course

## Big Data Problems and Applications

Clustering

k-NN

...

PCA

SVD

...

SGD

Matrix  
Factorization

...

PageRank

...

...

Similarity in High-  
Dimensional Spaces

Dimensionality  
Reduction

Optimization  
/Learning

Graph  
Analysis

...

## Big Data Infrastructure

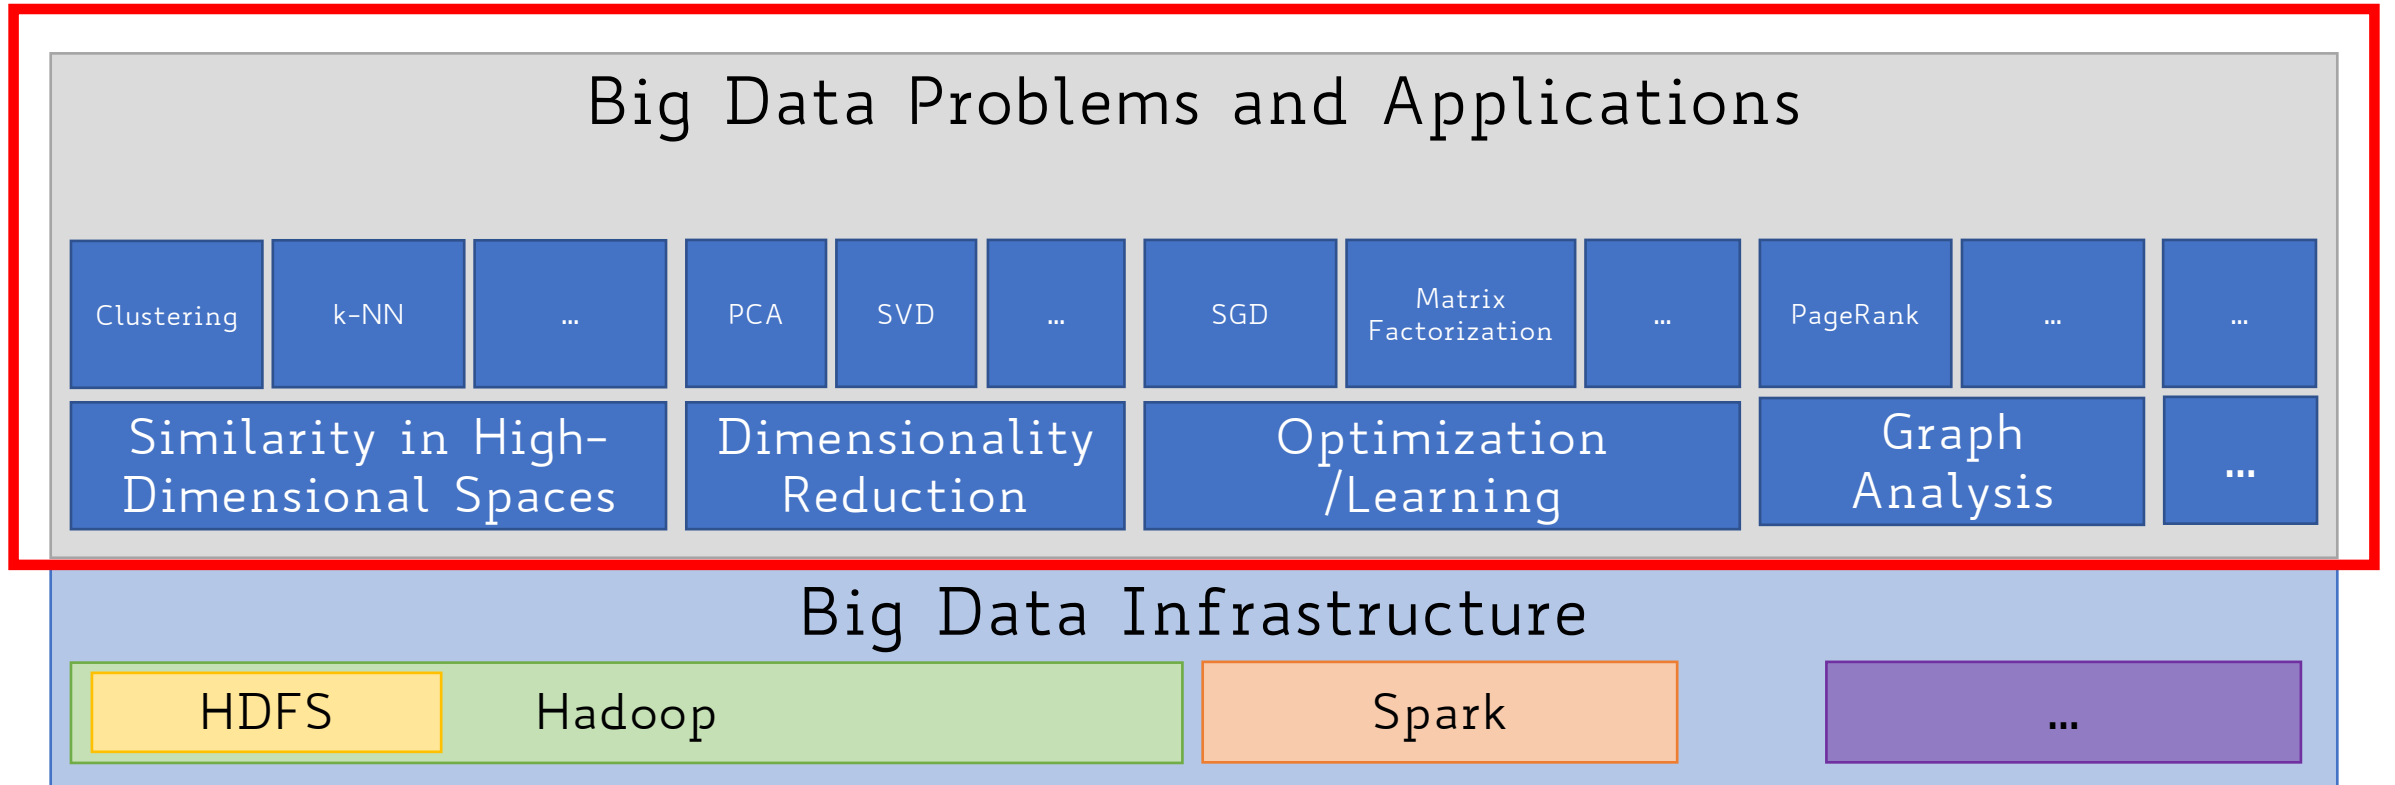
HDFS

Hadoop

Spark

...

# Outline of the Course



# Let's Get Started!

# The Search Problem

Given a collection  $S$  of  $N$  items, find if  $x \in S$

# The Search Problem

Given a collection  $S$  of  $N$  items, find if  $x \in S$

This is a very general problem that occurs frequently



# The Search Problem

Given a collection  $S$  of  $N$  items, find if  $x \in S$

This is a very general problem that occurs frequently

## Example:

Find if a number occurs in a list of  $N$  integers

# The Search Problem

If the list is **not** sorted, it takes  $O(N)$  steps to respond

# The Search Problem

If the list is **not** sorted, it takes  $O(N)$  steps to respond

If the list is sorted, it takes  $O(\log(N))$  steps to respond  
but we must pay the cost of sorting it  $O(N\log(N))$

# The Search Problem

If the list is **not** sorted, it takes  $O(N)$  steps to respond

If the list is sorted, it takes  $O(\log(N))$  steps to respond  
but we must pay the cost of sorting it  $O(N\log(N))$

Still, it might be beneficial to pre-sort the list if we  
repeat the find operation several times...

# Time Complexity

Let  $M$  be the number of times the find operation is called

# Time Complexity

Let  $M$  be the number of times the find operation is called

Without sorting, the total complexity is  $O(M*N)$  steps

# Time Complexity

Let  $M$  be the number of times the find operation is called

Without sorting, the total complexity is  $O(M*N)$  steps

With sorting, the complexity becomes  $O(N\log(N)) + O(M\log(N))$

# Time Complexity

Let  $M$  be the number of times the find operation is called

**Without sorting**, the total complexity is  $O(M*N)$  steps

**With sorting**, the complexity becomes  $O(N\log(N)) + O(M\log(N))$

We must find the value of  $M$  for which:

$$O(N\log(N)) + O(M\log(N)) < O(M*N)$$

with sorting

without sorting



# Time Complexity Trade-Off

Brutally:

$$N\log(N) + M\log(N) < MN$$

# Time Complexity Trade-Off

Brutally:

$$N\log(N) + M\log(N) < MN$$

$$M\log(N) - MN < -N\log(N)$$

# Time Complexity Trade-Off

Brutally:

$$N\log(N) + M\log(N) < MN$$

$$M\log(N) - MN < -N\log(N)$$

$$MN - M\log(N) > N\log(N)$$

# Time Complexity Trade-Off

Brutally:

$$N\log(N) + M\log(N) < MN$$

$$M\log(N) - MN < -N\log(N)$$

$$MN - M\log(N) > N\log(N)$$

$$M(N - \log(N)) > N\log(N)$$

# Time Complexity Trade-Off

Brutally:

$$N\log(N) + M\log(N) < MN$$

$$M\log(N) - MN < -N\log(N)$$

$$MN - M\log(N) > N\log(N)$$

$$M(N - \log(N)) > N\log(N)$$

$$M \geq \lceil N\log(N)/(N - \log(N)) \rceil \quad [\text{assuming } N \in \mathbb{Z}_{>0}]$$

# Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$

# Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$
1000			

# Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$
1000	2		



# Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$
1000	2	2000	

# Time Complexity Trade-Off: Example

N	M	$M*N$	$N\log N + M\log N$
1000	2	2000	$1000*3 + 2*3 = 3006$

# Time Complexity Trade-Off: Example

N	M	M*N	$N\log N + M\log N$
1000	2	2000	$1000*3 + 2*3 = 3006$

# Time Complexity Trade-Off: Example

N	M	M*N	$N\log N + M\log N$
1000	2	2000	$1000*3 + 2*3 = 3006$
1000	3	3000	$1000*3 + 3*3 = 3009$
1000	4	4000	$1000*3 + 4*3 = 3012$

# Time Complexity Trade-Off: Example

N	M	M*N	NlogN + MlogN
1000	2	2000	1000*3 + 2*3 = 3006
1000	3	3000	1000*3 + 3*3 = 3009
1000	4	4000	1000*3 + 4*3 = 3012

$$M \geq \lceil N \log(N) / (N - \log(N)) \rceil \text{ [assuming } N \in \mathbb{Z}_{>0}]$$

$$M \geq \lceil 1000 * \log(1000) / (1000 - \log(1000)) \rceil$$

$$M \geq \lceil 3000 / (1000 - 3) \rceil$$

$$M \geq \lceil 3.009 \rceil = 4$$

# 'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

# 'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If  $N$  grows large and the list is **not** sorted, a linear scan can be too costly

# 'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If  $N$  grows large and the list is **not** sorted, a linear scan can be too costly

On the other hand, if  $N$  **does not fit into main memory**, we will need some external  $R$ -way merge-sorting



# 'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If  $N$  grows large and the list is **not** sorted, a linear scan can be too costly

On the other hand, if  $N$  **does not fit into main memory**, we will need some external  $R$ -way merge-sorting

In both cases, the complexity may arise from the fact that we are dealing with a massive 'vertical' input size  $N$

# 'Vertical' vs. 'Horizontal' Scale

What could make the search problem hard?

If  $N$  grows large and the list is **not** sorted, a linear scan can be too costly

On the other hand, if  $N$  **does not fit into main memory**, we will need some external  $R$ -way merge-sorting

In both cases, the complexity may arise from the fact that we are dealing with a massive 'vertical' input size  $N$

'vertical' here means that the **number**  $N$  of input data points is large, yet their **representation** (e.g., 8-byte integers) is not!

# 'Vertical' vs. 'Horizontal' Scale

Suppose that  $S$  is a collection of *images*, not integers

# 'Vertical' vs. 'Horizontal' Scale

Suppose that  $S$  is a collection of `images`, not integers

Our goal is to find if a given image  $x$  is in  $S$   
(or retrieve an image  $x'$  in  $S$  that is 'most similar' to  $x$ )

# 'Vertical' vs. 'Horizontal' Scale

Suppose that  $S$  is a collection of *images*, not integers

Our goal is to find if a given image  $x$  is in  $S$   
(or retrieve an image  $x'$  in  $S$  that is 'most similar' to  $x$ )

We may have to cope both with the large number of items  $N$  in the collection *and* the complex representation of each item

# 'Vertical' vs. 'Horizontal' Scale

Suppose that  $S$  is a collection of **images**, not integers

Our goal is to find if a given image  $x$  is in  $S$   
(or retrieve an image  $x'$  in  $S$  that is 'most similar' to  $x$ )

We may have to cope both with the large number of items  $N$  in the collection **and** the complex representation of each item

Each image must be represented by a high-dimensional vector  
of RGB pixels

# 'Vertical' vs. 'Horizontal' Scale

Suppose that  $S$  is a collection of **images**, not integers

Our goal is to find if a given image  $x$  is in  $S$   
(or retrieve an image  $x'$  in  $S$  that is 'most similar' to  $x$ )

We may have to cope both with the large number of items  $N$  in the collection **and** the complex representation of each item

Each image must be represented by a high-dimensional vector  
of RGB pixels

For example, a 100x100 pixel image requires a 30000-dimensional  
real-value vector

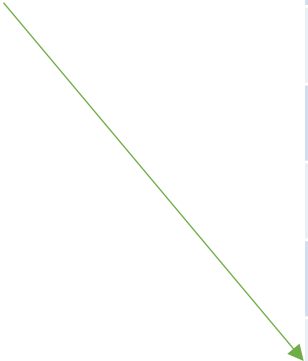
# 'Vertical' vs. 'Horizontal' Scale

27
1290
-45
...
32
-670
...
42
...
810



# 'Vertical' vs. 'Horizontal' Scale

`find(42)`



A vertical array of numbers, each in a light blue box. The numbers are: 27, 1290, -45, ..., 32, -670, ..., 42, ..., 810. A green arrow points from the text 'find(42)' to the box containing '42'.

27
1290
-45
...
32
-670
...
42
...
810

# 'Vertical' vs. 'Horizontal' Scale

`find(42)`

27
1290
-45
...
32
-670
...
42
...
810


**vertical**

# 'Vertical' vs. 'Horizontal' Scale

find(42)

27
1290
-45
...
32
-670
...
42
...
810

vertical



...



...




# 'Vertical' vs. 'Horizontal' Scale

`find(42)`

27
1290
-45
...
32
-670
...
42
...
810

**vertical**



...



...




`find(`  `)`

# 'Vertical' vs. 'Horizontal' Scale

`find(42)`

27
1290
-45
...
32
-670
...
42
...
810

**vertical**




...



...



`find()`



**horizontal**

# 'Vertical' vs. 'Horizontal' Scale

We can experience **both** types of scalability issues

# 'Vertical' vs. 'Horizontal' Scale

We can experience **both** types of scalability issues

Very large **number** of input data points  
(vertical)

**AND**

Very large **representation** of input data points  
(horizontal)

# 'Vertical' vs. 'Horizontal' Scale

We can experience **both** types of scalability issues

Very large **number** of input data points  
(vertical)

**AND**

Very large **representation** of input data points  
(horizontal)

We will focus mostly on the second and still assume we deal with very large number of input data points



# Similarity Measures

- What does "similar" mean?

# Similarity Measures

- What does "**similar**" mean?
- No single answer! It depends on what we want to find or emphasize in the data

# Similarity Measures

- What does "**similar**" mean?
- No single answer! It depends on what we want to find or emphasize in the data
- Domain and representation specific (e.g., similar images vs. similar text documents)

# Similarity Measures

- What does "**similar**" mean?
- No single answer! It depends on what we want to find or emphasize in the data
- Domain and representation specific (e.g., similar images vs. similar text documents)
- Crucial for several big data tasks (e.g., search/retrieval/filter, clustering, classification, etc.)

# Notion of Similarity

- We implicitly assumed data live in a  $d$ -dimensional Euclidean space

# Notion of Similarity

- We implicitly assumed data live in a  $d$ -dimensional **Euclidean space**
- Similarity between data is computed using:
  - **Euclidean metric** (i.e., distance)
  - **Cosine similarity**
  - **Jaccard coefficient**
  - ...

# Metric and Metric Space

$X$  is a set

$\delta$  is a function  $\delta : X \times X \rightarrow [0, \infty)$ , where:

1.  $\delta(x, y) \geq 0$  (**non-negativity**)
2.  $\delta(x, y) = 0 \Leftrightarrow x = y$  (**identity** of indiscernibles)
3.  $\delta(x, y) = \delta(y, x)$  (**symmetry**)
4.  $\delta(x, y) \leq \delta(x, z) + \delta(z, y)$  (**triangle inequality**)

Then  $\delta$  is called a **metric** (or distance function) and  $X$  a **metric space**

# Euclidean Metric (Distance) & Euclidean Space

$$X = \mathbb{R}^d$$

$$\delta : \mathbb{R}^d \times \mathbb{R}^d \rightarrow [0, \infty)$$

$\mathbf{x} = (x_1, \dots, x_d)$  and  $\mathbf{y} = (y_1, \dots, y_d)$  are 2 points in  $\mathbb{R}^d$

$$\delta(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + \dots + (x_d - y_d)^2} = \sqrt{\sum_{i=1}^d (x_i - y_i)^2}$$



# Euclidean Norm ( $L^2$ -Norm)

- The position of a point in a Euclidean  $d$ -space is a **Euclidean vector**

# Euclidean Norm ( $L^2$ -Norm)

- The position of a point in a Euclidean  $d$ -space is a **Euclidean vector**
- The **Euclidean norm** of a vector measures its length (from the origin)

# Euclidean Norm (L<sup>2</sup>-Norm)

- The position of a point in a Euclidean  $d$ -space is a **Euclidean vector**
- The **Euclidean norm** of a vector measures its length (from the origin)

$$||\mathbf{x}||_2 = \sqrt{x_1^2 + \dots + x_d^2} = \sqrt{\mathbf{x} \cdot \mathbf{x}}$$

where  $\cdot$  indicates the **dot product**

# Euclidean Norm (L<sup>2</sup>-Norm)

- The position of a point in a Euclidean  $d$ -space is a **Euclidean vector**
- The **Euclidean norm** of a vector measures its length (from the origin)

$$||\mathbf{x}||_2 = \sqrt{x_1^2 + \dots + x_d^2} = \sqrt{\mathbf{x} \cdot \mathbf{x}}$$

where  $\cdot$  indicates the **dot product**

This can be just seen as the Euclidean distance between vector's tail and tip

# Euclidean Norm & Euclidean Metric

Let  $\mathbf{x}-\mathbf{y} = (x_1-y_1, \dots, x_d-y_d)$  the **displacement vector** between  $\mathbf{x}$  and  $\mathbf{y}$

# Euclidean Norm & Euclidean Metric

Let  $\mathbf{x}-\mathbf{y} = (x_1-y_1, \dots, x_d-y_d)$  the **displacement vector** between  $\mathbf{x}$  and  $\mathbf{y}$

The Euclidean distance between  $\mathbf{x}$  and  $\mathbf{y}$  is just the Euclidean norm of the displacement vector

# Euclidean Norm & Euclidean Metric

Let  $\mathbf{x} - \mathbf{y} = (x_1 - y_1, \dots, x_d - y_d)$  the **displacement vector** between  $\mathbf{x}$  and  $\mathbf{y}$

The Euclidean distance between  $\mathbf{x}$  and  $\mathbf{y}$  is just the Euclidean norm of the displacement vector

$$\delta(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|_2 = \sqrt{(\mathbf{x} - \mathbf{y}) \cdot (\mathbf{x} - \mathbf{y})}$$

# Euclidean Distance: 1-dimensional Case

$$d = 1$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}$$

$\mathbf{x} = x, \mathbf{y} = y$  both  $\mathbf{x}$  and  $\mathbf{y}$  are scalars

$$\delta(\mathbf{x}, \mathbf{y}) = \delta(x, y) = \sqrt{(x - y)^2} = |x - y|$$



# Euclidean Distance: 1-dimensional Case

$$d = 1$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}$$

$\mathbf{x} = x, \mathbf{y} = y$  both  $\mathbf{x}$  and  $\mathbf{y}$  are scalars

$$\delta(\mathbf{x}, \mathbf{y}) = \delta(x, y) = \sqrt{(x - y)^2} = |x - y|$$

The Euclidean distance between any two 1-d points on the real line is the **absolute value** of the numerical difference of their coordinates

# Euclidean Distance: 2-dimensional Case

$$d = 2$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}^2$$

$$\mathbf{x} = (x_1, x_2), \mathbf{y} = (y_1, y_2)$$

$$\delta(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2} = \|\mathbf{x} - \mathbf{y}\|_2$$

# Euclidean Distance: 2-dimensional Case

$$d = 2$$

$$\mathbf{x}, \mathbf{y} \in \mathbb{R}^d = \mathbb{R}^2$$

$$\mathbf{x} = (x_1, x_2), \mathbf{y} = (y_1, y_2)$$

$$\delta(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2} = \|\mathbf{x} - \mathbf{y}\|_2$$

The Euclidean distance between any two 2-d points on the Euclidean plane equals to the **Pythagorean theorem**

# Minkowski Distance ( $L^p$ -Norm)

Generalization of the Euclidean distance

$$\mathbf{x} = (x_1, \dots, x_d) \text{ and } \mathbf{y} = (y_1, \dots, y_d) \in \mathbb{R}^d$$

$$\delta_p(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^d |x_i - y_i|^p \right)^{\frac{1}{p}}$$

# Minkowski Distance ( $L^p$ -Norm): $p=1$

$L^1$ -Norm or Manhattan Distance

$$\delta_1(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^d |x_i - y_i|^1 \right)^{\frac{1}{1}} = \sum_{i=1}^d |x_i - y_i|$$

# Minkowski Distance ( $L^p$ -Norm): $p=2$

$L^2$ -Norm or Euclidean Distance

$$\delta_2(\mathbf{x}, \mathbf{y}) = \left( \sum_{i=1}^d |x_i - y_i|^2 \right)^{\frac{1}{2}} = \sqrt{\sum_{i=1}^d |x_i - y_i|^2}$$

# Minkowski Distance ( $L^p$ -Norm): $p=\infty$

$L^\infty$ -Norm or Chebyshev Distance

$$\begin{aligned}\delta_\infty(\mathbf{x}, \mathbf{y}) &= \lim_{p \rightarrow \infty} \left( \sum_{i=1}^d |x_i - y_i|^p \right)^{\frac{1}{p}} = \\ &= \max\{|x_1 - y_1|, |x_2 - y_2|, \dots, |x_d - y_d|\}\end{aligned}$$

# Cosine Similarity

- A measure of similarity between two non-zero vectors of an inner product space



# Cosine Similarity

- A measure of similarity between two non-zero vectors of an inner product space
- Measures the **cosine of the angle** between vectors


# Cosine Similarity

- A measure of similarity between two non-zero vectors of an inner product space
- Measures the **cosine of the angle** between vectors
- It ranges between  $[-1,1]$

# Cosine Similarity


- A measure of similarity between two non-zero vectors of an inner product space
- Measures the **cosine of the angle** between vectors
- It ranges between  $[-1,1]$
- It captures the **orientation** and not the magnitude

# Cosine Similarity




$\theta$  is close to  $0^\circ$   
 $\cos(\theta) \approx 1$

similar vectors



$\theta$  is close to  $90^\circ$   
 $\cos(\theta) \approx 0$


orthogonal vectors



$\theta$  is close to  $180^\circ$   
 $\cos(\theta) \approx -1$

opposite vectors

# Cosine Similarity: 2-dimensional Case




$$\theta = \beta - \alpha$$

$$x = (\underbrace{\|x\|\cos\alpha}_{x_1}, \underbrace{\|x\|\sin\alpha}_{x_2})$$

$$y = (\underbrace{\|y\|\cos\beta}_{y_1}, \underbrace{\|y\|\sin\beta}_{y_2})$$

# Cosine Similarity: 2-dimensional Case


$$x \cdot y = x_1 y_1 + x_2 y_2 =$$




# Cosine Similarity: 2-dimensional Case

$$\mathbf{x} \cdot \mathbf{y} = x_1 y_1 + x_2 y_2 =$$

$$= \|\mathbf{x}\| \cos \alpha \|\mathbf{y}\| \cos \beta + \|\mathbf{x}\| \sin \alpha \|\mathbf{y}\| \sin \beta$$



# Cosine Similarity: 2-dimensional Case




$$x \cdot y = x_1 y_1 + x_2 y_2 =$$

$$= \|x\| \cos \alpha \|y\| \cos \beta + \|x\| \sin \alpha \|y\| \sin \beta$$

$$= \|x\| \|y\| (\cos \alpha \cos \beta + \sin \alpha \sin \beta)$$




# Cosine Similarity: 2-dimensional Case



$$\begin{aligned}
 x \cdot y &= x_1 y_1 + x_2 y_2 = \\
 &= \|x\| \cos \alpha \|y\| \cos \beta + \|x\| \sin \alpha \|y\| \sin \beta \\
 &= \|x\| \|y\| (\underbrace{\cos \alpha \cos \beta + \sin \alpha \sin \beta}_{\cos(\beta - \alpha)}) \\
 &\quad \underbrace{\hspace{1.5cm}}_{\theta}
 \end{aligned}$$

# Cosine Similarity: 2-dimensional Case



$$x \cdot y = x_1 y_1 + x_2 y_2 =$$

$$= \|x\| \cos \alpha \|y\| \cos \beta + \|x\| \sin \alpha \|y\| \sin \beta$$


$$= \|x\| \|y\| (\cos \alpha \cos \beta + \sin \alpha \sin \beta)$$

$$\cos(\beta - \alpha)$$

$$\theta$$


$$x \cdot y = \|x\| \|y\| \cos \theta$$

# Cosine Similarity: 2-dimensional Case



$$x \cdot y = \|x\| \|y\| \cos \theta$$

# Cosine Similarity: 2-dimensional Case



$$x \cdot y = \|x\| \|y\| \cos \theta$$



$$\cos \theta = x \cdot y / \|x\| \|y\|$$

# Cosine Similarity: $d$ -dimensional Case

- Computed as in the case of 2-dimensional vectors

# Cosine Similarity: $d$ -dimensional Case

- Computed as in the case of 2-dimensional vectors
- If two  $d$ -dimensional vectors are not collinear then they span a 2-dimensional plane  $E \subset \mathbb{R}^d$

# Cosine Similarity: $d$ -dimensional Case

- Computed as in the case of 2-dimensional vectors
- If two  $d$ -dimensional vectors are not collinear then they span a 2-dimensional plane  $E \subset \mathbb{R}^d$
- This plane  $E$  inherits the dot product in  $\mathbb{R}^d$  and so becomes an ordinary Euclidean plane

# Cosine Similarity: $d$ -dimensional Case

- Computed as in the case of 2-dimensional vectors
- If two  $d$ -dimensional vectors are not collinear then they span a 2-dimensional plane  $E \subset \mathbb{R}^d$
- This plane  $E$  inherits the dot product in  $\mathbb{R}^d$  and so becomes an ordinary Euclidean plane
- The angles in this plane are related to the dot product as they are in 2-dimensional vector geometry



# Jaccard Index (Coefficient)

Measures similarity between finite sample sets

# Jaccard Index (Coefficient)


Measures similarity between finite sample sets

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}$$

$$J(A, B) = 1 \text{ if } A = B = \emptyset$$

$$0 \leq J(A, B) \leq 1$$

# Jaccard Index (Coefficient): Interpretation



source: [Wikipedia](#)

# Jaccard Distance

Complementary to the Jaccard coefficient

$$\delta_J(A, B) = 1 - J(A, B) = \frac{|A \cup B| - |A \cap B|}{|A \cup B|}$$

This distance is a **metric** on the collection of all finite sets

# Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
  - number of input data points vs. high-dimensional representation

# Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
  - number of input data points vs. high-dimensional representation
- Many big data tasks require computing "similarity" between domain items

# Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
  - number of input data points vs. high-dimensional representation
- Many big data tasks require computing "similarity" between domain items
- Different similarity measures:
  - Euclidean, Cosine, Jaccard, etc.

# Take-Home Message of Today

- Vertical vs. Horizontal scale of data:
  - number of input data points vs. high-dimensional representation
- Many big data tasks require computing "similarity" between domain items
- Different similarity measures:
  - Euclidean, Cosine, Jaccard, etc.
- We'll see the notion of similarity can be flawed in high-dimensional spaces